

BigQuery_Pipeline

July 1, 2025

1 Configuration and Imports

```
[1]: from google.colab import auth
      auth.authenticate_user()
```

WARNING: google.colab.auth.authenticate_user() is not supported in Colab Enterprise.

```
[2]: import bigframes as bf
      import bigframes.pandas as bpd
      import matplotlib.pyplot as plt
      import seaborn as sns
      import pandas as pd
      import numpy as np
      import time
      from sklearn.model_selection import train_test_split
      from sklearn.utils.class_weight import compute_sample_weight
      from sklearn.utils.multiclass import type_of_target
      from sklearn.preprocessing import StandardScaler, OneHotEncoder,
      ↪FunctionTransformer
      from sklearn.feature_extraction import FeatureHasher
      from sklearn.compose import ColumnTransformer
      from sklearn.pipeline import Pipeline
      from sklearn.metrics import accuracy_score, classification_report,
      ↪confusion_matrix, cohen_kappa_score
      from sklearn.base import BaseEstimator, TransformerMixin, clone
      from collections import Counter
      from xgboost import XGBClassifier
```

```
[3]: from google.cloud import bigquery

      project_id = "reliable-jet-452114-s2"

      client = bigquery.Client(project=project_id)
```

2 Datasets

The datasets used in these project were the ones we considered most relevant for the task at hand. They were uploaded to the BigQuery project `reliable-jet-452114 s-2` via Google Cloud Storage. Below we perform some queries to vizualize and ensure data quality in all of these the datasets.

2.1 Chartevents Reduced

Contains time-stamped clinical data and measurements recorded for patients during their hospital stay.

- `ROW_ID`: Unique identifier for the row.
- `SUBJECT_ID`: Foreign key to the `PATIENTS` table.
- `HADM_ID`: Foreign key to the `ADMISSIONS` table.
- `ITEMID`: Foreign key to the `D_ITEMS` table, indicating the type of measurement.
- `CHARTTIME`: Timestamp when the measurement was recorded.
- `STORETIME`: Timestamp when the measurement was stored.
- `VALUE`: The value of the measurement (can be numeric or text).
- `VALUENUM`: Numeric value of the measurement, if applicable.
- `VALUEUOM`: Unit of measurement.
- `WARNING`: Flags indicating potential issues with the data.

2.1.1 Reducing the original Chartevents dataset

We decided to create a more small and acessible table using the original `chartevents`. This new table, `chartevents_reduced`, that will be used from now on, only contains the measurements of the most common disease in the original table.

```
[ ]: start_time = time.time()

query = """
-- Step 1: Create a new table for disease-related data
CREATE OR REPLACE TABLE `reliable-jet-452114-s2.table.chartevents_reduced` AS

-- Step 2: Identify and select only disease-related measurements
WITH disease_measurements AS (
  SELECT *
  FROM `reliable-jet-452114-s2.table.chartevents`
  WHERE ITEMID IN (
    -- Cardiovascular
    220045, -- Heart Rate
    220050, -- Blood Pressure Systolic
    220051, -- Blood Pressure Diastolic

    -- Metabolic/Endocrine
    220179, -- Glucose
    50912,  -- Creatinine
    50809,  -- Glucose (serum)
```

```

-- Respiratory
220277, -- SpO2
224690, -- Respiratory Rate

-- Infection/Inflammation
50813, -- Lactate (sepsis marker)

-- Liver
50821 -- Bilirubin
)
),

-- Step 3: Find the top 3 most common disease measurements
top_disease_measurements AS (
    SELECT
        ITEMID,
        COUNT(*) AS measurement_count
    FROM disease_measurements
    GROUP BY ITEMID
    ORDER BY measurement_count DESC
    LIMIT 3
)

-- Step 4: Create final table with only top disease measurements
SELECT d.*
FROM disease_measurements d
JOIN top_disease_measurements t ON d.ITEMID = t.ITEMID;
"""

query_job = client.query(query)
print("Dataset reduced successfully")

end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")

```

Dataset reduced successfully
Query Execution Time: 0.29 seconds

2.1.2 Visualization

```

[ ]: start_time = time.time()

dataset_id = "table"
table_id = "chartevents_reduced"

```

```

table_ref = client.dataset(dataset_id).table(table_id)

chartevents = client.list_rows(table_ref).to_dataframe()

end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")

chartevents.head()

```

Query Execution Time: 5.15 seconds

```

[ ]:
  ROW_ID  SUBJECT_ID  HADM_ID  ICUSTAY_ID  ITEMID  \
0  19184679      66298   152072    200105  220045
1  19184805      66298   152072    200105  220045
2    334575       3952   112643    200563  220179
3   2259032      20173   154817    202537  220045
4   30271971     90629   100197    203563  220045

      CHARTTIME                STORETIME  CGID  VALUE  VALUENUM  \
0  2104-10-23 23:57:00+00:00             NaT  <NA>      0      0.0
1  2104-10-24 00:03:00+00:00             NaT  <NA>      0      0.0
2  2128-03-04 23:00:00+00:00  2128-03-04 23:15:00+00:00  17446      0      0.0
3  2108-09-27 05:05:00+00:00  2108-09-27 05:32:00+00:00  20622      0      0.0
4  2136-10-24 12:55:00+00:00  2136-10-24 13:03:00+00:00  21386      0      0.0

  VALUEUOM  WARNING  ERROR  RESULTSTATUS  STOPPED
0      bpm         0      0           None     None
1      bpm         0      0           None     None
2     mmHg         1      0           None     None
3      bpm         1      0           None     None
4      bpm         0      0           None     None

```

2.1.3 Data quality check

```

[ ]: start_time = time.time()

query = """
-- Data Quality Assessment for chartevents_reduced
WITH stats AS (
  SELECT
    COUNT(*) AS total_records,
    COUNT(DISTINCT SUBJECT_ID) AS unique_patients,
    COUNT(DISTINCT HADM_ID) AS unique_admissions,
    COUNT(DISTINCT ITEMID) AS unique_measurement_types
  FROM `reliable-jet-452114-s2.table.chartevents_reduced`
),

```

```

measurement_analysis AS (
  SELECT
    ITEMID,
    COUNT(*) AS record_count,
    ROUND(COUNT(*)*100/(SELECT total_records FROM stats), 2) AS_
    percentage_of_total,
    MIN(VALUENUM) AS min_value,
    MAX(VALUENUM) AS max_value,
    AVG(VALUENUM) AS avg_value,
    COUNT(CASE WHEN VALUENUM IS NULL THEN 1 END) AS null_value_counts,
    COUNT(CASE WHEN VALUE = '' THEN 1 END) AS empty_string_counts,
    MIN(CHARTTIME) AS earliest_measurement,
    MAX(CHARTTIME) AS latest_measurement
  FROM `reliable-jet-452114-s2.table.chartevents_reduced`
  GROUP BY ITEMID
),

temporal_analysis AS (
  SELECT
    EXTRACT(YEAR FROM CHARTTIME) AS year,
    EXTRACT(MONTH FROM CHARTTIME) AS month,
    COUNT(*) AS measurements_count
  FROM `reliable-jet-452114-s2.table.chartevents_reduced`
  GROUP BY year, month
  ORDER BY year, month
)

SELECT
  -- Basic Statistics
  s.total_records,
  s.unique_patients,
  s.unique_admissions,
  s.unique_measurement_types,

  -- Measurement-specific quality metrics
  ARRAY(
    SELECT AS STRUCT * FROM measurement_analysis
    ORDER BY record_count DESC
  ) AS measurement_quality,

  -- Temporal distribution
  (SELECT COUNT(*) FROM temporal_analysis) AS months_with_data,
  (SELECT MIN(year) FROM temporal_analysis) AS first_year,
  (SELECT MAX(year) FROM temporal_analysis) AS last_year,

  -- Data completeness

```

```

    (SELECT COUNT(*) FROM `reliable-jet-452114-s2.table.chartevents_reduced`
    ↪WHERE CHARTTIME IS NULL) AS null_timestamps,
    (SELECT COUNT(*) FROM `reliable-jet-452114-s2.table.chartevents_reduced`
    ↪WHERE SUBJECT_ID IS NULL) AS null_patient_ids,

    -- Clinical validity checks
    (SELECT COUNT(*) FROM `reliable-jet-452114-s2.table.chartevents_reduced`
    WHERE ITEMID = 220045 AND (VALUENUM < 20 OR VALUENUM > 250)) AS
    ↪abnormal_heart_rates,

    (SELECT COUNT(*) FROM `reliable-jet-452114-s2.table.chartevents_reduced`
    WHERE ITEMID = 220050 AND (VALUENUM < 50 OR VALUENUM > 300)) AS
    ↪abnormal_bp_readings
FROM stats s
"""

query_job = client.query(query)
data_quality = query_job.to_dataframe()

end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")

data_quality

```

Query Execution Time: 1.93 seconds

```

[ ]:  total_records  unique_patients  unique_admissions  \
0      6724529      17717      21927

      unique_measurement_types  \
0              3

                                measurement_quality  months_with_data  \
0  [{'ITEMID': 220045, 'record_count': 2762225, '...'      1289

      first_year  last_year  null_timestamps  null_patient_ids  \
0      2100      2209      0      0

      abnormal_heart_rates  abnormal_bp_readings
0              1177      0

```

2.2 Admissions

Contains information about patient admissions to the hospital, such as admission and discharge times, admission type, and insurance information.

- ROW_ID: Unique identifier for each row

- SUBJECT_ID: Foreign key to the PATIENTS table.
- HADM_ID: Unique identifier for the hospital admission
- ADMITTIME: Timestamp for hospital admission.
- DISCHTIME: Timestamp for hospital discharge.
- DEATHTIME: Timestamp for patient death, if applicable.
- ADMISSION_TYPE: Type of admission, e.g., EMERGENCY, ELECTIVE, URGENT.
- ADMISSION_LOCATION: Location of the patient prior to admission.
- DISCHARGE_LOCATION: Location to which the patient was discharged.
- INSURANCE: The patient's insurance provider.
- LANGUAGE: The patient's primary language.
- RELIGION: The patient's religious affiliation.
- MARITAL_STATUS: The patient's marital status.
- ETHNICITY: The patient's ethnicity.
- EDREGTIME: Emergency Department registration time
- EDOUTTIME: Emergency Department departure time
- DIAGNOSIS: The patient's primary diagnosis.
- HOSPITAL_EXPIRE_FLAG: Indicates if the patient died in the hospital.
- HAS_CHARTEVENTS_DATA: Flag indicating if there is chart event data.

2.2.1 Visualization

```
[ ]: dataset_id = "table"
table_id = "admissions"

table_ref = client.dataset(dataset_id).table(table_id)

admissions = client.list_rows(table_ref).to_dataframe()

admissions.head()
```

```
[ ]:  ROW_ID  SUBJECT_ID  HADM_ID  ADMITTIME  \
0    4060      3369    126808  2111-01-24  11:53:00+00:00
1    50952     74869    123152  2150-09-05  17:49:00+00:00
2    12812     10484    113233  2190-09-18  22:39:00+00:00
3    13573     11091    164694  2101-06-07  13:57:00+00:00
4    33654     27527    155091  2131-08-27  18:01:00+00:00

      DISCHTIME  DEATHTIME  ADMISSION_TYPE  \
0  2111-01-25  22:40:00+00:00  2111-01-25  22:40:00+00:00  EMERGENCY
1  2150-09-12  18:30:00+00:00  2150-09-12  18:30:00+00:00  EMERGENCY
2  2190-09-24  20:40:00+00:00  2190-09-24  20:40:00+00:00  EMERGENCY
3  2101-09-18  07:20:00+00:00  2101-09-18  07:20:00+00:00  EMERGENCY
4  2131-10-03  05:30:00+00:00  2131-10-03  05:30:00+00:00  EMERGENCY

      ADMISSION_LOCATION  DISCHARGE_LOCATION  INSURANCE  LANGUAGE  \
0  PHYS REFERRAL/NORMAL DELI  DEAD/EXPIRED  Private  None
1  TRANSFER FROM HOSP/EXTRAM  DEAD/EXPIRED  Medicare  ENGL
```

2	TRANSFER FROM HOSP/EXTRAM	DEAD/EXPIRED	Medicaid	None
3	CLINIC REFERRAL/PREMATURE	DEAD/EXPIRED	Private	None
4	CLINIC REFERRAL/PREMATURE	DEAD/EXPIRED	Private	PTUN

	RELIGION	MARITAL_STATUS	ETHNICITY	EDREGTIME	EDOUTTIME \
0	UNOBTAINABLE	SINGLE	WHITE	NaT	NaT
1	CATHOLIC	WIDOWED	ASIAN - JAPANESE	NaT	NaT
2	CATHOLIC	MARRIED	WHITE	NaT	NaT
3	CATHOLIC	SEPARATED	WHITE	NaT	NaT
4	NOT SPECIFIED	MARRIED	WHITE	NaT	NaT

	DIAGNOSIS	HOSPITAL_EXPIRE_FLAG \
0	None	1
1	? SEROTONIN SYNDROME	1
2	(AML) ACUTE MYELOGENOUS LEUKEMIA	1
3	(AML) ACUTE MYELOGENOUS LEUKEMIA	1
4	(AML) ACUTE MYELOGENOUS LEUKEMIA	1

	HAS_CHARTEVENTS_DATA
0	1
1	1
2	1
3	1
4	1

2.2.2 Data quality check


```
[ ]: query = """
-- Data Quality Assessment for admissions
SELECT
    COUNT(*) AS total_records,
    COUNT(DISTINCT HADM_ID) AS unique_admissions,
    COUNT(DISTINCT SUBJECT_ID) AS unique_patients,
    COUNT(CASE WHEN ADMITTIME IS NULL THEN 1 END) AS null_admit_times,
    COUNT(CASE WHEN DISCHTIME IS NULL THEN 1 END) AS null_discharge_times,
    COUNT(CASE WHEN ADMISSION_TYPE IS NULL THEN 1 END) AS null_admission_types,
    COUNT(CASE WHEN ADMISSION_LOCATION IS NULL THEN 1 END) AS_
↪null_admission_locations,
    COUNT(CASE WHEN INSURANCE IS NULL THEN 1 END) AS null_insurance_info,
    MIN(ADMITTIME) AS first_admission,
    MAX(ADMITTIME) AS last_admission
FROM `reliable-jet-452114-s2.table.admissions`;
"""

query_job = client.query(query)
data_quality = query_job.to_dataframe()
data_quality
```

```
[ ]:      total_records  unique_admissions  unique_patients  null_admit_times  \
0           58976           58976           46520           0

      null_discharge_times  null_admission_types  null_admission_locations  \
0                0                0                0

      null_insurance_info      first_admission      last_admission
0                0 2100-06-07 19:59:00+00:00 2210-08-17 17:13:00+00:00
```

2.3 Callout

Contains information about requests for services or consultations for patients.

- ROW_ID: Unique identifier for the row.
- SUBJECT_ID: Foreign key to the PATIENTS table.
- HADM_ID: Foreign key to the ADMISSIONS table.
- CALLOUT_ID: Unique identifier for the callout request.
- CALLOUTTIME: Timestamp for the callout request.
- SERVICE_ID: ID of the service requested.
- LOCATION: Location of the patient when the callout was placed.
- STATUS: Status of the callout request.
- OUTCOME: Outcome of the callout request.
- ACKNOWLEDGE_TIME: Timestamp when the callout was acknowledged.
- OUTCOMETIME: Timestamp when the callout outcome was recorded.
- FIRSTRESERVATIONTIME: of the first reservation.
- CURRENTRESERVATIONTIME: Timestamp of the current reservation.
- CREATETIME: Timestamp when the row was created.

- UPDATETIME: Timestamp when the row was updated.
- CALLOUT_WARDID: Ward ID of the callout.
- CALLOUT_SERVICEREQUEST: Service requested.
- CALLOUT_TELEPHONE: Telephone number for the callout.
- REQUEST_TELE: Telephone request.
- REQUEST_RESP: Respiratory reqTimestampuest.
- REQUEST_CDIF: C. difficile request.
- REQUEST_MRSA: MRSA request.
- REQUEST_VRE: VRE request.
- DISCHARGE_WARDID: Discharge ward ID.
- ACKNOWLEDGE_STATUS: Acknowledge status.

2.3.1 Visualization

```
[ ]: dataset_id = "table"
table_id = "callout"

table_ref = client.dataset(dataset_id).table(table_id)

callout = client.list_rows(table_ref).to_dataframe()

callout.head()
```

```
[ ]:  ROW_ID  SUBJECT_ID  HADM_ID  SUBMIT_WARDID  SUBMIT_CAREUNIT  CURR_WARDID  \
0    15115      31974   144780          <NA>          None          <NA>
1      161        309   162308           7          None           2
2      169        333   160548           7          None           2
3      197        383   173723           7          None           2
4      136        253   176189           7          None           2

  CURR_CAREUNIT  CALLOUT_WARDID  CALLOUT_SERVICE  REQUEST_TELE  ...  \
0          None              1          MED           1  ...
1          CCU              2          CCU           1  ...
2          CCU              2          CCU           1  ...
3          CCU              2          CCU           1  ...
4          CCU              2          CCU           1  ...

  CALLOUT_STATUS  CALLOUT_OUTCOME  DISCHARGE_WARDID  ACKNOWLEDGE_STATUS  \
0      Inactive      Discharged           0      Unacknowledged
1      Inactive      Discharged           2      Acknowledged
2      Inactive      Discharged           2      Acknowledged
3      Inactive      Discharged           2      Unacknowledged
4      Inactive      Discharged           2      Acknowledged

          CREATETIME          UPGRADETIME  \
0  2191-01-26 13:55:10+00:00  2191-01-26 13:55:10+00:00
1  2160-06-05 10:22:04+00:00  2160-06-05 10:22:04+00:00
```

```

2 2137-09-30 09:42:12+00:00 2137-09-30 09:42:12+00:00
3 2143-09-08 10:53:04+00:00 2143-09-08 10:53:04+00:00
4 2174-01-23 09:57:24+00:00 2174-01-23 10:44:12+00:00

```

	ACKNOWLEDGETIME	OUTCOMETIME	FIRSTRESERVATIONTIME	\
0	NaT	2191-01-26 14:10:04+00:00		NaT
1	2160-06-05 11:20:06+00:00	2160-06-05 19:25:01+00:00		NaT
2	2137-09-30 09:45:08+00:00	2137-10-01 14:40:02+00:00		NaT
3	NaT	2143-09-08 11:55:02+00:00		NaT
4	2174-01-23 11:10:50+00:00	2174-01-23 13:40:02+00:00		NaT

	CURRENTRESERVATIONTIME
0	NaT
1	NaT
2	NaT
3	NaT
4	NaT

[5 rows x 24 columns]

2.3.2 Data quality check

```

[ ]: query = """
-- Data Quality Assessment for callout
WITH basic_stats AS (
    SELECT
        COUNT(*) AS total_records,
        COUNT(DISTINCT SUBJECT_ID) AS unique_patients,
        COUNT(DISTINCT HADM_ID) AS unique_admissions,
        COUNT(DISTINCT ROW_ID) AS unique_row_ids,
        COUNT(DISTINCT SUBMIT_WARDID) AS unique_submit_ward_ids,
        COUNT(DISTINCT SUBMIT_CAREUNIT) AS unique_submit_care_units,
        COUNT(DISTINCT CURR_WARDID) AS unique_current_ward_ids,
        COUNT(DISTINCT CURR_CAREUNIT) AS unique_current_care_units,
        COUNT(DISTINCT CALLOUT_WARDID) AS unique_callout_ward_ids,
        COUNT(DISTINCT CALLOUT_SERVICE) AS unique_callout_service,
        COUNT(DISTINCT REQUEST_TELE) AS unique_request_telephones,
        COUNT(DISTINCT REQUEST_RESP) AS unique_request_resp,
        COUNT(DISTINCT REQUEST_CDIFF) AS unique_request_cdifff,
        COUNT(DISTINCT REQUEST_MRSA) AS unique_request_mrsa,
        COUNT(DISTINCT REQUEST_VRE) AS unique_request_vre,
        COUNT(DISTINCT CALLOUT_STATUS) AS unique_callout_statuses,
        COUNT(DISTINCT CALLOUT_OUTCOME) AS unique_callout_outcomes,
        COUNT(DISTINCT DISCHARGE_WARDID) AS unique_discharge_ward_ids,
        COUNT(DISTINCT ACKNOWLEDGE_STATUS) AS unique_acknowledge_statuses
    FROM `reliable-jet-452114-s2.table.callout`
),

```

```

completeness AS (
    SELECT
        COUNT(CASE WHEN SUBJECT_ID IS NULL THEN 1 END) AS null_subject_ids,
        COUNT(CASE WHEN HADM_ID IS NULL THEN 1 END) AS null_hadm_ids,
        COUNT(CASE WHEN SUBMIT_WARDID IS NULL THEN 1 END) AS
↪null_submit_ward_ids,
        COUNT(CASE WHEN SUBMIT_CAREUNIT IS NULL THEN 1 END) AS
↪null_submit_care_units,
        COUNT(CASE WHEN SUBMIT_CAREUNIT = '' THEN 1 END) AS
↪empty_submit_care_units,
        COUNT(CASE WHEN CURR_WARDID IS NULL THEN 1 END) AS null_curr_ward_ids,
        COUNT(CASE WHEN CURR_CAREUNIT IS NULL THEN 1 END) AS
↪null_curr_care_units,
        COUNT(CASE WHEN CURR_CAREUNIT = '' THEN 1 END) AS empty_curr_care_units,
        COUNT(CASE WHEN CALLOUT_WARDID IS NULL THEN 1 END) AS
↪null_callout_ward_ids,
        COUNT(CASE WHEN CALLOUT_SERVICE IS NULL THEN 1 END) AS
↪null_callout_service,
        COUNT(CASE WHEN CALLOUT_SERVICE = '' THEN 1 END) AS
↪empty_callout_service,
        COUNT(CASE WHEN REQUEST_TELE IS NULL THEN 1 END) AS
↪null_request_telephones,
        COUNT(CASE WHEN REQUEST_RESP IS NULL THEN 1 END) AS null_request_resp,
        COUNT(CASE WHEN REQUEST_CDIF IS NULL THEN 1 END) AS null_request_cdif,
        COUNT(CASE WHEN REQUEST_MRSA IS NULL THEN 1 END) AS null_request_mrsa,
        COUNT(CASE WHEN REQUEST_VRE IS NULL THEN 1 END) AS null_request_vre,
        COUNT(CASE WHEN CALLOUT_STATUS IS NULL THEN 1 END) AS
↪null_callout_statuses,
        COUNT(CASE WHEN CALLOUT_STATUS = '' THEN 1 END) AS
↪empty_callout_statuses,
        COUNT(CASE WHEN CALLOUT_OUTCOME IS NULL THEN 1 END) AS
↪null_callout_outcomes,
        COUNT(CASE WHEN CALLOUT_OUTCOME = '' THEN 1 END) AS
↪empty_callout_outcomes,
        COUNT(CASE WHEN DISCHARGE_WARDID IS NULL THEN 1 END) AS
↪null_discharge_ward_ids,
        COUNT(CASE WHEN ACKNOWLEDGE_STATUS IS NULL THEN 1 END) AS
↪null_acknowledge_statuses,
        COUNT(CASE WHEN ACKNOWLEDGE_STATUS = '' THEN 1 END) AS
↪empty_acknowledge_statuses,
        COUNT(CASE WHEN CREATETIME IS NULL THEN 1 END) AS null_createtimes,
        COUNT(CASE WHEN UPDATETIME IS NULL THEN 1 END) AS null_updatetimes,
        COUNT(CASE WHEN ACKNOWLEDGETIME IS NULL THEN 1 END) AS
↪null_acknowledgetimes,
        COUNT(CASE WHEN OUTCOMETIME IS NULL THEN 1 END) AS null_outcometimes,

```

```

        COUNT(CASE WHEN FIRSTRESERVATIONTIME IS NULL THEN 1 END) AS_
↪null_firstreservationtimes,
        COUNT(CASE WHEN CURRENTRESERVATIONTIME IS NULL THEN 1 END) AS_
↪null_currentreservationtimes
    FROM `reliable-jet-452114-s2.table.callout`
),
temporal_analysis AS (
    SELECT
        MIN(CREATETIME) AS first_creation_time,
        MAX(CREATETIME) AS last_creation_time,
        MIN(UPDATETIME) AS first_update_time,
        MAX(UPDATETIME) AS last_update_time,
        MIN(ACKNOWLEDGETIME) AS first_acknowledge_time,
        MAX(ACKNOWLEDGETIME) AS last_acknowledge_time,
        MIN(OUTCOMETIME) AS first_outcome_time,
        MAX(OUTCOMETIME) AS last_outcome_time,
        MIN(FIRSTRESERVATIONTIME) AS first_reservation_time,
        MAX(FIRSTRESERVATIONTIME) AS last_reservation_time,
        MIN(CURRENTRESERVATIONTIME) AS first_current_reservation_time,
        MAX(CURRENTRESERVATIONTIME) AS last_current_reservation_time
    FROM `reliable-jet-452114-s2.table.callout`
)
SELECT
    bs.*,
    c.*,
    t.*
FROM basic_stats bs
CROSS JOIN completeness c
CROSS JOIN temporal_analysis t;
"""

query_job = client.query(query)
data_quality = query_job.to_dataframe()
data_quality

```

```

[ ]:  total_records  unique_patients  unique_admissions  unique_row_ids  \
0      34499          22871          28732          34499

      unique_submit_ward_ids  unique_submit_care_units  unique_current_ward_ids  \
0              9              5              35

      unique_current_care_units  unique_callout_ward_ids  unique_callout_service  \
0              5              37              21

...      first_update_time      last_update_time  \
0 ... 2100-06-08 12:58:29+00:00 2210-08-20 16:05:16+00:00

```

```

      first_acknowledge_time      last_acknowledge_time  \
0 2100-06-08 12:58:32+00:00 2210-08-20 16:05:27+00:00

      first_outcome_time      last_outcome_time  \
0 2100-06-08 15:10:26+00:00 2210-08-20 18:55:15+00:00

      first_reservation_time      last_reservation_time  \
0 2100-06-08 11:55:26+00:00 2210-08-20 16:25:16+00:00

      first_current_reservation_time  last_current_reservation_time
0          2100-08-09 14:42:25+00:00          2209-08-05 14:01:23+00:00

[1 rows x 60 columns]

```

2.4 ICD Diagnoses Description

Contains descriptions for ICD-9 diagnosis codes.

- ROW_ID: Unique identifier for the row.
- ICD9_CODE: The ICD-9 diagnosis code.
- SHORT_TITLE: Short description of the diagnosis.
- LONG_TITLE: Long description of the diagnosis.

2.4.1 Visualization

```

[ ]: dataset_id = "table"
      table_id = "d_icd_diagnoses"

      table_ref = client.dataset(dataset_id).table(table_id)

      d_icd_diagnoses = client.list_rows(table_ref).to_dataframe()

      d_icd_diagnoses.head()

```

```

[ ]:  ROW_ID  ICD9_CODE      SHORT_TITLE  \
0      5120      4957  "ventilation" pneumonit
1      11159      94416  1 deg burn back of hand
2      11157      94414  1 deg burn fingr w thumb
3       3658      36911  1 eye-sev/oth-blind NOS
4      12505      94811  10-19% bdy brn/10-19% 3d

                                LONG_TITLE
0                                "Ventilation" pneumonitis
1                Erythema [first degree] of back of hand
2  Erythema [first degree] of two or more digits ...
3  Better eye: severe vision impairment; lesser e...
4  Burn [any degree] involving 10-19 percent of b...

```

2.4.2 Data quality check

```
[ ]: query = """
-- Data Quality Assessment for d_icd_diagnoses
SELECT
    COUNT(*) AS total_records,
    COUNT(DISTINCT ROW_ID) AS unique_row_ids,
    COUNT(DISTINCT ICD9_CODE) AS unique_icd9_codes,
    COUNT(DISTINCT SHORT_TITLE) AS unique_short_titles,
    COUNT(DISTINCT LONG_TITLE) AS unique_long_titles,

    -- Completeness checks
    COUNT(CASE WHEN ROW_ID IS NULL THEN 1 END) AS null_row_ids,
    COUNT(CASE WHEN ICD9_CODE IS NULL THEN 1 END) AS null_icd9_codes,
    COUNT(CASE WHEN SHORT_TITLE IS NULL THEN 1 END) AS null_short_titles,
    COUNT(CASE WHEN LONG_TITLE IS NULL THEN 1 END) AS null_long_titles,

    COUNT(CASE WHEN ICD9_CODE = '' THEN 1 END) AS empty_icd9_codes,
    COUNT(CASE WHEN SHORT_TITLE = '' THEN 1 END) AS empty_short_titles,
    COUNT(CASE WHEN LONG_TITLE = '' THEN 1 END) AS empty_long_titles,

    -- Basic Analysis of ICD9 Code Length
    AVG(LENGTH(ICD9_CODE)) AS avg_icd9_code_length,
    MIN(LENGTH(ICD9_CODE)) AS min_icd9_code_length,
    MAX(LENGTH(ICD9_CODE)) AS max_icd9_code_length,
    COUNT(CASE WHEN LENGTH(ICD9_CODE) NOT BETWEEN 3 AND 5 THEN 1 END) AS
    potential_invalid_icd9_length_count,

    -- Potential data inconsistencies
    COUNT(CASE WHEN SHORT_TITLE LIKE '%NOS%' AND LONG_TITLE NOT LIKE '%not_
    otherwise specified%' THEN 1 END) AS short_title_nos_long_title_mismatch,
    COUNT(CASE WHEN SHORT_TITLE LIKE '%NEC%' AND LONG_TITLE NOT LIKE '%not_
    elsewhere classified%' THEN 1 END) AS short_title_nec_long_title_mismatch

FROM
    `reliable-jet-452114-s2.table.d_icd_diagnoses`;
"""
query_job = client.query(query)
data_quality = query_job.to_dataframe()
data_quality
```

```
[ ]:   total_records  unique_row_ids  unique_icd9_codes  unique_short_titles  \
0          14567          14567          14567          14328

   unique_long_titles  null_row_ids  null_icd9_codes  null_short_titles  \
0          14562           0           0           0
```

0	0	0	0	0
0	4.686483	3	5	
0	0	2195		
0	2138			

2.5 D Items

Contains metadata about different medical items (measurements, procedures, medications) recorded in the database.

- ROW_ID: Unique identifier for the row.
- ITEMID: Unique identifier for the item.
- LABEL: Label or name of the item.
- ABBREVIATION: Abbreviation for the item.
- DBSOURCE: Source database for the item.
- LINKSTO: Table linked to.
- CATEGORY: Category of the item.
- UNITNAME: Unit of measurement for the item.
- PARAM_TYPE: Parameter type.
- CONCEPTID: Identifier for the concept.

2.5.1 Visualization

```
[ ]: dataset_id = "table"
      table_id = "d_items"

      table_ref = client.dataset(dataset_id).table(table_id)

      d_items = client.list_rows(table_ref).to_dataframe()

      d_items.head()
```

```
[ ]:  ROW_ID  ITEMID  LABEL ABBREVIATION \
0     457    497  Patient controlled analgesia (PCA) [Inject]  None
1     458    498                PCA Lockout (Min)             None
2     459    499                PCA Medication                None
3     460    500                PCA Total Dose                 None
4     461    501                PCV Exh Vt (Obser)             None

      DBSOURCE  LINKSTO  CATEGORY  UNITNAME  PARAM_TYPE  CONCEPTID
0  carevue    chartevents      None      None      None      None
```


1	carevue	chartevents	None	None	None	None
2	carevue	chartevents	None	None	None	None
3	carevue	chartevents	None	None	None	None
4	carevue	chartevents	None	None	None	None

2.5.2 Data quality check

```
[ ]: query = """
-- Data Quality Assessment for d_items
SELECT
    COUNT(*) AS total_records,
    COUNT(DISTINCT ROW_ID) AS unique_row_ids,
    COUNT(DISTINCT ITEMID) AS unique_itemids,
    COUNT(DISTINCT LABEL) AS unique_labels,
    COUNT(DISTINCT ABBREVIATION) AS unique_abbreviations,
    COUNT(DISTINCT DBSOURCE) AS unique_dbsources,
    COUNT(DISTINCT LINKSTO) AS unique_linkstos,
    COUNT(DISTINCT CATEGORY) AS unique_categories,
    COUNT(DISTINCT UNITNAME) AS unique_unitnames,
    COUNT(DISTINCT PARAM_TYPE) AS unique_param_types,
    COUNT(DISTINCT CONCEPTID) AS unique_conceptids,

    -- Completeness checks (NULL values)
    COUNT(CASE WHEN ROW_ID IS NULL THEN 1 END) AS null_row_ids,
    COUNT(CASE WHEN ITEMID IS NULL THEN 1 END) AS null_itemids,
    COUNT(CASE WHEN LABEL IS NULL THEN 1 END) AS null_labels,
    COUNT(CASE WHEN ABBREVIATION IS NULL THEN 1 END) AS null_abbreviations,
    COUNT(CASE WHEN DBSOURCE IS NULL THEN 1 END) AS null_dbsources,
    COUNT(CASE WHEN LINKSTO IS NULL THEN 1 END) AS null_linkstos,
    COUNT(CASE WHEN CATEGORY IS NULL THEN 1 END) AS null_categories,
    COUNT(CASE WHEN UNITNAME IS NULL THEN 1 END) AS null_unitnames,
    COUNT(CASE WHEN PARAM_TYPE IS NULL THEN 1 END) AS null_param_types,
    COUNT(CASE WHEN CONCEPTID IS NULL THEN 1 END) AS null_conceptids,

    -- Completeness checks (Empty strings)
    COUNT(CASE WHEN LABEL = '' THEN 1 END) AS empty_labels,
    COUNT(CASE WHEN ABBREVIATION = '' THEN 1 END) AS empty_abbreviations,
    COUNT(CASE WHEN DBSOURCE = '' THEN 1 END) AS empty_dbsources,
    COUNT(CASE WHEN LINKSTO = '' THEN 1 END) AS empty_linkstos,
    COUNT(CASE WHEN CATEGORY = '' THEN 1 END) AS empty_categories,
    COUNT(CASE WHEN UNITNAME = '' THEN 1 END) AS empty_unitnames,
    COUNT(CASE WHEN PARAM_TYPE = '' THEN 1 END) AS empty_param_types,
    COUNT(CASE WHEN CONCEPTID = '' THEN 1 END) AS empty_conceptids

FROM
    `reliable-jet-452114-s2.table.d_items`;
"""
```

```
query_job = client.query(query)
data_quality = query_job.to_dataframe()
data_quality
```

```
[ ]:   total_records  unique_row_ids  unique_itemids  unique_labels  \
0           12487           12487           12487           11847

      unique_abbreviations  unique_dbsources  unique_linkstos  unique_categories  \
0                2907                3                7                94

      unique_unitnames  unique_param_types  ...  null_param_types  \
0                53                7  ...                9495

      null_conceptids  empty_labels  empty_abbreviations  empty_dbsources  \
0                12487                0                0                0

      empty_linkstos  empty_categories  empty_unitnames  empty_param_types  \
0                0                0                0                0

      empty_conceptids
0                0

[1 rows x 29 columns]
```

2.6 ICD Diagnoses

Contains ICD-9 diagnosis codes assigned to patients during their hospital admissions.

- ROW_ID: Unique identifier for the row.
- SUBJECT_ID: Foreign key to the PATIENTS table.
- HADM_ID: Foreign key to the ADMISSIONS table.
- SEQ_NUM: Sequence number for the diagnosis within the admission.
- ICD9_CODE: The ICD-9 diagnosis code.

2.6.1 Visualization

```
[ ]: dataset_id = "table"
      table_id = "diagnoses_icd"

      table_ref = client.dataset(dataset_id).table(table_id)

      diagnoses_icd = client.list_rows(table_ref).to_dataframe()

      diagnoses_icd.head()
```

```
[ ]:   ROW_ID  SUBJECT_ID  HADM_ID  SEQ_NUM  ICD9_CODE
0     1297         109    172335         1     40301
1     1298         109    172335         2         486
```

2	1299	109	172335	3	58281
3	1300	109	172335	4	5855
4	1301	109	172335	5	4254

2.6.2 Data quality check

```
[ ]: query = """
-- Data Quality Assessment for diagnoses_icd
SELECT
    COUNT(*) AS total_records,
    COUNT(DISTINCT ROW_ID) AS unique_row_ids,
    COUNT(DISTINCT SUBJECT_ID) AS unique_patients,
    COUNT(DISTINCT HADM_ID) AS unique_admissions,
    COUNT(DISTINCT SEQ_NUM) AS unique_sequence_numbers,
    COUNT(DISTINCT ICD9_CODE) AS unique_icd9_codes,

    -- Completeness checks (NULL values)
    COUNT(CASE WHEN ROW_ID IS NULL THEN 1 END) AS null_row_ids,
    COUNT(CASE WHEN SUBJECT_ID IS NULL THEN 1 END) AS null_subject_ids,
    COUNT(CASE WHEN HADM_ID IS NULL THEN 1 END) AS null_hadm_ids,
    COUNT(CASE WHEN SEQ_NUM IS NULL THEN 1 END) AS null_sequence_numbers,
    COUNT(CASE WHEN ICD9_CODE IS NULL THEN 1 END) AS null_icd9_codes,

    -- Completeness checks (Empty strings)
    COUNT(CASE WHEN ICD9_CODE = '' THEN 1 END) AS empty_icd9_codes,

    -- Analysis of ICD9 Code Length
    AVG(LENGTH(ICD9_CODE)) AS avg_icd9_code_length,
    MIN(LENGTH(ICD9_CODE)) AS min_icd9_code_length,
    MAX(LENGTH(ICD9_CODE)) AS max_icd9_code_length,
    COUNT(CASE WHEN LENGTH(ICD9_CODE) NOT BETWEEN 3 AND 5 THEN 1 END) AS_
    ↪potential_invalid_icd9_length_count

FROM
    `reliable-jet-452114-s2.table.diagnoses_icd`;
"""
query_job = client.query(query)
data_quality = query_job.to_dataframe()
data_quality
```

```
[ ]: total_records  unique_row_ids  unique_patients  unique_admissions  \
0          651047          651047          46520          58976

    unique_sequence_numbers  unique_icd9_codes  null_row_ids  null_subject_ids  \
0                   39                6984                0                0

    null_hadm_ids  null_sequence_numbers  null_icd9_codes  empty_icd9_codes  \
```

0	0	47	47	0
avg_icd9_code_length	min_icd9_code_length	max_icd9_code_length	\	
0	4.448883	3	5	
potential_invalid_icd9_length_count				
0	0			

2.7 Icustays

Contains information about patient stays in the intensive care unit (ICU).

- ROW_ID: Unique identifier for the row.
- SUBJECT_ID: Foreign key to the PATIENTS table.
- HADM_ID: Foreign key to the ADMISSIONS table.
- ICUSTAY_ID: Unique identifier for the ICU stay.
- DBSOURCE: Source database.
- FIRST_CAREUNIT: First care unit the patient was in.
- LAST_CAREUNIT: Last care unit the patient was in.
- FIRST_WARDID: First ward ID.
- LAST_WARDID: Last ward ID.
- INTIME: Timestamp for ICU admission.
- OUTTIME: Timestamp for ICU discharge.
- LOS: Length of ICU stay.

2.7.1 Visualization

```
[ ]: dataset_id = "table"
table_id = "icustays"

table_ref = client.dataset(dataset_id).table(table_id)

icustays = client.list_rows(table_ref).to_dataframe()

icustays.head()
```

```
[ ]:  ROW_ID  SUBJECT_ID  HADM_ID  ICUSTAY_ID  DBSOURCE  FIRST_CAREUNIT  \
0      372          275   129886    219649    carevue          CCU
1      389          291   113649    256641    carevue          CCU
2      390          291   125726    275109  metavision          CCU
3      394          294   152578    222074    carevue          CCU
4      401          301   160332    288401    carevue          CCU

      LAST_CAREUNIT  FIRST_WARDID  LAST_WARDID  INTIME  \
0          CCU          7          7  2170-10-07 11:28:53+00:00
1          CCU          7          7  2102-04-08 23:05:28+00:00
2          CCU          7          7  2106-04-17 12:26:17+00:00
3          CCU          7          7  2118-01-17 21:45:05+00:00
```

4 CCU 7 7 2189-11-11 12:12:33+00:00

		OUTTIME	LOS
0	2170-10-14	14:38:07+00:00	7.1314
1	2102-04-09	11:20:11+00:00	0.5102
2	2106-04-18	22:05:39+00:00	1.4023
3	2118-01-20	11:12:45+00:00	2.5609
4	2189-11-13	22:11:28+00:00	2.4159

2.7.2 Data quality check

```
[ ]: query = ""
-- Data Quality Assessment for icustays
SELECT
    COUNT(*) AS total_records,
    COUNT(DISTINCT ROW_ID) AS unique_row_ids,
    COUNT(DISTINCT SUBJECT_ID) AS unique_patients,
    COUNT(DISTINCT HADM_ID) AS unique_admissions,
    COUNT(DISTINCT ICUSTAY_ID) AS unique_icustay_ids,
    COUNT(DISTINCT DBSOURCE) AS unique_dbsources,
    COUNT(DISTINCT FIRST_CAREUNIT) AS unique_first_careunits,
    COUNT(DISTINCT LAST_CAREUNIT) AS unique_last_careunits,
    COUNT(DISTINCT FIRST_WARDID) AS unique_first_wardids,
    COUNT(DISTINCT LAST_WARDID) AS unique_last_wardids,

    -- Completeness checks (NULL values)
    COUNT(CASE WHEN ROW_ID IS NULL THEN 1 END) AS null_row_ids,
    COUNT(CASE WHEN SUBJECT_ID IS NULL THEN 1 END) AS null_subject_ids,
    COUNT(CASE WHEN HADM_ID IS NULL THEN 1 END) AS null_hadm_ids,
    COUNT(CASE WHEN ICUSTAY_ID IS NULL THEN 1 END) AS null_icustay_ids,
    COUNT(CASE WHEN DBSOURCE IS NULL THEN 1 END) AS null_dbsources,
    COUNT(CASE WHEN FIRST_CAREUNIT IS NULL THEN 1 END) AS null_first_careunits,
    COUNT(CASE WHEN LAST_CAREUNIT IS NULL THEN 1 END) AS null_last_careunits,
    COUNT(CASE WHEN FIRST_WARDID IS NULL THEN 1 END) AS null_first_wardids,
    COUNT(CASE WHEN LAST_WARDID IS NULL THEN 1 END) AS null_last_wardids,
    COUNT(CASE WHEN INTIME IS NULL THEN 1 END) AS null_intimes,
    COUNT(CASE WHEN OUTTIME IS NULL THEN 1 END) AS null_outtimes,
    COUNT(CASE WHEN LOS IS NULL THEN 1 END) AS null_los,

    -- Completeness checks (Empty strings)
    COUNT(CASE WHEN FIRST_CAREUNIT = '' THEN 1 END) AS empty_first_careunits,
    COUNT(CASE WHEN LAST_CAREUNIT = '' THEN 1 END) AS empty_last_careunits,
    COUNT(CASE WHEN DBSOURCE = '' THEN 1 END) AS empty_dbsources,

    -- Basic Analysis of LOS (Length of ICU Stay)
    AVG(LOS) AS avg_los,
    MIN(LOS) AS min_los,
```

```

        MAX(LOS) AS max_los,
        COUNT(CASE WHEN LOS < 0 THEN 1 END) AS negative_los_count -- Potential data_
↪issue

FROM
    `reliable-jet-452114-s2.table.icustays`;
"""
query_job = client.query(query)
data_quality = query_job.to_dataframe()
data_quality

```

```

[ ]:   total_records  unique_row_ids  unique_patients  unique_admissions  \
0           61532           61532           46476           57786

      unique_icustay_ids  unique_dbsources  unique_first_careunits  \
0           61532                3                6

      unique_last_careunits  unique_first_wardids  unique_last_wardids  ...  \
0                6                16                17  ...

      null_intimes  null_outtimes  null_los  empty_first_careunits  \
0                0                10        10                0

      empty_last_careunits  empty_dbsources  avg_los  min_los  max_los  \
0                0                0  4.917972  0.0001  173.0725

      negative_los_count
0                0

[1 rows x 29 columns]

```

2.8 Patients

Contains demographic information about the patients in the database.

- ROW_ID: Unique identifier for the row.
- SUBJECT_ID: Unique identifier for the patient.
- GENDER: Patient's gender.
- DOB: Patient's date of birth.
- DOD: Patient's date of death, if applicable.
- DOD_HOSP: Date of death in hospital.
- DOD_SSN: Date of death according to Social Security records.
- EXPIRE_FLAG: Flag indicating if the patient expired.

2.8.1 Visualization

```
[ ]: dataset_id = "table"
      table_id = "patients"

      table_ref = client.dataset(dataset_id).table(table_id)

      patients = client.list_rows(table_ref).to_dataframe()

      patients.head()
```

```
[ ]:  ROW_ID  SUBJECT_ID  GENDER  DOB  DOD  DOD_HOSP  DOD_SSN  \
0      234      249      F  2075-03-13  00:00:00+00:00  NaT      NaT      NaT
1      238      253      F  2089-11-26  00:00:00+00:00  NaT      NaT      NaT
2      242      258      F  2124-09-19  00:00:00+00:00  NaT      NaT      NaT
3      243      260      F  2105-03-23  00:00:00+00:00  NaT      NaT      NaT
4      247      264      F  2162-11-30  00:00:00+00:00  NaT      NaT      NaT

      EXPIRE_FLAG
0              0
1              0
2              0
3              0
4              0
```

2.8.2 Data quality check

```
[ ]: query = """
-- Data Quality Assessment for patients
SELECT
    COUNT(*) AS total_records,
    COUNT(DISTINCT ROW_ID) AS unique_row_ids,
    COUNT(DISTINCT SUBJECT_ID) AS unique_patients,
    COUNT(DISTINCT GENDER) AS unique_genders,
    COUNT(DISTINCT DOB) AS unique_dates_of_birth,
    COUNT(DISTINCT DOD) AS unique_dates_of_death,
    COUNT(DISTINCT DOD_HOSP) AS unique_dates_of_death_hospital,
    COUNT(DISTINCT DOD_SSN) AS unique_dates_of_death_ssn,
    COUNT(DISTINCT EXPIRE_FLAG) AS unique_expire_flags,

    -- Completeness checks (NULL values)
    COUNT(CASE WHEN ROW_ID IS NULL THEN 1 END) AS null_row_ids,
    COUNT(CASE WHEN SUBJECT_ID IS NULL THEN 1 END) AS null_subject_ids,
    COUNT(CASE WHEN GENDER IS NULL THEN 1 END) AS null_genders,
    COUNT(CASE WHEN DOB IS NULL THEN 1 END) AS null_dates_of_birth,
    COUNT(CASE WHEN DOD IS NULL THEN 1 END) AS null_dates_of_death,
```

```

COUNT(CASE WHEN DOD_HOSP IS NULL THEN 1 END) AS
↳null_dates_of_death_hospital,
COUNT(CASE WHEN DOD_SSN IS NULL THEN 1 END) AS null_dates_of_death_ssn,
COUNT(CASE WHEN EXPIRE_FLAG IS NULL THEN 1 END) AS null_expire_flags,

-- Completeness checks (Empty strings)
COUNT(CASE WHEN GENDER = '' THEN 1 END) AS empty_genders,

-- Basic Analysis of Dates
MIN(DOB) AS first_dob,
MAX(DOB) AS last_dob,
MIN(DOD) AS first_dod,
MAX(DOD) AS last_dod,

-- Potential Data Inconsistencies
COUNT(CASE WHEN DOD IS NOT NULL AND EXPIRE_FLAG = 0 THEN 1 END) AS
↳died_but_not_expired_flag,
COUNT(CASE WHEN DOD IS NULL AND EXPIRE_FLAG = 1 THEN 1 END) AS
↳expired_flag_but_no_dod

FROM
    `reliable-jet-452114-s2.table.patients`;
"""
query_job = client.query(query)
data_quality = query_job.to_dataframe()
data_quality

```

```

[ ]:  total_records  unique_row_ids  unique_patients  unique_genders  \
0          46520          46520          46520          2

      unique_dates_of_birth  unique_dates_of_death  \
0          32540          12911

      unique_dates_of_death_hospital  unique_dates_of_death_ssn  \
0          8747          11301

      unique_expire_flags  null_row_ids  ...  null_dates_of_death_hospital  \
0          2          0  ...          36546

      null_dates_of_death_ssn  null_expire_flags  empty_genders  \
0          33142          0          0

              first_dob              last_dob  \
0  1800-07-02 00:00:00+00:00  2201-07-24 00:00:00+00:00

              first_dod              last_dod  \
0  2100-06-19 00:00:00+00:00  2211-06-10 00:00:00+00:00

```


died_but_not_expired_flag	expired_flag_but_no_dod
0	0

[1 rows x 24 columns]

3 Junction of tables

To make the pre-processing and Length of ICU Stay prediction more efficient we decided to create a table, using a query, that contain the most relevant columns of each of the datasets.

The most relevant columns choosen were: - Patient Info: SUBJECT_ID, GENDER, DOB, EXPIRE_FLAG, AGE_AT_ADMISSION - Admission Info: HADM_ID, ADMITTIME, DISCHTIME, ADMISSION_TYPE, ADMISSION_LOCATION, INSURANCE, ETHNICITY, ADMISSION_DIAGNOSIS_TEXT, ADMISSION_HOUR - ICU Stay Info: ICUSTAY_ID, ICU_INTIME, ICU_OUTTIME, ICU_LOS (Target), FIRST_CAREUNIT - Diagnosis Info: PRIMARY_ICD9_CODE, PRIMARY_ICD9_TITLE - Callout Info: NUM_CALLOUTS

```
[ ]: query = """
-- Creating a junction table with aggregated diagnoses, callouts and engineered_
↳features

CREATE OR REPLACE TABLE `reliable-jet-452114-s2.table.junction_table` AS

-- Define CTEs within the AS clause
WITH
  -- CTE to select only the primary diagnosis for each hospital admission
  PrimaryDiagnosis AS (
    SELECT
      di.HADM_ID,
      di.ICD9_CODE,
      dd.SHORT_TITLE AS ICD9_SHORT_TITLE
    FROM
      `reliable-jet-452114-s2.table.diagnoses_icd` AS di
    LEFT JOIN
      `reliable-jet-452114-s2.table.d_icd_diagnoses` AS dd ON di.ICD9_CODE = dd.
↳ICD9_CODE
    WHERE
      di.SEQ_NUM = 1 -- Filter for primary diagnosis only
  ),

  -- CTE to count the number of callouts for each hospital admission
  CalloutCount AS (
    SELECT
      HADM_ID,
      COUNT(*) AS NUM_CALLOUTS -- Count callouts per admission
    FROM
      `reliable-jet-452114-s2.table.callout`
```

```

        GROUP BY
            HADM_ID
    )

-- Main SELECT statement that uses the CTEs
SELECT
    -- Patient demographics & calculated age
    p.SUBJECT_ID,
    p.GENDER,
    p.DOB,
    p.EXPIRE_FLAG,
    DATE_DIFF(DATE(a.ADMITTIME), DATE(p.DOB), YEAR) AS AGE_AT_ADMISSION, --□
    ↪Calculate age

    -- Admission details & extracted features
    a.HADM_ID,
    a.ADMITTIME,
    a.DISCHTIME,
    a.ADMISSION_TYPE,
    a.ADMISSION_LOCATION,
    a.INSURANCE,
    a.ETHNICITY,
    a.DIAGNOSIS AS ADMISSION_DIAGNOSIS_TEXT,
    EXTRACT(HOUR FROM a.ADMITTIME) AS ADMISSION_HOUR,

    -- ICU stay details (TARGET VARIABLE HERE)
    icu.ICUSTAY_ID,
    icu.INTIME AS ICU_INTIME,
    icu.OUTTIME AS ICU_OUTTIME,
    icu.LOS AS ICU_LOS,                -- TARGET VARIABLE
    icu.FIRST_CAREUNIT,

    -- Primary Diagnosis details (from CTE)
    pd.ICD9_CODE AS PRIMARY_ICD9_CODE,
    pd.ICD9_SHORT_TITLE AS PRIMARY_ICD9_TITLE,

    -- Aggregated Callout details (from CTE)
    cc.NUM_CALLOUTS

FROM
    -- Start with patients table
    `reliable-jet-452114-s2.table.patients` AS p
LEFT JOIN
    -- Join with admissions using SUBJECT_ID
    `reliable-jet-452114-s2.table.admissions` AS a ON p.SUBJECT_ID = a.
    ↪SUBJECT_ID
LEFT JOIN

```

```

-- Join with ICU stays using HADM_ID
`reliable-jet-452114-s2.table.icustays` AS icu ON a.HADM_ID = icu.HADM_ID
LEFT JOIN
-- Join with pre-filtered primary diagnosis using HADM_ID
PrimaryDiagnosis AS pd ON a.HADM_ID = pd.HADM_ID
LEFT JOIN
-- Join with pre-aggregated callout counts using HADM_ID
CalloutCount AS cc ON a.HADM_ID = cc.HADM_ID

"""

query_job = client.query(query)
query_job.result()

```

```
[ ]: <google.cloud.bigquery.table._EmptyRowIterator at 0x7a18ad91ac50>
```

3.0.1 Visualization

```

[ ]: dataset_id = "table"
table_id = "junction_table"

table_ref = client.dataset(dataset_id).table(table_id)

junction_table = client.list_rows(table_ref).to_dataframe()

junction_table.head()

```

```
[ ]:  SUBJECT_ID  GENDER      DOB  EXPIRE_FLAG  AGE_AT_ADMISSION  \
0         71093      M  2147-07-29  00:00:00+00:00          0          54
1         75536      F  2060-01-02  00:00:00+00:00          0          57
2         70191      M  2088-07-26  00:00:00+00:00          1          83
3         26942      F  2110-08-25  00:00:00+00:00          0          81
4          5890      M  2104-03-22  00:00:00+00:00          1          73
```

```

      HADM_ID      ADMITTIME      DISCHTIME  ADMISSION_TYPE  \
0  161963  2201-02-25  00:55:00+00:00  2201-02-27  13:05:00+00:00  EMERGENCY
1  139446  2117-11-25  00:30:00+00:00  2117-12-01  16:30:00+00:00  EMERGENCY
2  116326  2171-09-14  00:14:00+00:00  2171-09-23  15:00:00+00:00  EMERGENCY
3  176388  2191-09-02  00:54:00+00:00  2191-09-09  15:00:00+00:00  EMERGENCY
4  158408  2177-11-16  00:33:00+00:00  2177-12-06  19:04:00+00:00  EMERGENCY

```

```

      ADMISSION_LOCATION  ...      ETHNICITY  \
0  EMERGENCY ROOM ADMIT  ...      WHITE
1  EMERGENCY ROOM ADMIT  ...      WHITE
2  CLINIC REFERRAL/PREMAT  ...      WHITE
3  EMERGENCY ROOM ADMIT  ...      WHITE

```

4 EMERGENCY ROOM ADMIT ... BLACK/AFRICAN AMERICAN

	ADMISSION_DIAGNOSIS_TEXT	ADMISSION_HOUR	ICUSTAY_ID	\
0	VOMITING AND DIARRHEA	0	221142	
1	FEVER	0	223298	
2	FEVER	0	294563	
3	DIARRHEA-HYPOTENSION	0	255416	
4	SEPSIS	0	230915	

	ICU_INTIME	ICU_OUTTIME	ICU_LOS	\
0	2201-02-25 00:56:12+00:00	2201-02-26 09:15:43+00:00	1.3469	
1	2117-11-25 00:31:16+00:00	2117-11-27 21:13:44+00:00	2.8628	
2	2171-09-14 00:15:32+00:00	2171-09-23 17:29:20+00:00	9.7179	
3	2191-09-02 00:54:49+00:00	2191-09-03 16:02:41+00:00	1.6305	
4	2177-11-16 00:34:30+00:00	2177-12-06 20:14:23+00:00	20.8194	

	FIRST_CAREUNIT	PRIMARY_ICD9_CODE	PRIMARY_ICD9_TITLE
0	MICU	0059	Food poisoning NOS
1	SICU	00845	Int inf clstrdium dfcile
2	CSRU	00845	Int inf clstrdium dfcile
3	TSICU	00845	Int inf clstrdium dfcile
4	MICU	00845	Int inf clstrdium dfcile

[5 rows x 21 columns]

4 Data Pre-Processing - BigQuery

In this section of the project, we will use **BigQuery** to perform some pre-processing on the `junction_table`, to deal with repeated rows and null values

4.1 Duplicate Rows Analysis

```
[ ]: start_time = time.time()
      query = """
      SELECT
        ICUSTAY_ID,
        COUNT(*) AS number_of_rows
      FROM
        `reliable-jet-452114-s2.table.junction_table`
      GROUP BY
        ICUSTAY_ID
      HAVING
        COUNT(*) > 1
      ORDER BY
        number_of_rows DESC;
      """
```

```

query_job = client.query(query)
results = query_job.result()

print("ICUSTAY_ID | number_of_rows")
print("-----")
for row in results:
    print(f"{row.ICUSTAY_ID} | {row.number_of_rows}")

end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")

```

```

ICUSTAY_ID | number_of_rows
-----
Query Execution Time: 0.88 seconds

```

4.1.1 Conclusion

Based on the results of the query, we can conclude that besides the rows that contain the value None that are not any duplicate rows.

4.2 Null Values Analysis

```

[ ]: query = """
SELECT
    -- Total number of rows in the table
    COUNT(*) AS total_rows,

    -- Count of NULLs for each column
    COUNTIF(SUBJECT_ID IS NULL) AS null_SUBJECT_ID,
    COUNTIF(GENDER IS NULL) AS null_GENDER,
    COUNTIF(DOB IS NULL) AS null_DOB,
    COUNTIF(EXPIRE_FLAG IS NULL) AS null_EXPIRE_FLAG,
    COUNTIF(AGE_AT_ADMISSION IS NULL) AS null_AGE_AT_ADMISSION,
    COUNTIF(HADM_ID IS NULL) AS null_HADM_ID,
    COUNTIF(ADMITTIME IS NULL) AS null_ADMITTIME,
    COUNTIF(DISCHTIME IS NULL) AS null_DISCHTIME,
    COUNTIF(ADMISSION_TYPE IS NULL) AS null_ADMISSION_TYPE,
    COUNTIF(ADMISSION_LOCATION IS NULL) AS null_ADMISSION_LOCATION,
    COUNTIF(INSURANCE IS NULL) AS null_INSURANCE,
    COUNTIF(ETHNICITY IS NULL) AS null_ETHNICITY,
    COUNTIF(ADMISSION_DIAGNOSIS_TEXT IS NULL) AS null_ADMISSION_DIAGNOSIS_TEXT,
    COUNTIF(ADMISSION_HOUR IS NULL) AS null_ADMISSION_HOUR,
    COUNTIF(ICUSTAY_ID IS NULL) AS null_ICUSTAY_ID,
    COUNTIF(ICU_INTIME IS NULL) AS null_ICU_INTIME,
    COUNTIF(ICU_OUTTIME IS NULL) AS null_ICU_OUTTIME,
    COUNTIF(ICU_LOS IS NULL) AS null_ICU_LOS,

```

```

COUNTIF(FIRST_CAREUNIT IS NULL) AS null_FIRST_CAREUNIT,
COUNTIF(PRIMARY_ICD9_CODE IS NULL) AS null_PRIMARY_ICD9_CODE,
COUNTIF(PRIMARY_ICD9_TITLE IS NULL) AS null_PRIMARY_ICD9_TITLE,
COUNTIF(NUM_CALLOUTS IS NULL) AS null_NUM_CALLOUTS,

-- Percentage of NULLs for each column (formatted to 2 decimal places)
ROUND(100.0 * COUNTIF(SUBJECT_ID IS NULL) / COUNT(*), 2) AS_
↳pct_null_SUBJECT_ID,
ROUND(100.0 * COUNTIF(GENDER IS NULL) / COUNT(*), 2) AS pct_null_GENDER,
ROUND(100.0 * COUNTIF(DOB IS NULL) / COUNT(*), 2) AS pct_null_DOB,
ROUND(100.0 * COUNTIF(EXPIRE_FLAG IS NULL) / COUNT(*), 2) AS_
↳pct_null_EXPIRE_FLAG,
ROUND(100.0 * COUNTIF(AGE_AT_ADMISSION IS NULL) / COUNT(*), 2) AS_
↳pct_null_AGE_AT_ADMISSION,
ROUND(100.0 * COUNTIF(HADM_ID IS NULL) / COUNT(*), 2) AS pct_null_HADM_ID,
ROUND(100.0 * COUNTIF(ADMITTIME IS NULL) / COUNT(*), 2) AS pct_null_ADMITTIME,
ROUND(100.0 * COUNTIF(DISCHTIME IS NULL) / COUNT(*), 2) AS pct_null_DISCHTIME,
ROUND(100.0 * COUNTIF(ADMISSION_TYPE IS NULL) / COUNT(*), 2) AS_
↳pct_null_ADMISSION_TYPE,
ROUND(100.0 * COUNTIF(ADMISSION_LOCATION IS NULL) / COUNT(*), 2) AS_
↳pct_null_ADMISSION_LOCATION,
ROUND(100.0 * COUNTIF(INSURANCE IS NULL) / COUNT(*), 2) AS pct_null_INSURANCE,
ROUND(100.0 * COUNTIF(ETHNICITY IS NULL) / COUNT(*), 2) AS pct_null_ETHNICITY,
ROUND(100.0 * COUNTIF(ADMISSION_DIAGNOSIS_TEXT IS NULL) / COUNT(*), 2) AS_
↳pct_null_ADMISSION_DIAGNOSIS_TEXT,
ROUND(100.0 * COUNTIF(ADMISSION_HOUR IS NULL) / COUNT(*), 2) AS_
↳pct_null_ADMISSION_HOUR,
ROUND(100.0 * COUNTIF(ICUSTAY_ID IS NULL) / COUNT(*), 2) AS_
↳pct_null_ICUSTAY_ID,
ROUND(100.0 * COUNTIF(ICU_INTIME IS NULL) / COUNT(*), 2) AS_
↳pct_null_ICU_INTIME,
ROUND(100.0 * COUNTIF(ICU_OUTTIME IS NULL) / COUNT(*), 2) AS_
↳pct_null_ICU_OUTTIME,
ROUND(100.0 * COUNTIF(ICU_LOS IS NULL) / COUNT(*), 2) AS pct_null_ICU_LOS,
ROUND(100.0 * COUNTIF(FIRST_CAREUNIT IS NULL) / COUNT(*), 2) AS_
↳pct_null_FIRST_CAREUNIT,
ROUND(100.0 * COUNTIF(PRIMARY_ICD9_CODE IS NULL) / COUNT(*), 2) AS_
↳pct_null_PRIMARY_ICD9_CODE,
ROUND(100.0 * COUNTIF(PRIMARY_ICD9_TITLE IS NULL) / COUNT(*), 2) AS_
↳pct_null_PRIMARY_ICD9_TITLE,
ROUND(100.0 * COUNTIF(NUM_CALLOUTS IS NULL) / COUNT(*), 2) AS_
↳pct_null_NUM_CALLOUTS
FROM
`reliable-jet-452114-s2.table.junction_table`;
"""

```

```

# Execute the query and load the results into a Pandas DataFrame
null_analysis_df = client.query(query).to_dataframe()

# Display the DataFrame
# The result is a single row, so transposing makes it easier to read
print("Null Value Analysis (Transposed View):")
print(null_analysis_df.T)

```

Null Value Analysis (Transposed View):

	0
total_rows	62722
null_SUBJECT_ID	0
null_GENDER	0
null_DOB	0
null_EXPIRE_FLAG	0
null_AGE_AT_ADMISSION	0
null_HADM_ID	0
null_ADMITTIME	0
null_DISCHTIME	0
null_ADMISSION_TYPE	0
null_ADMISSION_LOCATION	0
null_INSURANCE	0
null_ETHNICITY	0
null_ADMISSION_DIAGNOSIS_TEXT	25
null_ADMISSION_HOUR	0
null_ICUSTAY_ID	1190
null_ICU_INTIME	1190
null_ICU_OUTTIME	1200
null_ICU_LOS	1200
null_FIRST_CAREUNIT	1190
null_PRIMARY_ICD9_CODE	47
null_PRIMARY_ICD9_TITLE	839
null_NUM_CALLOUTS	31366
pct_null_SUBJECT_ID	0.0
pct_null_GENDER	0.0
pct_null_DOB	0.0
pct_null_EXPIRE_FLAG	0.0
pct_null_AGE_AT_ADMISSION	0.0
pct_null_HADM_ID	0.0
pct_null_ADMITTIME	0.0
pct_null_DISCHTIME	0.0
pct_null_ADMISSION_TYPE	0.0
pct_null_ADMISSION_LOCATION	0.0
pct_null_INSURANCE	0.0
pct_null_ETHNICITY	0.0
pct_null_ADMISSION_DIAGNOSIS_TEXT	0.04
pct_null_ADMISSION_HOUR	0.0

pct_null_ICUSTAY_ID	1.9
pct_null_ICU_INTIME	1.9
pct_null_ICU_OUTTIME	1.91
pct_null_ICU_LOS	1.91
pct_null_FIRST_CAREUNIT	1.9
pct_null_PRIMARY_ICD9_CODE	0.07
pct_null_PRIMARY_ICD9_TITLE	1.34
pct_null_NUM_CALLOUTS	50.01

4.2.1 Conclusion

There are 62722 rows. The core patient and admission data is complete. However, approximately 1.9% of rows have nulls for ICU-related features, including the target variable ICU_LOS, primarily representing non-ICU admissions or missing discharge times. The primary diagnosis title has minor missingness (~1.3%). The biggest problem is the NUM_CALLOUTS feature, that has a high null rate (50%), requiring specific handling.

4.2.2 Approach to Missing Values in NUM_CALLOUTS

Knowing that the feature NUM_CALLOUTS is not too essential to predict the length of stay, we will remove this column from the table to avoid future processing problems.

```
[ ]: query = """
CREATE OR REPLACE TABLE `reliable-jet-452114-s2.table.junction_table` AS
SELECT
    * EXCEPT(NUM_CALLOUTS) -- Selects all columns except NUM_CALLOUTS
FROM
    `reliable-jet-452114-s2.table.junction_table`;
"""
query_job = client.query(query)
query_job.result()
```

```
[ ]: <google.cloud.bigquery.table._EmptyRowIterator at 0x7eefb0c17eb0>
```

4.2.3 Approach to other Missing Values

Knowing that the rows with missing values in other columns, including the target column ICU_LOS represent a very small percentage of the dataset, we decided the simplest and most effective solution would be to remove these rows.

```
[ ]: query = """
CREATE OR REPLACE TABLE `reliable-jet-452114-s2.table.junction_table` AS
SELECT
    * -- Select all columns from the filtered rows
FROM
    `reliable-jet-452114-s2.table.junction_table`
WHERE
    -- Ensure all potentially nullable columns are NOT NULL
    ADMISSION_DIAGNOSIS_TEXT IS NOT NULL AND
```



```

ICUSTAY_ID IS NOT NULL AND
ICU_INTIME IS NOT NULL AND
ICU_OUTTIME IS NOT NULL AND
ICU_LOS IS NOT NULL AND
FIRST_CAREUNIT IS NOT NULL AND
PRIMARY_ICD9_CODE IS NOT NULL AND
PRIMARY_ICD9_TITLE IS NOT NULL;
"""
query_job = client.query(query)
query_job.result()

```

```
[ ]: <google.cloud.bigquery.table._EmptyRowIterator at 0x7eefb0db8f10>
```

4.2.4 Running the Null Values Query again

In order to access if our changes were successful, we decided to run the Null Values query again.

```

[ ]: query = """
SELECT
    -- Total number of rows in the table
    COUNT(*) AS total_rows,

    -- Count of NULLs for each column
    COUNTIF(SUBJECT_ID IS NULL) AS null_SUBJECT_ID,
    COUNTIF(GENDER IS NULL) AS null_GENDER,
    COUNTIF(DOB IS NULL) AS null_DOB,
    COUNTIF(EXPIRE_FLAG IS NULL) AS null_EXPIRE_FLAG,
    COUNTIF(AGE_AT_ADMISSION IS NULL) AS null_AGE_AT_ADMISSION,
    COUNTIF(HADM_ID IS NULL) AS null_HADM_ID,
    COUNTIF(ADMITTIME IS NULL) AS null_ADMITTIME,
    COUNTIF(DISCHTIME IS NULL) AS null_DISCHTIME,
    COUNTIF(ADMISSION_TYPE IS NULL) AS null_ADMISSION_TYPE,
    COUNTIF(ADMISSION_LOCATION IS NULL) AS null_ADMISSION_LOCATION,
    COUNTIF(INSURANCE IS NULL) AS null_INSURANCE,
    COUNTIF(ETHNICITY IS NULL) AS null_ETHNICITY,
    COUNTIF(ADMISSION_DIAGNOSIS_TEXT IS NULL) AS null_ADMISSION_DIAGNOSIS_TEXT,
    COUNTIF(ADMISSION_HOUR IS NULL) AS null_ADMISSION_HOUR,
    COUNTIF(ICUSTAY_ID IS NULL) AS null_ICUSTAY_ID,
    COUNTIF(ICU_INTIME IS NULL) AS null_ICU_INTIME,
    COUNTIF(ICU_OUTTIME IS NULL) AS null_ICU_OUTTIME,
    COUNTIF(ICU_LOS IS NULL) AS null_ICU_LOS,
    COUNTIF(FIRST_CAREUNIT IS NULL) AS null_FIRST_CAREUNIT,
    COUNTIF(PRIMARY_ICD9_CODE IS NULL) AS null_PRIMARY_ICD9_CODE,
    COUNTIF(PRIMARY_ICD9_TITLE IS NULL) AS null_PRIMARY_ICD9_TITLE,

    -- Percentage of NULLs for each column (formatted to 2 decimal places)

```

```

ROUND(100.0 * COUNTIF(SUBJECT_ID IS NULL) / COUNT(*), 2) AS
↳pct_null_SUBJECT_ID,
ROUND(100.0 * COUNTIF(GENDER IS NULL) / COUNT(*), 2) AS pct_null_GENDER,
ROUND(100.0 * COUNTIF(DOB IS NULL) / COUNT(*), 2) AS pct_null_DOB,
ROUND(100.0 * COUNTIF(EXPIRE_FLAG IS NULL) / COUNT(*), 2) AS
↳pct_null_EXPIRE_FLAG,
ROUND(100.0 * COUNTIF(AGE_AT_ADMISSION IS NULL) / COUNT(*), 2) AS
↳pct_null_AGE_AT_ADMISSION,
ROUND(100.0 * COUNTIF(HADM_ID IS NULL) / COUNT(*), 2) AS pct_null_HADM_ID,
ROUND(100.0 * COUNTIF(ADMITTIME IS NULL) / COUNT(*), 2) AS pct_null_ADMITTIME,
ROUND(100.0 * COUNTIF(DISCHTIME IS NULL) / COUNT(*), 2) AS pct_null_DISCHTIME,
ROUND(100.0 * COUNTIF(ADMISSION_TYPE IS NULL) / COUNT(*), 2) AS
↳pct_null_ADMISSION_TYPE,
ROUND(100.0 * COUNTIF(ADMISSION_LOCATION IS NULL) / COUNT(*), 2) AS
↳pct_null_ADMISSION_LOCATION,
ROUND(100.0 * COUNTIF(INSURANCE IS NULL) / COUNT(*), 2) AS pct_null_INSURANCE,
ROUND(100.0 * COUNTIF(ETHNICITY IS NULL) / COUNT(*), 2) AS pct_null_ETHNICITY,
ROUND(100.0 * COUNTIF(ADMISSION_DIAGNOSIS_TEXT IS NULL) / COUNT(*), 2) AS
↳pct_null_ADMISSION_DIAGNOSIS_TEXT,
ROUND(100.0 * COUNTIF(ADMISSION_HOUR IS NULL) / COUNT(*), 2) AS
↳pct_null_ADMISSION_HOUR,
ROUND(100.0 * COUNTIF(ICUSTAY_ID IS NULL) / COUNT(*), 2) AS
↳pct_null_ICUSTAY_ID,
ROUND(100.0 * COUNTIF(ICU_INTIME IS NULL) / COUNT(*), 2) AS
↳pct_null_ICU_INTIME,
ROUND(100.0 * COUNTIF(ICU_OUTTIME IS NULL) / COUNT(*), 2) AS
↳pct_null_ICU_OUTTIME,
ROUND(100.0 * COUNTIF(ICU_LOS IS NULL) / COUNT(*), 2) AS pct_null_ICU_LOS,
ROUND(100.0 * COUNTIF(FIRST_CAREUNIT IS NULL) / COUNT(*), 2) AS
↳pct_null_FIRST_CAREUNIT,
ROUND(100.0 * COUNTIF(PRIMARY_ICD9_CODE IS NULL) / COUNT(*), 2) AS
↳pct_null_PRIMARY_ICD9_CODE,
ROUND(100.0 * COUNTIF(PRIMARY_ICD9_TITLE IS NULL) / COUNT(*), 2) AS
↳pct_null_PRIMARY_ICD9_TITLE
FROM
    `reliable-jet-452114-s2.table.junction_table`;
"""

# Execute the query and load the results into a Pandas DataFrame
null_analysis_df = client.query(query).to_dataframe()

# Display the DataFrame
# The result is a single row, so transposing it makes it easier to read
print("Null Value Analysis (Transposed View - NUM_CALLOUTS Excluded):")
print(null_analysis_df.T)

```

Null Value Analysis (Transposed View - NUM_CALLOUTS Excluded):

	0
total_rows	60748
null_SUBJECT_ID	0
null_GENDER	0
null_DOB	0
null_EXPIRE_FLAG	0
null_AGE_AT_ADMISSION	0
null_HADM_ID	0
null_ADMITTIME	0
null_DISCHTIME	0
null_ADMISSION_TYPE	0
null_ADMISSION_LOCATION	0
null_INSURANCE	0
null_ETHNICITY	0
null_ADMISSION_DIAGNOSIS_TEXT	0
null_ADMISSION_HOUR	0
null_ICUSTAY_ID	0
null_ICU_INTIME	0
null_ICU_OUTTIME	0
null_ICU_LOS	0
null_FIRST_CAREUNIT	0
null_PRIMARY_ICD9_CODE	0
null_PRIMARY_ICD9_TITLE	0
pct_null_SUBJECT_ID	0.0
pct_null_GENDER	0.0
pct_null_DOB	0.0
pct_null_EXPIRE_FLAG	0.0
pct_null_AGE_AT_ADMISSION	0.0
pct_null_HADM_ID	0.0
pct_null_ADMITTIME	0.0
pct_null_DISCHTIME	0.0
pct_null_ADMISSION_TYPE	0.0
pct_null_ADMISSION_LOCATION	0.0
pct_null_INSURANCE	0.0
pct_null_ETHNICITY	0.0
pct_null_ADMISSION_DIAGNOSIS_TEXT	0.0
pct_null_ADMISSION_HOUR	0.0
pct_null_ICUSTAY_ID	0.0
pct_null_ICU_INTIME	0.0
pct_null_ICU_OUTTIME	0.0
pct_null_ICU_LOS	0.0
pct_null_FIRST_CAREUNIT	0.0
pct_null_PRIMARY_ICD9_CODE	0.0
pct_null_PRIMARY_ICD9_TITLE	0.0

4.2.5 Conclusion

There are 60748 rows (~96.8% of the original table) and there are no missing values, as expected.

5 Dataset Analysis - BigQuery

In this section, we use **BigQuery** to perform a dataset analysis on three different topics, **Patients**, **Medical Data** and **Correlations**.

In **Patients** and **Medical Data** we will analyse the distribution of several features and try to find connections between them and the target feature (**Length of ICU Stay**), to get a better understanding of what features will be more relevant to our ML model.

In **Correlations** we will try to find connections between different features of the dataset, to gain an even better insight of the data.

5.1 Patients

We will start this analysis with an in-depth analysis of the features regarding patients characteristics, which include:

- Gender
- Age
- Insurance
- Ethnicity

5.1.1 Gender Analysis

```
[ ]: start_time = time.time()

# 1. Define the BigQuery SQL query to get gender counts
query = """
SELECT
    GENDER,
    COUNT(*) AS count
FROM
    `reliable-jet-452114-s2.table.junction_table`
GROUP BY
    GENDER
ORDER BY
    GENDER;
"""

# 2. Execute the query and load results into a Pandas DataFrame
gender_distribution_df = client.query(query).to_dataframe()

# 3. Create the visualization using Matplotlib and Seaborn
if not gender_distribution_df.empty:
    plt.figure(figsize=(6, 4))
```

```

sns.barplot(x='GENDER', y='count', data=gender_distribution_df,
↪palette='viridis') # Use Seaborn for nicer plotting

# Add labels and title
plt.xlabel("Gender")
plt.ylabel("Number of Patients")
plt.title("Distribution of Patients by Gender")

# Add count labels on top of bars
for index, row in gender_distribution_df.iterrows():
    plt.text(index, row['count'], row['count'], color='black', ha="center",
↪va='bottom')

# Display the plot
plt.tight_layout()
plt.show()
else:
    print("No data returned from the query to plot.")

# 4. Print the Dataframe
print("\nGender Distribution Data:")
print(gender_distribution_df)

end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")

```


<ipython-input-40-ebfdb9af863e>:22: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```

sns.barplot(x='GENDER', y='count', data=gender_distribution_df,
palette='viridis') # Use Seaborn for nicer plotting

```



Gender Distribution Data:

	GENDER	count
0	F	25013
1	M	33026

Query Execution Time: 2.12 seconds

5.1.2 Gender & Length of ICU Stay

```
[ ]: start_time = time.time()

# 1. Define the BigQuery SQL query to get Gender and Length of ICU Stay
query = """
SELECT
    GENDER,
    ICU_LOS
FROM
    `reliable-jet-452114-s2.table.junction_table`
WHERE
    ICU_LOS IS NOT NULL -- Ensure LOS is not null
    AND GENDER IS NOT NULL -- Ensure Gender is not null
ORDER BY GENDER;
"""
```

```

# 2. Execute the query and load results into a Pandas DataFrame
gender_los_df = client.query(query).to_dataframe()

# 3. Create the visualization using Matplotlib and Seaborn
if not gender_los_df.empty:
    plt.figure(figsize=(8, 6))

    # Create the box plot
    sns.boxplot(x='GENDER', y='ICU_LOS', data=gender_los_df, palette='viridis',
        showfliers=True)
    # showfliers=True to includes outliers

    # Add labels and title
    plt.xlabel("Gender")
    plt.ylabel("Length of ICU Stay (Days)")
    plt.title("Distribution of Length of ICU Stay by Gender")
    plt.ylim(bottom=0) # Ensure y-axis starts at 0

    # Display the plot
    plt.tight_layout()
    plt.show()

    # Calculate and print summary statistics
    print("\nSummary Statistics for ICU LOS by Gender:")
    print(gender_los_df.groupby('GENDER')['ICU_LOS'].describe())

else:
    print("No data returned from the query to plot.")

end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")

```

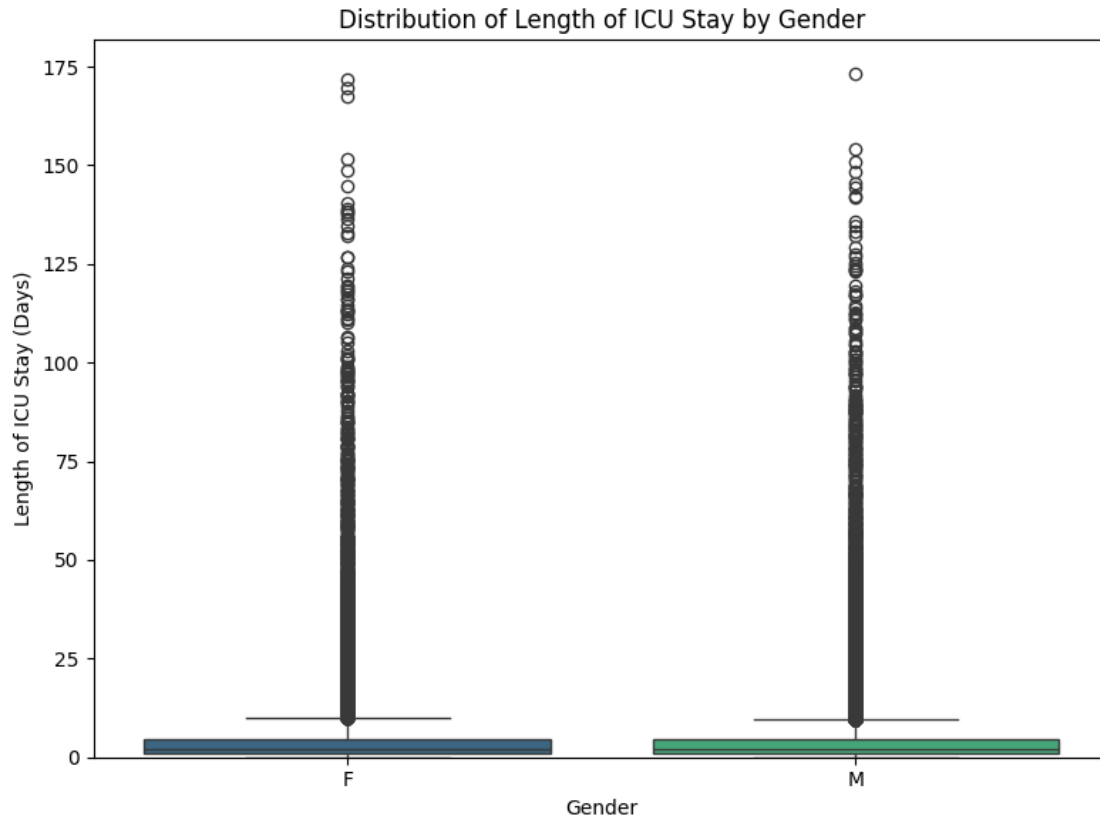
<ipython-input-41-3e283b9877ca>:24: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```

sns.boxplot(x='GENDER', y='ICU_LOS', data=gender_los_df, palette='viridis',
showfliers=True)

```



Summary Statistics for ICU LOS by Gender:

	count	mean	std	min	25%	50%	75%	\
GENDER								
F	25013.0	5.066738	10.104802	0.0001	1.090800	2.1168	4.6438	
M	33026.0	4.928852	9.606139	0.0002	1.111725	2.0739	4.4723	

max
GENDER

F 171.6227

M 173.0725

Query Execution Time: 3.26 seconds

5.1.3 Conclusion

The distributions of patients by gender is fairly simetric, with a slight imbalance for MALE (~56.0%).

There is no discernible connection between Gender and Length of ICU Stay.

5.1.4 Age Analysis

```
[ ]: # 1. Define the BigQuery SQL query to get the Age at Admission for each row
query = """
SELECT
    AGE_AT_ADMISSION
FROM
    `reliable-jet-452114-s2.table.junction_table`
WHERE
    AGE_AT_ADMISSION IS NOT NULL; -- Ensure age is not null
"""

# 2. Execute the query and load results into a Pandas DataFrame
age_distribution_df = client.query(query).to_dataframe()


# 3. Create the visualization using Matplotlib and Seaborn
if not age_distribution_df.empty:
    plt.figure(figsize=(10, 6))

    # Create a histogram with a Kernel Density Estimate (KDE) overlay
    sns.histplot(data=age_distribution_df, x='AGE_AT_ADMISSION', kde=True,
    ↪bins=30)
    # kde=True adds a smooth line representing the distribution shape

    # Add labels and title
    plt.xlabel("Age at Admission (Years)")
    plt.ylabel("Number of Patients")
    plt.title("Distribution of Patient Age at Admission")

    # Display the plot
    plt.tight_layout()
    plt.show()
else:
    print("No data returned from the query to plot.")

# 4. Print descriptive statistics for Age
print("\nAge at Admission Statistics:")
print(age_distribution_df['AGE_AT_ADMISSION'].describe())
```



Age at Admission Statistics:

```
count      58039.0
mean       53.943693
std        26.5604
min         0.0
25%        43.0
50%        61.0
75%        74.0
max        89.0
```

Name: AGE_AT_ADMISSION, dtype: Float64

5.1.5 Eliminating rows with Age bigger then 120 years

Given the histogram results, we decided to eliminate the rows with outliers that most definitely represent input errors, given that there are no people near 300 years old.

```
[ ]: query = """
CREATE OR REPLACE TABLE `reliable-jet-452114-s2.table.junction_table` AS
SELECT
    * -- Select all columns from the filtered rows
FROM
    `reliable-jet-452114-s2.table.junction_table`
WHERE
    AGE_AT_ADMISSION <= 120;
"""
```

```
query_job = client.query(query)
query_job.result()
```

```
[ ]: <google.cloud.bigquery.table._EmptyRowIterator at 0x7eef8976be80>
```

```
[ ]: table_id = "reliable-jet-452114-s2.table.junction_table"

# Get the table metadata from BigQuery
table = client.get_table(table_id)

# Access the num_rows attribute from the metadata
num_rows = table.num_rows

print(f"The new number of rows in the table is: {num_rows}")
```

The new number of rows in the table is: 58039

After eliminating these input error mistakes, we ended up with a table containing 58039 rows (~95.5% of the previous table).

5.1.6 Age & Length of ICU Stay

```
[ ]: # 1. Define the BigQuery SQL query to get Age and ICU Length of ICU Stay
query = """
SELECT
    AGE_AT_ADMISSION,
    ICU_LOS
FROM
    `reliable-jet-452114-s2.table.junction_table`
WHERE
    ICU_LOS IS NOT NULL
    AND AGE_AT_ADMISSION IS NOT NULL
"""

# 2. Execute the query and load results into a Pandas DataFrame
age_los_df = client.query(query).to_dataframe()

# 3. Create the visualization using Matplotlib and Seaborn
if not age_los_df.empty:
    plt.figure(figsize=(10, 7))

    # Create a scatter plot
    sns.scatterplot(x='AGE_AT_ADMISSION', y='ICU_LOS', data=age_los_df, alpha=0.
↪3, s=15)

    # Add labels and title
    plt.xlabel("Age at Admission (Years)")
    plt.ylabel("ICU Length of ICU Stay (Days)")
```

```

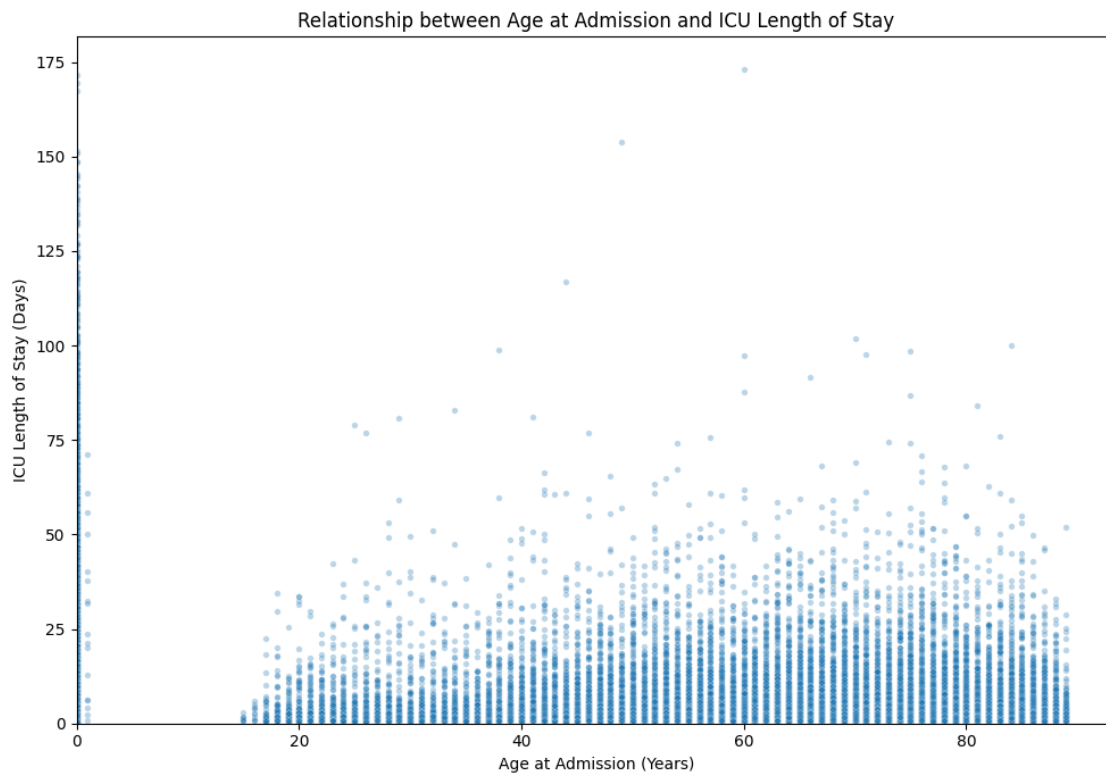
plt.title("Relationship between Age at Admission and ICU Length of ICU_
↳Stay")
plt.ylim(bottom=0) # Ensure y-axis starts at 0
plt.xlim(left=0)   # Ensure x-axis starts at 0

# Display the plot
plt.tight_layout()
plt.show()

# Calculate correlation
correlation = age_los_df['AGE_AT_ADMISSION'].corr(age_los_df['ICU_LOS'])
print(f"\nCorrelation between Age at Admission and ICU LOS: {correlation:.
↳3f}")

else:
    print("No data returned from the query to plot.")

```



Correlation between Age at Admission and ICU LOS: -0.161

5.1.7 Conclusion

By analysing the histogram of **Age**, we can see that the most common age intervals in the dataset are 0 and 50–80. The dataset proposely omimits data from minor patients with the excpetion of newborns.

Regarding the connection between **Age** and **Length of ICU Stay**, we can conclude that: 1. The vast majority of ICU stays are relatively short (concentrated roughly below 25 days) across all adult age groups. 2. While short stays are common for all ages, the spread of the values of **Length of ICU Stay** appears to widen slightly for older patients, which means that very long stays seem slightly more prevalent among older age groups.

5.1.8 Insurance Analysis

```
[ ]: # 1. Define the BigQuery SQL query to get the Insurance for each row
query = """
SELECT
    INSURANCE,
    COUNT(*) AS count
FROM
    `reliable-jet-452114-s2.table.junction_table`
WHERE
    INSURANCE IS NOT NULL
GROUP BY
    INSURANCE
ORDER BY
    count DESC; -- Order by count to see most common types first
"""

# 2. Execute the query and load results into a Pandas DataFrame
insurance_distribution_df = client.query(query).to_dataframe()

# 3. Create the distribution visualization using Matplotlib and SeaBorn
if not insurance_distribution_df.empty:
    plt.figure(figsize=(10, 5))
    sns.barplot(x='count', y='INSURANCE', data=insurance_distribution_df,
        ↪palette='magma', orient='h') # Horizontal bar chart


    # Add labels and title
    plt.xlabel("Number of Patients")
    plt.ylabel("Insurance Type")
    plt.title("Distribution of Patients by Insurance Type")

    # Display the plot
    plt.tight_layout()
    plt.show()
else:
    print("No data returned for insurance distribution query.")
```

<ipython-input-4-bde86873813b>:22: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x='count', y='INSURANCE', data=insurance_distribution_df,
palette='magma', orient='h') # Horizontal bar chart
```



5.1.9 Insurance & Length of ICU Stay

```
[ ]: # 1. Define the BigQuery SQL query to get Insurance and ICU Length of ICU Stay
query = """
SELECT
    INSURANCE,
    ICU_LOS
FROM
    `reliable-jet-452114-s2.table.junction_table`
WHERE
    ICU_LOS IS NOT NULL
    AND INSURANCE IS NOT NULL
ORDER BY INSURANCE;
"""

# 2. Execute the query and load results into a Pandas DataFrame
insurance_los_df = client.query(query).to_dataframe()

# 3. Create the visualization using Matplotlib and Seaborn
if not insurance_los_df.empty:
```

```

plt.figure(figsize=(12, 7))

# Create the box plot
sns.boxplot(x='INSURANCE', y='ICU_LOS', data=insurance_los_df,
palette='magma', showfliers=False)

# Add labels and title
plt.xlabel("Insurance Type")
plt.ylabel("Length of ICU Stay (Days)")
plt.title("Distribution of Length of ICU Stay by Insurance Type")
plt.xticks(rotation=45, ha='right')
plt.ylim(bottom=0)

# Display the plot
plt.tight_layout()
plt.show()

# Print summary statistics
print("\nSummary Statistics for ICU LOS by Insurance Type:")
print(insurance_los_df.groupby('INSURANCE')['ICU_LOS'].describe())

else:
    print("No data returned for insurance vs. LOS query.")

```

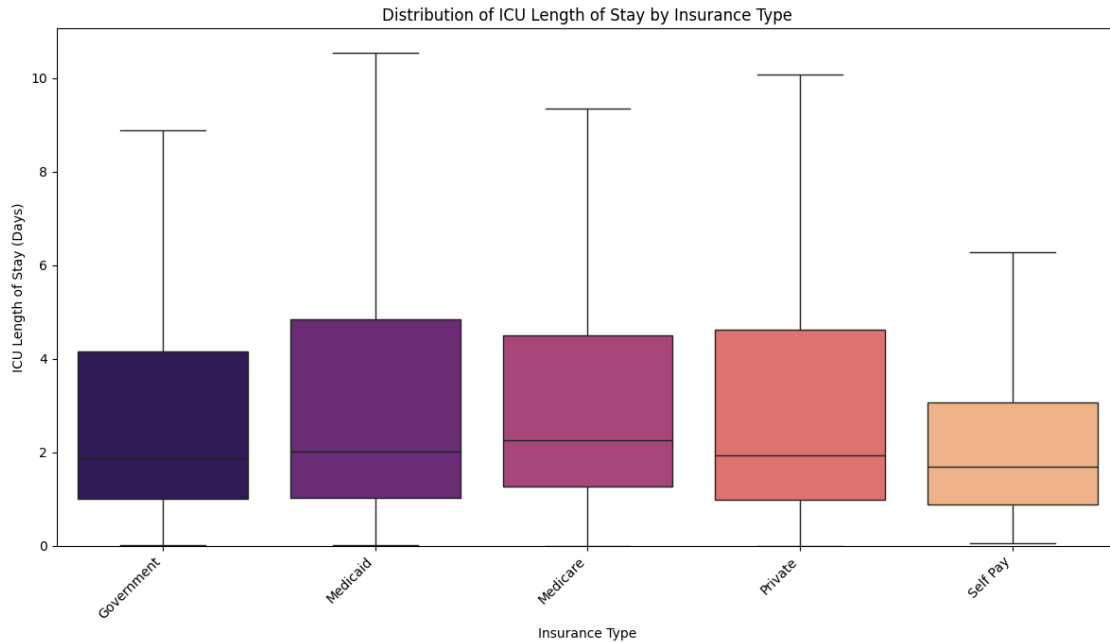
<ipython-input-5-d83a46a0aebb>:28: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```

sns.boxplot(x='INSURANCE', y='ICU_LOS', data=insurance_los_df,
palette='magma', showfliers=False) # Vertical, hide outliers for clarity

```



Summary Statistics for ICU LOS by Insurance Type:

	count	mean	std	min	25%	50%	75%	\
INSURANCE								
Government	1822.0	4.895602	9.793321	0.0057	0.998850	1.87805	4.167150	
Medicaid	5919.0	5.900825	12.800130	0.0036	1.025900	2.02050	4.835200	
Medicare	26843.0	4.306665	6.252719	0.0001	1.260800	2.26430	4.501300	
Private	22845.0	5.611555	12.140707	0.0002	0.982100	1.92950	4.625200	
Self Pay	610.0	3.062356	4.486948	0.0614	0.888425	1.68900	3.057675	

	max
INSURANCE	
Government	101.8397
Medicaid	169.4202
Medicare	173.0725
Private	171.6227
Self Pay	43.1465

5.1.10 Conclusion

By analysing the bar chart of **Insurance**, we can see that the most common insurance systems are **Medicare** followed by **Private**, and that there is a significantly lower number of patients using other systems.

Regarding the connection between **Insurance** and **Length of ICU Stay**, we can conclude that all systems have a similar distribution and mean in the comparison of the two metrics, with the exception of **Self Pay**, whose average value for **Length of ICU Stay** is significantly lower.

5.1.11 Ethnicity Analysis

```
[ ]: # 1. Define the BigQuery SQL query to get the top 5 ethnicity counts
query = """
SELECT
    ETHNICITY,
    COUNT(*) AS count
FROM
    `reliable-jet-452114-s2.table.junction_table`
WHERE
    ETHNICITY IS NOT NULL
GROUP BY
    ETHNICITY
ORDER BY
    count DESC -- Order by count to see most common groups first
LIMIT 5; -- Limit the results to the top 5
"""

# 2. Execute the query and load results into a Pandas DataFrame
ethnicity_distribution_df = client.query(query).to_dataframe()

# 3. Create the distribution visualization using Matplotlib and Seaborn
if not ethnicity_distribution_df.empty:
    plt.figure(figsize=(10, 6))
    sns.barplot(x='count', y='ETHNICITY', data=ethnicity_distribution_df,
                palette='crest', orient='h') # Horizontal bar chart

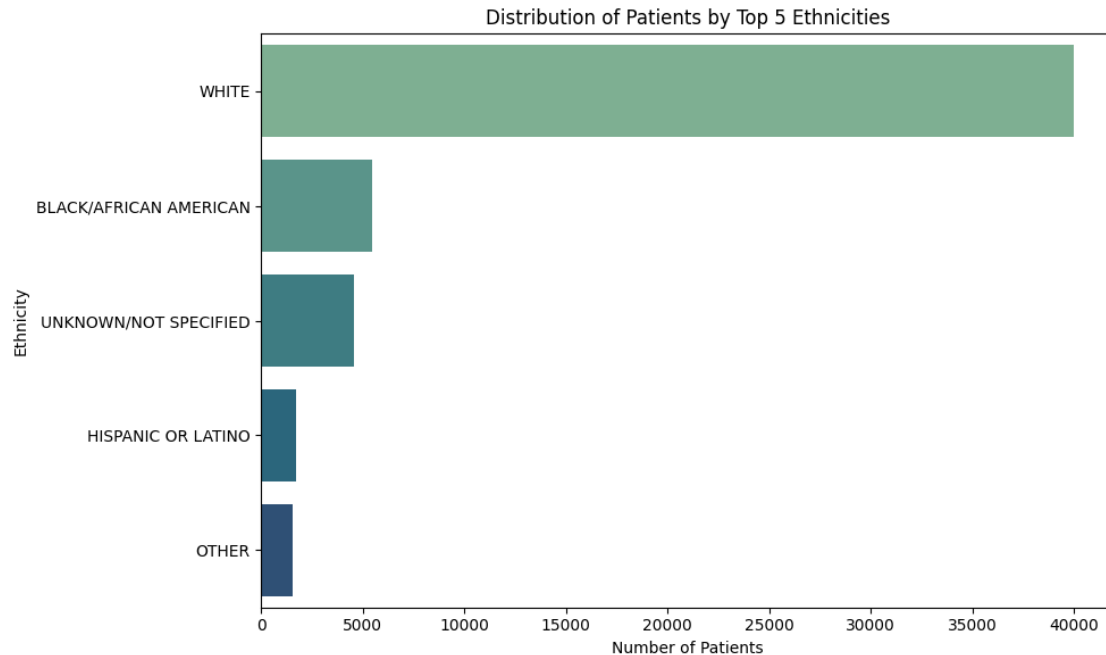
    # Add labels and title
    plt.xlabel("Number of Patients")
    plt.ylabel("Ethnicity")
    plt.title("Distribution of Patients by Top 5 Ethnicities")

    # Display the plot
    plt.tight_layout()
    plt.show()
else:
    print("No data returned for ethnicity distribution query.")
```

<ipython-input-16-68e434b7d292>:23: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x='count', y='ETHNICITY', data=ethnicity_distribution_df,
palette='crest', orient='h') # Horizontal bar chart
```



5.1.12 Ethnicity & Length of ICU Stay

```
[ ]: # 1. Define the BigQuery SQL query to get the Top 5 Ethnicities first
query_top_ethnicities = """
SELECT
    ETHNICITY
FROM
    `reliable-jet-452114-s2.table.junction_table`
WHERE
    ETHNICITY IS NOT NULL
GROUP BY
    ETHNICITY
ORDER BY
    COUNT(*) DESC
LIMIT 5;
"""

# Execute the query to get the top 5 list
top_ethnicities_df = client.query(query_top_ethnicities).to_dataframe()
top_ethnicities_list = top_ethnicities_df['ETHNICITY'].tolist()

# Check if the list is not empty before proceeding
if top_ethnicities_list:
    # Format the list for the SQL IN clause
```

```

    formatted_top_ethnicities = ", ".join([f'"{eth}"' for eth in top_ethnicities_list])

    # 2. Define the BigQuery SQL query to get Ethnicity and ICU Length of Stay
    for the Top 5
    query = f"""
    SELECT
        ETHNICITY,
        ICU_LOS
    FROM
        `reliable-jet-452114-s2.table.junction_table`
    WHERE
        ICU_LOS IS NOT NULL
        AND ETHNICITY IN ({formatted_top_ethnicities}) -- Filter for top 5
    ethnicities
        AND ICU_LOS < 25 -- Filter for patients with 25 or less days in ICU
    ORDER BY ETHNICITY;
    """

    # 3. Execute the query and load results into a Pandas DataFrame
    ethnicity_los_df = client.query(query).to_dataframe()

    # 4. Create the visualization using Matplotlib and Seaborn
    if not ethnicity_los_df.empty:
        plt.figure(figsize=(12, 7))

        # Create the box plot, explicitly using the top_ethnicities_list for
        order
        sns.boxplot(x='ETHNICITY', y='ICU_LOS', data=ethnicity_los_df,
        palette='crest', order=top_ethnicities_list, showfliers=False)

        # Add labels and title
        plt.xlabel("Ethnicity (Top 5)")
        plt.ylabel("Length of ICU Stay (Days)")
        plt.title("Distribution of Length of ICU Stay by Top 5 Ethnicities")
        plt.xticks(rotation=45, ha='right')
        plt.ylim(bottom=0)

        # Display the plot
        plt.tight_layout()
        plt.show()

        # Print summary statistics for the Top 5
        print("\nSummary Statistics for ICU LOS by Top 5 Ethnicities:")
        print(ethnicity_los_df.groupby('ETHNICITY')['ICU_LOS'].describe())

    else:

```

```

        print("No data returned for ethnicity vs. LOS query for top 5.")
    else:
        print("Could not determine the top 5 ethnicities.")

```

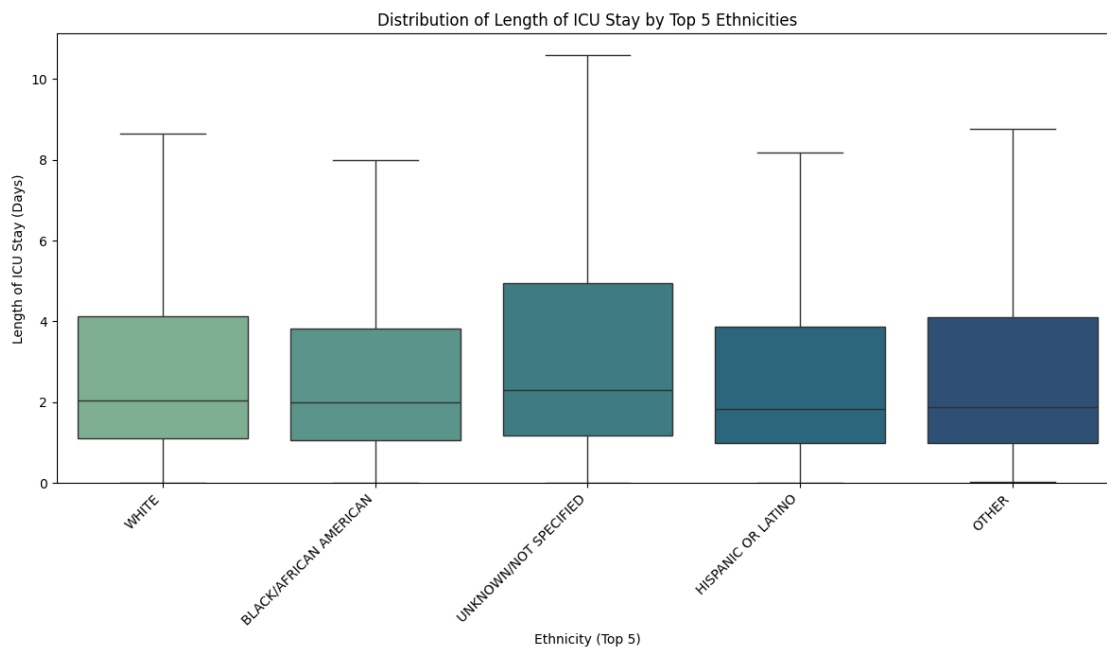
<ipython-input-18-0fa546c6be4d>:47: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```

sns.boxplot(x='ETHNICITY', y='ICU_LOS', data=ethnicity_los_df,
palette='crest', order=top_ethnicities_list, showfliers=False)

```



Summary Statistics for ICU LOS by Top 5 Ethnicities:

ETHNICITY	count	mean	std	min	25%	\
BLACK/AFRICAN AMERICAN	5247.0	3.315243	3.933401	0.0025	1.046900	
HISPANIC OR LATINO	1656.0	3.463486	4.396406	0.0019	0.980425	
OTHER	1454.0	3.584422	4.574565	0.0209	0.983450	
UNKNOWN/NOT SPECIFIED	4364.0	4.075112	4.571024	0.0014	1.181575	
WHITE	38802.0	3.594078	4.203717	0.0001	1.105100	
	50%	75%	max			
ETHNICITY						
BLACK/AFRICAN AMERICAN	1.98800	3.829750	24.9090			

HISPANIC OR LATINO	1.84185	3.869450	24.4639
OTHER	1.87655	4.109275	24.8130
UNKNOWN/NOT SPECIFIED	2.28950	4.951025	24.6889
WHITE	2.04480	4.122675	24.9968

5.1.13 Conclusion

There is a clear imbalance in the `Ethnicity` distribution, with most of the patients being `WHITE`.

There isn't a clear connection between `Ethnicity` and `ICU Length of Stay`, with the mean values for the `ICU Length of Stay` of each `Ethnicity` being similar.

5.2 Medical Data

We will continue the dataset analysis with an in-depth analysis of the features regarding patients medical records while they were hospitalized, which include:

- Type of Admission
- Hour of Admission
- Difference between Hour of Admission and Hour of Icu Admission
- Diagnostic Code
- First Care Unit
- Death
- Length of ICU Stay (Target)

5.2.1 Type of Admission Analysis

```
[ ]: # 1. Define the BigQuery SQL query to get admission type counts
query = """
SELECT
    ADMISSION_TYPE,
    COUNT(*) AS count
FROM
    `reliable-jet-452114-s2.table.junction_table`
WHERE
    ADMISSION_TYPE IS NOT NULL
GROUP BY
    ADMISSION_TYPE
ORDER BY
    count DESC; -- Order by count to see most common types first
"""

# 2. Execute the query and load results into a Pandas DataFrame
adm_type_distribution_df = client.query(query).to_dataframe()

# 3. Create the distribution visualization using Matplotlib and Seaborn
if not adm_type_distribution_df.empty:
    plt.figure(figsize=(8, 5))
```

```

    ax = sns.barplot(x='ADMISSION_TYPE', y='count',
data=adm_type_distribution_df, palette='cubehelix')

    # Add labels and title
    plt.xlabel("Admission Type")
    plt.ylabel("Number of Patients")
    plt.title("Distribution of Patients by Admission Type")
    plt.xticks(rotation=45, ha='right')

    # Add count labels on top of bars
    for container in ax.containers:
        ax.bar_label(container)

    # Display the plot
    plt.tight_layout()
    plt.show()
else:
    print("No data returned for admission type distribution query.")

```


<ipython-input-4-7d33f20a73c8>:22: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```

ax = sns.barplot(x='ADMISSION_TYPE', y='count', data=adm_type_distribution_df,
palette='cubehelix')

```



5.2.2 Type of Admission & Length of ICU Stay

```
[ ]: # 1. Define the BigQuery SQL query to get Admission Type and Length of ICU Stay
query = """
SELECT
    ADMISSION_TYPE,
    ICU_LOS
FROM
    `reliable-jet-452114-s2.table.junction_table`
WHERE
    ICU_LOS IS NOT NULL
    AND ADMISSION_TYPE IS NOT NULL
ORDER BY ADMISSION_TYPE;
"""

# 2. Execute the query and load results into a Pandas DataFrame
adm_type_los_df = client.query(query).to_dataframe()

# 3. Create the relationship visualization using Matplotlib and Seaborn
if not adm_type_los_df.empty:
    plt.figure(figsize=(10, 6))

    # Create the box plot
    sns.boxplot(x='ADMISSION_TYPE', y='ICU_LOS', data=adm_type_los_df,
                palette='cubehelix', showfliers=False) # Vertical, hide outliers

    # Add labels and title
    plt.xlabel("Admission Type")
    plt.ylabel("ICU Length of Stay (Days)")
    plt.title("Distribution of ICU Length of Stay by Admission Type")
    plt.xticks(rotation=45, ha='right')
    plt.ylim(bottom=0) # Ensure y-axis starts at 0


    # Display the plot
    plt.tight_layout()
    plt.show()

    # Print summary statistics
    print("\nSummary Statistics for ICU LOS by Admission Type:")
    print(adm_type_los_df.groupby('ADMISSION_TYPE')['ICU_LOS'].describe())
else:
    print("No data returned for admission type vs. LOS query.")
```

<ipython-input-5-168f0654216f>:22: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='ADMISSION_TYPE', y='ICU_LOS', data=adm_type_los_df,
palette='cubehelix', showfliers=False) # Vertical, hide outliers
```



Summary Statistics for ICU LOS by Admission Type:

	count	mean	std	min	25%	50% \
ADMISSION_TYPE						
ELECTIVE	7226.0	3.503513	6.008552	0.0004	1.140675	1.9895
EMERGENCY	41601.0	4.309662	6.436865	0.0001	1.229700	2.1897
NEWBORN	7866.0	9.878190	20.483139	0.0008	0.137675	0.7253
URGENT	1346.0	5.356694	8.359788	0.0025	1.241075	2.4351
		75%	max			
ADMISSION_TYPE						
ELECTIVE	3.283750	173.0725				
EMERGENCY	4.420000	169.4202				
NEWBORN	9.500975	171.6227				
URGENT	5.380250	97.4897				

5.2.3 Conclusion

The distributions of patients by Type of Admission is uneven, with most admissions being classified as EMERGENCY. The least recurrent type of admissions is Urgent.

There is an evident connection between Type of Admission and Length of ICU Stay, with NEWBORN having the highest mean values (close to 10 days) while the other types of admission have mean values between 3 to 5 days.

5.2.4 Hour of Admission Analysis

```
[ ]: start_time = time.time()

# 1. Define the BigQuery SQL query to get admission hour counts
query = """
SELECT
    ADMISSION_HOUR,
    COUNT(*) AS count
FROM
    `reliable-jet-452114-s2.table.junction_table`
WHERE
    ADMISSION_HOUR IS NOT NULL
GROUP BY
    ADMISSION_HOUR
ORDER BY
    ADMISSION_HOUR; -- Order by hour (0-23)
"""

# 2. Execute the query and load results into a Pandas DataFrame
adm_hour_distribution_df = client.query(query).to_dataframe()

# 3. Create the visualization using Matplotlib and Seaborn
if not adm_hour_distribution_df.empty:
    plt.figure(figsize=(12, 6))
    ax = sns.barplot(x='ADMISSION_HOUR', y='count',
                    data=adm_hour_distribution_df, palette='rocket')

    # Add labels and title
    plt.xlabel("Hour of Admission (0-23)")
    plt.ylabel("Number of Patients")
    plt.title("Distribution of Patients by Admission Hour")
    plt.xticks(range(0, 24))

    # Add count labels on top of bars
    for container in ax.containers:
        ax.bar_label(container, fmt='%.0f', label_type='edge', rotation=90,
                    padding=3)
```

```

# Display the plot
plt.tight_layout()
plt.show()
else:
    print("No data returned for admission hour distribution query.")

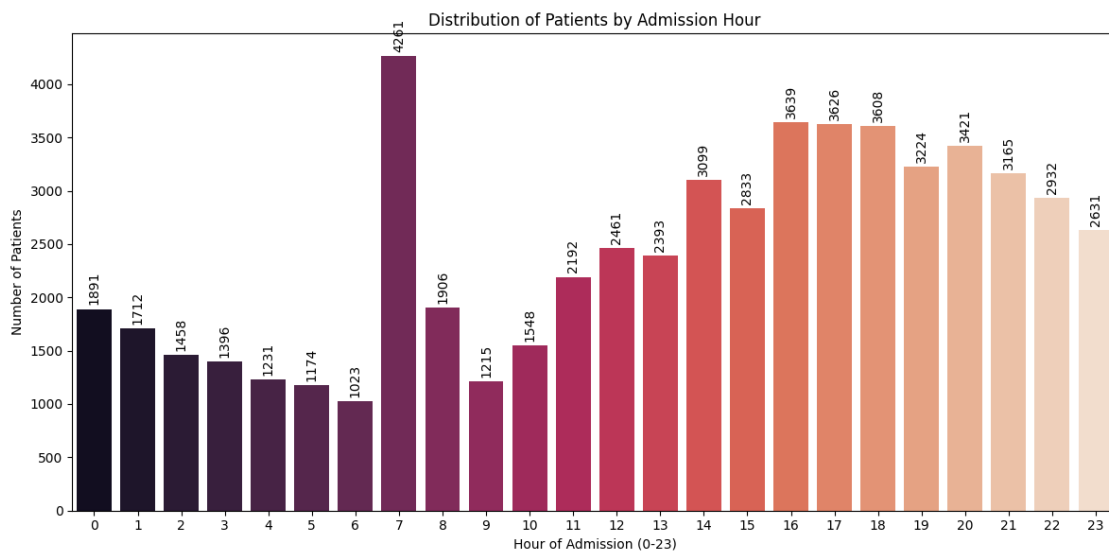
end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")

```

<ipython-input-48-beb0f04ab633>:24: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
ax = sns.barplot(x='ADMISSION_HOUR', y='count', data=adm_hour_distribution_df,
palette='rocket')
```



Query Execution Time: 2.33 seconds

5.2.5 Hour of Admission & Length of ICU Stay

```

[ ]: start_time = time.time()

# 1. Define the BigQuery SQL query to get Admission Hour and ICU Length of Stay
query = """
SELECT
    ADMISSION_HOUR,

```

```

        ICU_LOS
FROM
    `reliable-jet-452114-s2.table.junction_table`
WHERE
    ICU_LOS IS NOT NULL
    AND ADMISSION_HOUR IS NOT NULL
ORDER BY ADMISSION_HOUR;
"""

# 2. Execute the query and load results into a Pandas DataFrame
adm_hour_los_df = client.query(query).to_dataframe()

# 3. Create the visualization using Matplotlib and Seaborn
if not adm_hour_los_df.empty:
    plt.figure(figsize=(15, 7))

    # Create the box plot
    sns.boxplot(x='ADMISSION_HOUR', y='ICU_LOS', data=adm_hour_los_df,
        palette='rocket', showfliers=False)

    # Add labels and title
    plt.xlabel("Hour of Admission (0-23)")
    plt.ylabel("ICU Length of Stay (Days)")
    plt.title("Distribution of ICU Length of Stay by Admission Hour")
    plt.xticks(range(0, 24)) # Ensure all hours are labeled
    plt.ylim(bottom=0) # Ensure y-axis starts at 0

    # Display the plot
    plt.tight_layout()
    plt.show()

    # Print summary statistics
    print("\nSummary Statistics for ICU LOS by Admission Hour:")
    print(adm_hour_los_df.groupby('ADMISSION_HOUR')['ICU_LOS'].describe())

else:
    print("No data returned for admission hour vs. LOS query.")

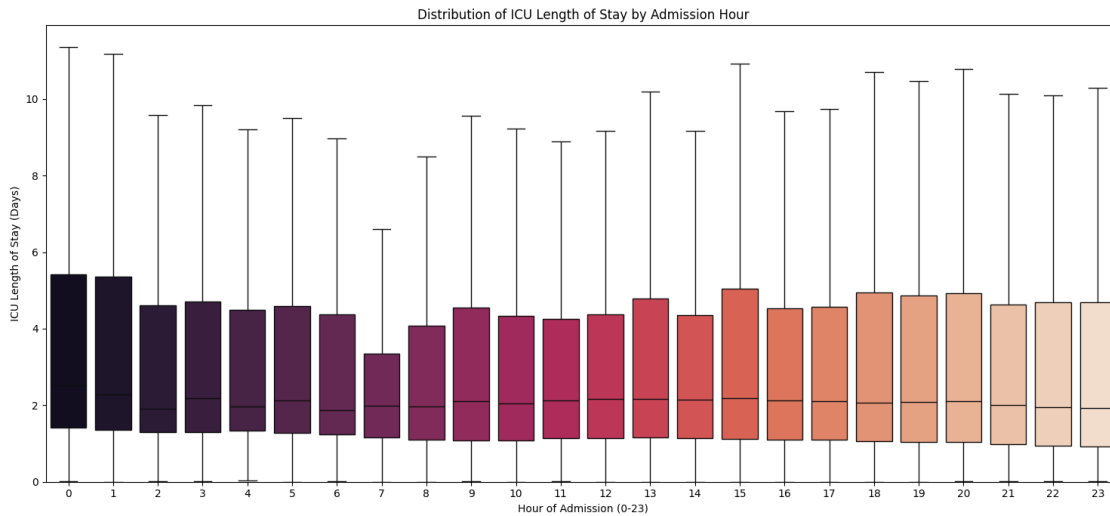
end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")

```

<ipython-input-49-277b40886152>:24: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='ADMISSION_HOUR', y='ICU_LOS', data=adm_hour_los_df,
palette='rocket', showfliers=False)
```



Summary Statistics for ICU LOS by Admission Hour:

	count	mean	std	min	25%	50% \
ADMISSION_HOUR						
0	1891.0	5.395593	9.814776	0.0160	1.420350	2.52960
1	1712.0	5.512634	10.433807	0.0003	1.346150	2.28265
2	1458.0	4.813653	9.620965	0.0206	1.298175	1.91085
3	1396.0	5.047518	10.259017	0.0047	1.290625	2.19230
4	1231.0	5.023875	10.586523	0.0245	1.328450	1.97100
5	1174.0	5.256344	10.591170	0.0002	1.275700	2.12510
6	1023.0	5.368494	12.690322	0.0203	1.244300	1.87670
7	4261.0	3.832038	7.981137	0.0036	1.162000	1.99170
8	1906.0	5.067366	12.268948	0.0014	1.096200	1.97030
9	1215.0	5.884291	11.624294	0.0218	1.084150	2.09620
10	1548.0	5.094745	9.783121	0.0006	1.073025	2.05040
11	2192.0	5.066186	10.892462	0.0169	1.133875	2.13450
12	2461.0	4.985880	9.973139	0.0004	1.139800	2.17170
13	2393.0	5.326986	10.293288	0.0029	1.164300	2.16970
14	3099.0	4.621340	8.345987	0.0001	1.132650	2.15390
15	2833.0	5.411682	10.503084	0.0008	1.111800	2.18470
16	3639.0	4.786286	8.466581	0.0042	1.106500	2.12790
17	3626.0	5.036128	9.968254	0.0004	1.091900	2.11340
18	3608.0	4.972881	8.580742	0.0012	1.052000	2.07585
19	3224.0	4.931366	9.064584	0.0025	1.043700	2.08265
20	3421.0	5.414299	10.158946	0.0055	1.030800	2.10340
21	3165.0	4.844583	9.644968	0.0079	0.972100	2.00920

22	2932.0	5.067636	9.973443	0.0094	0.943025	1.95310
23	2631.0	4.914707	9.962603	0.0047	0.919700	1.91860

	75%	max
ADMISSION_HOUR		
0	5.416900	134.6656
1	5.352625	136.6061
2	4.619400	138.4001
3	4.718525	145.5447
4	4.490700	167.5077
5	4.582175	135.5667
6	4.365950	140.5160
7	3.347300	169.4202
8	4.076275	148.3041
9	4.543550	97.1804
10	4.334750	105.2416
11	4.260750	151.4215
12	4.367800	153.9280
13	4.783100	125.3683
14	4.346350	113.1004
15	5.042200	129.1025
16	4.540650	122.9740
17	4.564125	144.7200
18	4.937575	108.0229
19	4.870700	148.7011
20	4.932100	131.9745
21	4.640800	173.0725
22	4.684750	171.6227
23	4.699800	150.7993

Query Execution Time: 3.52 seconds

5.2.6 Conclusion

There is clearly an outlier (7 AM) in the distribution of patients by Hour of Admission. Besides that, this distribution follows an increasing trend during the day [8-18] and a decreasing trend during the night [8-6].

There isn't an apparent connection between Hour of Admission and ICU Length of Stay, with the distribution of the mean values for the correlation being uniform.

5.2.7 Difference between Hour of Admission and ICU Hour of Admission Analysis

```
[ ]: # 1. Define the BigQuery SQL query to calculate and get the time difference
query = """
SELECT
    TIMESTAMP_DIFF(ICU_INTIME, ADMITTIME, HOUR) AS HOSP_TO_ICU_HOURS
FROM
    `reliable-jet-452114-s2.table.junction_table`
```

```

WHERE
    ADMITTIME IS NOT NULL
    AND ICU_INTIME IS NOT NULL
    AND TIMESTAMP_DIFF(ICU_INTIME, ADMITTIME, HOUR) >= 0 -- Exclude cases where ICU time is before admission time
    -- Add upper limit for visualization
    AND TIMESTAMP_DIFF(ICU_INTIME, ADMITTIME, HOUR) < 5
"""

# 2. Execute the query and load results into a Pandas DataFrame
hosp_to_icu_dist_df = client.query(query).to_dataframe()

# 3. Create the visualization using Matplotlib and Seaborn
if not hosp_to_icu_dist_df.empty:
    plt.figure(figsize=(12, 6))

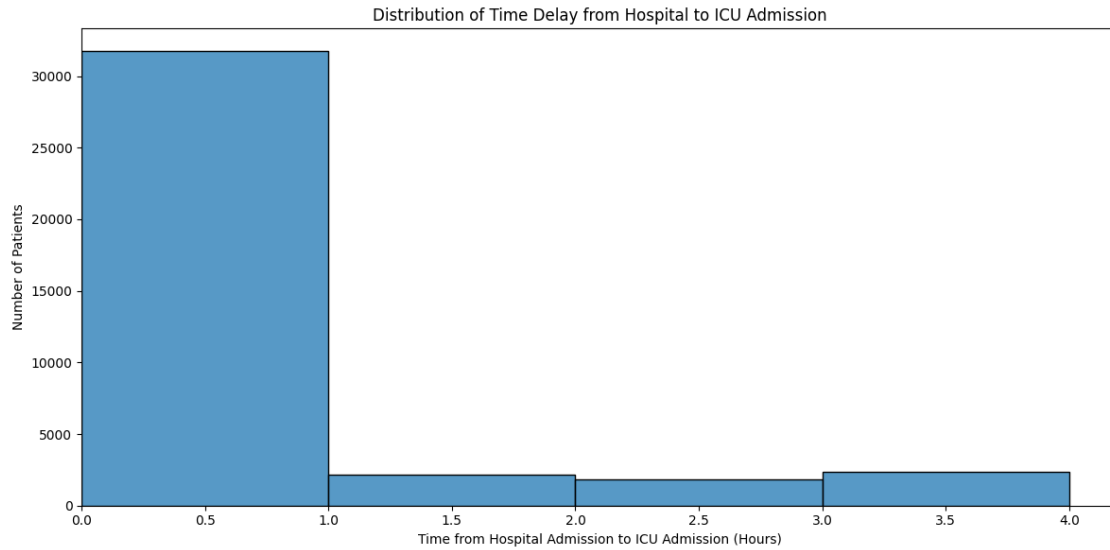
    # Create a histogram
    sns.histplot(data=hosp_to_icu_dist_df, x='HOSP_TO_ICU_HOURS', kde=False, bins=4)

    # Add labels and title
    plt.xlabel("Time from Hospital Admission to ICU Admission (Hours)")
    plt.ylabel("Number of Patients")
    plt.title("Distribution of Time Delay from Hospital to ICU Admission")
    plt.xlim(left=0) # Start x-axis at 0

    # Display the plot
    plt.tight_layout()
    plt.show()

    # Print descriptive statistics for the time difference
    print("\nHospital to ICU Admission Time Difference Statistics (Hours):")
    print(hosp_to_icu_dist_df['HOSP_TO_ICU_HOURS'].describe())
else:
    print("No data returned for Hospital to ICU admission time difference distribution query.")

```



Hospital to ICU Admission Time Difference Statistics (Hours):

```
count      38117.0
mean       0.364352
std        0.919322
min         0.0
25%         0.0
50%         0.0
75%         0.0
max         4.0
Name: HOSP_TO_ICU_HOURS, dtype: Float64
```

5.2.8 Conclusion

Most of the Admissions in ICU Stay occur within the first hour in the hospital, with the mean time for the difference between Hour of Admission and ICU Hour of Admission being ~22 minutes.

5.2.9 Diagnosis Analysis

```
[ ]: # 1. Define the BigQuery SQL query to get the counts of the top 20 primary ICD-9 codes
query = """
WITH TopCodes AS (
    SELECT
        PRIMARY_ICD9_CODE,
        COUNT(*) AS count
    FROM
        `reliable-jet-452114-s2.table.junction_table`
    WHERE
```

```

        PRIMARY_ICD9_CODE IS NOT NULL
    GROUP BY
        PRIMARY_ICD9_CODE
    ORDER BY
        count DESC
    LIMIT 20 -- Limit to the top 20 most frequent codes
)
SELECT
    jt.PRIMARY_ICD9_CODE,
    tc.count,
    ANY_VALUE(jt.PRIMARY_ICD9_TITLE) AS PRIMARY_ICD9_TITLE -- Get a
    ↳representative title
FROM
    `reliable-jet-452114-s2.table.junction_table` jt
JOIN
    TopCodes tc ON jt.PRIMARY_ICD9_CODE = tc.PRIMARY_ICD9_CODE
GROUP BY
    jt.PRIMARY_ICD9_CODE, tc.count
ORDER BY
    tc.count DESC;
"""

# 2. Execute the query and load results into a Pandas DataFrame
icd9_distribution_df = client.query(query).to_dataframe()

# 3. Create the visualization using Matplotlib and Seaborn
if not icd9_distribution_df.empty:
    plt.figure(figsize=(10, 10))
    # Combine code and title for better y-axis labels
    icd9_distribution_df['CODE_TITLE'] =
    ↳icd9_distribution_df['PRIMARY_ICD9_CODE'] + ': ' +
    ↳icd9_distribution_df['PRIMARY_ICD9_TITLE'].fillna('N/A')
    sns.barplot(x='count', y='CODE_TITLE', data=icd9_distribution_df,
    ↳palette='viridis', orient='h')

    # Add labels and title
    plt.xlabel("Number of Patients")
    plt.ylabel("Primary ICD-9 Code")
    plt.title("Distribution of Top 20 Primary ICD-9 Codes at Admission")

    # Display the plot
    plt.tight_layout()
    plt.show()
else:
    print("No data returned for primary ICD-9 code distribution query.")


```



```
<ipython-input-14-f3c32900dad3>:39: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x='count', y='CODE_TITLE', data=icd9_distribution_df,
palette='viridis', orient='h')
```



5.2.10 Diagnosis & Length of ICU Stay

```
[ ]: # 1. Define the BigQuery SQL query to get ICU Length of Stay for the top 20
↳ICD-9 Codes
# We reuse the list of top codes identified above.
top_codes_list = icd9_distribution_df['PRIMARY_ICD9_CODE'].tolist()

# Check if the list is not empty before the query
if top_codes_list:
    # Format the list of codes for the SQL IN clause
    formatted_top_codes = ", ".join([f'"{code}"' for code in top_codes_list])

    query = f"""
    SELECT
        PRIMARY_ICD9_CODE,
        ICU_LOS
    FROM
        `reliable-jet-452114-s2.table.junction_table`
    WHERE
        ICU_LOS IS NOT NULL
        AND PRIMARY_ICD9_CODE IN ({formatted_top_codes}) -- Filter for top 20
↳codes
    ORDER BY PRIMARY_ICD9_CODE;
    """

# 2. Execute the query and load results into a Pandas DataFrame
icd9_los_df = client.query(query).to_dataframe()

# 3. Create the visualization using Matplotlib and Seaborn
if not icd9_los_df.empty:
    plt.figure(figsize=(15, 8))

    # Create the box plot, ordering by the frequency from the first query
    sns.boxplot(x='PRIMARY_ICD9_CODE', y='ICU_LOS', data=icd9_los_df,
↳palette='viridis', order=top_codes_list, showfliers=False) # Vertical, hide
↳outliers

    # Add labels and title
    plt.xlabel("Primary ICD-9 Code")
    plt.ylabel("ICU Length of Stay (Days)")
    plt.title("Distribution of ICU Length of Stay by Top 20 Primary ICD-9
↳Codes")
    plt.xticks(rotation=90)
    plt.ylim(bottom=0) # Ensure y-axis starts at 0

    # Display the plot
    plt.tight_layout()
```

```

plt.show()

# Print summary statistics
print("\nSummary Statistics for ICU LOS by Top 20 Primary ICD-9 Codes:")
print(icd9_los_df.groupby('PRIMARY_ICD9_CODE')['ICU_LOS'].describe())

else:
    print("No data returned for top ICD-9 codes vs. LOS query.")
else:
    print("Cannot proceed with LOS analysis as no top ICD-9 codes were_
    ↪identified.")

```

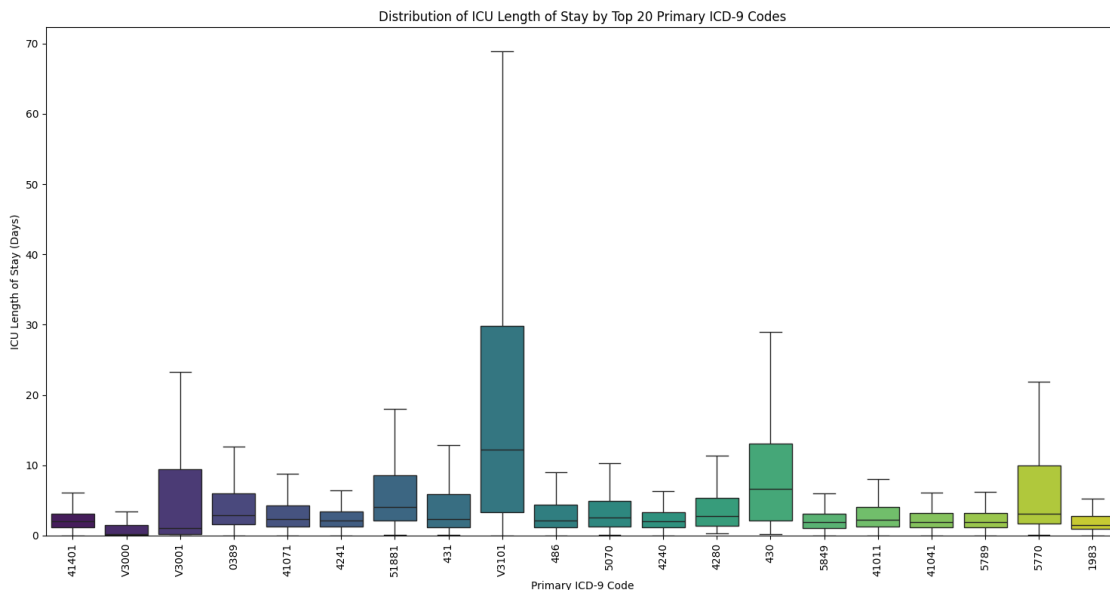
<ipython-input-15-24c6b6b5d6e2>:30: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```

sns.boxplot(x='PRIMARY_ICD9_CODE', y='ICU_LOS', data=icd9_los_df,
palette='viridis', order=top_codes_list, showfliers=False) # Vertical, hide
outliers

```



Summary Statistics for ICU LOS by Top 20 Primary ICD-9 Codes:

	count	mean	std	min	25%	50% \
PRIMARY_ICD9_CODE						
0389	1988.0	5.329782	7.084884	0.0079	1.614200	2.87470
1983	343.0	2.509505	3.203271	0.0079	1.005550	1.51880

41011	473.0	3.895759	5.718157	0.0435	1.314000	2.20950
41041	471.0	3.449264	4.951962	0.0033	1.230300	1.94030
41071	1803.0	4.026161	5.369919	0.0016	1.288900	2.35840
41401	3617.0	2.876750	3.885769	0.0048	1.154600	1.99070
4240	592.0	3.466497	5.780785	0.0077	1.159450	2.02135
4241	1168.0	3.624113	5.788327	0.0014	1.233175	2.14675
4280	562.0	4.949970	6.628354	0.3021	1.387300	2.73745
430	542.0	9.060081	8.904613	0.1733	2.170775	6.62515
431	1031.0	4.795295	5.835963	0.1118	1.213400	2.30710
486	684.0	3.930527	4.905842	0.0280	1.196175	2.15545
5070	596.0	4.541065	5.972763	0.1101	1.334350	2.55865
51881	1124.0	6.957914	7.858505	0.1524	2.096675	4.12695
5770	399.0	8.341084	12.501480	0.0566	1.757150	3.08880
5789	415.0	2.812222	3.300042	0.0214	1.149050	1.96770
5849	523.0	3.024175	4.131413	0.0012	1.110750	1.90300
V3000	3534.0	4.589990	14.277310	0.0037	0.104725	0.20955
V3001	2792.0	10.340239	21.146993	0.0008	0.162275	1.02360
V3101	1008.0	21.719032	26.430318	0.0098	3.314175	12.19845

	75%	max
PRIMARY_ICD9_CODE		
0389	6.028650	91.5726
1983	2.743500	30.7172
41011	4.038900	76.9211
41041	3.198900	40.9936
41071	4.309700	53.0333
41401	3.134600	68.2052
4240	3.360900	75.9919
4241	3.437825	98.6446
4280	5.372050	68.9517
430	13.097175	54.5119
431	5.869350	51.7955
486	4.354575	41.5576
5070	4.938750	59.4319
51881	8.548500	71.0056
5770	9.929350	101.7390
5789	3.202750	36.3250
5849	3.082300	43.2606
V3000	1.457850	167.5077
V3001	9.443825	171.6227
V3101	29.786825	142.3605

5.2.11 Conclusion

There is a very unsymmetric distribution of Diagnosis. In the top-20 most common diagnosis, with the most common being 41401 (Coronary atherosclerosis of native vessel). The correlation between some Diagnosis and ICU Length of Stay is evident, with V3101 (Twin mate, liveborn, delivered in hospital, with cesarean section) having a mean value very high

compared to other types of Diagnosis. There are also diagnosis who have an average of ICU Length of Stay very low in comparison, like V3000 (Single liveborn, delivered in hospital, without cesarean section).

5.2.12 First Care Unit Analysis

```
[ ]: # 1. Define the BigQuery SQL query to get first care unit counts
query = """
SELECT
    FIRST_CAREUNIT,
    COUNT(*) AS count
FROM
    `reliable-jet-452114-s2.table.junction_table`
WHERE
    FIRST_CAREUNIT IS NOT NULL
GROUP BY
    FIRST_CAREUNIT
ORDER BY
    count DESC; -- Order by count to see most common units first
"""

# 2. Execute the query and load results into a Pandas DataFrame
careunit_distribution_df = client.query(query).to_dataframe()

# 3. Create the visualization using Matplotlib and Seaborn
if not careunit_distribution_df.empty:
    plt.figure(figsize=(10, 6))
    ax = sns.barplot(x='FIRST_CAREUNIT', y='count',
                    data=careunit_distribution_df, palette='Spectral')

    # Add labels and title
    plt.xlabel("First Care Unit")
    plt.ylabel("Number of Patients")
    plt.title("Distribution of Patients by First Care Unit")
    plt.xticks(rotation=45, ha='right')


    # Add count labels on top of bars
    for container in ax.containers:
        ax.bar_label(container)

    # Display the plot
    plt.tight_layout()
    plt.show()
else:
    print("No data returned for first care unit distribution query.")
```

<ipython-input-16-2e951009bbdf>:22: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
ax = sns.barplot(x='FIRST_CAREUNIT', y='count', data=careunit_distribution_df,
palette='Spectral')
```



5.2.13 First Care of Unit & Length of ICU Stay

```
[ ]: # 1. Define the BigQuery SQL query to get First Care Unit and ICU Length of Stay
query = """
SELECT
    FIRST_CAREUNIT,
    ICU_LOS
FROM
    `reliable-jet-452114-s2.table.junction_table`
WHERE
    ICU_LOS IS NOT NULL
    AND FIRST_CAREUNIT IS NOT NULL
ORDER BY FIRST_CAREUNIT;
"""

# 2. Execute the query and load results into a Pandas DataFrame
careunit_los_df = client.query(query).to_dataframe()
```

```

# 3. Create the visualization using Matplotlib and Seaborn
if not careunit_los_df.empty:
    plt.figure(figsize=(12, 7))

    # Create the box plot
    sns.boxplot(x='FIRST_CAREUNIT', y='ICU_LOS', data=careunit_los_df,
palette='Spectral', showfliers=False)

    # Add labels and title
    plt.xlabel("First Care Unit")
    plt.ylabel("ICU Length of Stay (Days)")
    plt.title("Distribution of ICU Length of Stay by First Care Unit")
    plt.xticks(rotation=45, ha='right')
    plt.ylim(bottom=0) # Ensure y-axis starts at 0

    # Display the plot
    plt.tight_layout()
    plt.show()

    # Print summary statistics
    print("\nSummary Statistics for ICU LOS by First Care Unit:")
    print(careunit_los_df.groupby('FIRST_CAREUNIT')['ICU_LOS'].describe())

else:
    print("No data returned for first care unit vs. LOS query.")

```

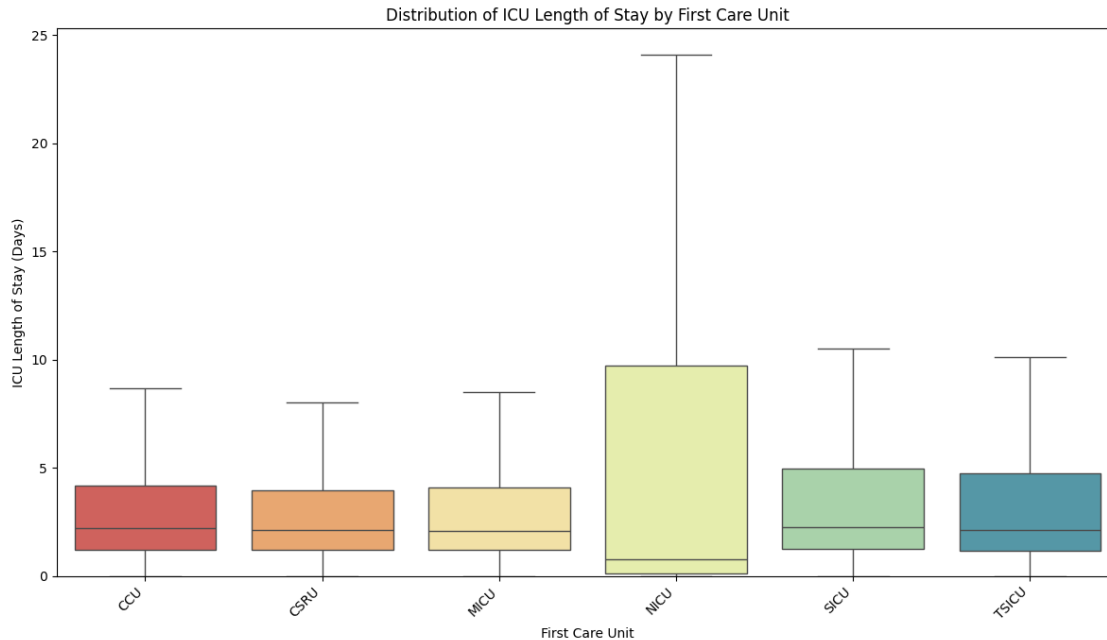
<ipython-input-17-c08169d1c5ed>:22: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```

sns.boxplot(x='FIRST_CAREUNIT', y='ICU_LOS', data=careunit_los_df,
palette='Spectral', showfliers=False) # Vertical, hide outliers

```



Summary Statistics for ICU LOS by First Care Unit:

	count	mean	std	min	25%	50% \
FIRST_CAREUNIT						
CCU	7132.0	3.953257	5.619626	0.0012	1.213125	2.19915
CSRU	9087.0	3.888229	6.097731	0.0001	1.213850	2.14770
MICU	19340.0	4.042175	5.890149	0.0004	1.189600	2.09700
NICU	8075.0	10.004395	20.644582	0.0008	0.140500	0.78470
SICU	8374.0	4.749462	6.997957	0.0003	1.258775	2.26470
TSICU	6031.0	4.519062	6.760158	0.0016	1.174100	2.12800

	75%	max
FIRST_CAREUNIT		
CCU	4.200700	100.1225
CSRU	3.943200	153.9280
MICU	4.109775	116.8327
NICU	9.723350	171.6227
SICU	4.964775	101.7390
TSICU	4.755900	173.0725

5.2.14 Conclusion

There is a relative uniform distribution in the number of patients by First Care Unit with the exception being MICU (Medical Intensive Care Unit), that has a count of patients that is more than double of each other type of First Care Unit.

The correlation between Type of Care Unit and ICU Length of Stay is also relatively uniform

with an exception, NICU (Neonatal Intensive Care Unit).

5.2.15 Death Analysis

```
[ ]: # 1. Define the BigQuery SQL query to get expire flag counts
query = """
SELECT
    EXPIRE_FLAG,
    COUNT(*) AS count
FROM
    `reliable-jet-452114-s2.table.junction_table`
WHERE
    EXPIRE_FLAG IS NOT NULL
GROUP BY
    EXPIRE_FLAG
ORDER BY
    EXPIRE_FLAG; -- Order by flag value (0, 1)
"""

# 2. Execute the query and load results into a Pandas DataFrame
expire_flag_distribution_df = client.query(query).to_dataframe()

# 3. Create the visualization using Matplotlib and Seaborn
if not expire_flag_distribution_df.empty:
    plt.figure(figsize=(6, 4))
    # Map 0/1 to more descriptive labels for the plot
    expire_flag_distribution_df['Outcome'] = \
    ↪expire_flag_distribution_df['EXPIRE_FLAG'].map({0: 'Survived', 1: 'Expired'})
    ax = sns.barplot(x='Outcome', y='count', data=expire_flag_distribution_df, \
    ↪palette='coolwarm')

    # Add labels and title
    plt.xlabel("Patient Outcome During Hospital Stay")
    plt.ylabel("Number of Patients")
    plt.title("Distribution of Patient Outcome")


    # Add count labels on top of bars
    for container in ax.containers:
        ax.bar_label(container)

    # Display the plot
    plt.tight_layout()
    plt.show()
else:
    print("No data returned for expire flag distribution query.")
```

<ipython-input-19-8827588862fa>:24: FutureWarning:

Passing ``palette`` without assigning ``hue`` is deprecated and will be removed in v0.14.0. Assign the ``x`` variable to ``hue`` and set ``legend=False`` for the same effect.

```
ax = sns.barplot(x='Outcome', y='count', data=expire_flag_distribution_df,
palette='coolwarm')
```



5.2.16 Death & Length of ICU Stay

```
[ ]: # 1. Define the BigQuery SQL query to get Expire Flag and ICU Length of Stay
query = """
SELECT
    EXPIRE_FLAG,
    ICU_LOS
FROM
    `reliable-jet-452114-s2.table.junction_table`
WHERE
    ICU_LOS IS NOT NULL
    AND EXPIRE_FLAG IS NOT NULL
ORDER BY EXPIRE_FLAG;
"""

# 2. Execute the query and load results into a Pandas DataFrame
```

```

expire_flag_los_df = client.query(query).to_dataframe()

# 3. Create the visualization using Matplotlib and Seaborn
if not expire_flag_los_df.empty:
    plt.figure(figsize=(8, 6))

    # Map 0/1 to more descriptive labels for the plot
    expire_flag_los_df['Outcome'] = expire_flag_los_df['EXPIRE_FLAG'].map({0: 'Survived', 1: 'Expired'})

    # Create the box plot
    sns.boxplot(x='Outcome', y='ICU_LOS', data=expire_flag_los_df,
        palette='coolwarm', showfliers=True)

    # Add labels and title
    plt.xlabel("Patient Outcome During Hospital Stay")
    plt.ylabel("ICU Length of Stay (Days)")
    plt.title("Distribution of ICU Length of Stay by Patient Outcome")
    plt.ylim(bottom=0) # Ensure y-axis starts at 0

    # Display the plot
    plt.tight_layout()
    plt.show()

    # Print summary statistics
    print("\nSummary Statistics for ICU LOS by Patient Outcome:")
    print(expire_flag_los_df.groupby('Outcome')['ICU_LOS'].describe())

else:
    print("No data returned for expire flag vs. LOS query.")

```


<ipython-input-19-ae139f396912>:25: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```

sns.boxplot(x='Outcome', y='ICU_LOS', data=expire_flag_los_df,
palette='coolwarm', showfliers=True) # Show outliers might be relevant here

```



Summary Statistics for ICU LOS by Patient Outcome:

	count	mean	std	min	25%	50%	75%	\
Outcome								
Expired	21736.0	5.139688	7.542676	0.0001	1.375575	2.6904	5.521625	
Survived	36303.0	4.897620	10.964588	0.0003	1.021650	1.8874	3.945950	

	max
Outcome	
Expired	173.0725
Survived	171.6227

5.2.17 Conclusion

The number of patients that **Survived** is much higher than the number of patients that **Expired**. The mean of the correlation between **Death** and **ICU Length of Stay** is similar for **Survived** and **Expired** but the standard deviation of **Survived** is significantly higher.

5.2.18 Length of ICU Stay Analysis

```
[ ]: # 1. Define the BigQuery SQL query to get ICU Length of Stay values
query = """
SELECT
    ICU_LOS
FROM
    `reliable-jet-452114-s2.table.junction_table`
WHERE
    ICU_LOS IS NOT NULL
"""

# 2. Execute the query and load results into a Pandas DataFrame
icu_los_distribution_df = client.query(query).to_dataframe()

# 3. Create the visualization using Matplotlib and Seaborn
if not icu_los_distribution_df.empty:
    plt.figure(figsize=(10, 6))

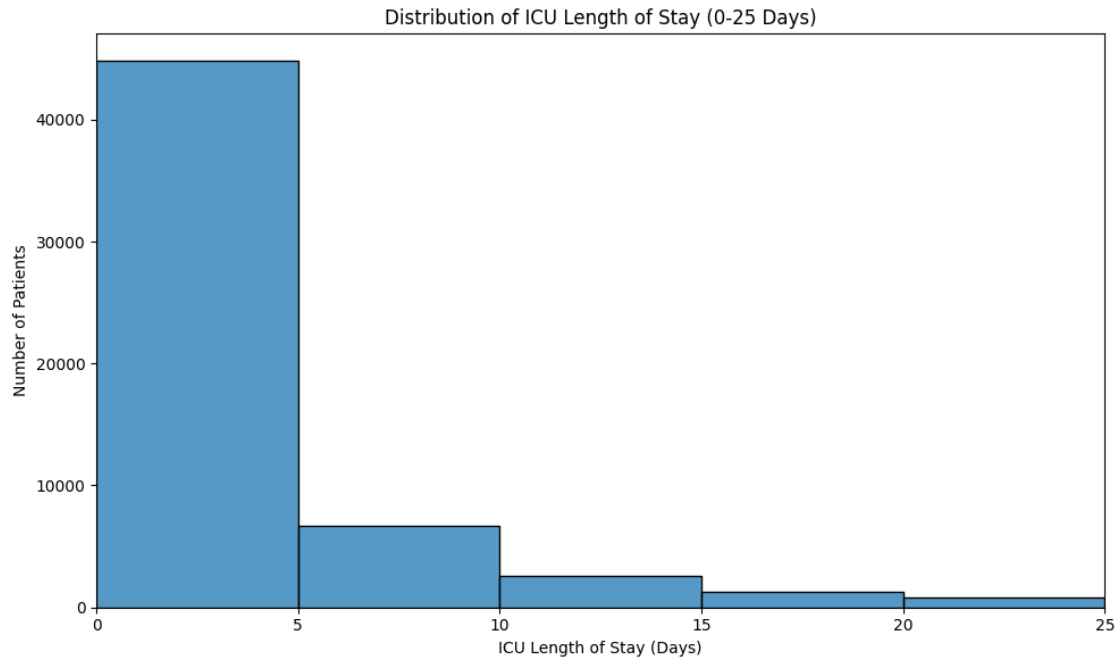
    # Filter data for plotting
    plot_data = icu_los_distribution_df[(icu_los_distribution_df['ICU_LOS'] >= 0
    ↪ 0) & (icu_los_distribution_df['ICU_LOS'] <= 25)]

    # Create a histogram
    sns.histplot(data=plot_data, x='ICU_LOS', kde=False, bins=5)

    # Add labels and title
    plt.xlabel("ICU Length of Stay (Days)")
    plt.ylabel("Number of Patients")
    plt.title("Distribution of ICU Length of Stay (0-25 Days)")
    plt.xlim(left=0, right=25) # Keep x-axis limit from 0 to 25

    # Display the plot
    plt.tight_layout()
    plt.show()

    # Print descriptive statistics
    print("\nICU Length of Stay Statistics:")
    print(icu_los_distribution_df['ICU_LOS'].describe())
else:
    print("No data returned for ICU LOS distribution query.")
```



ICU Length of Stay Statistics:

```
count    58039.000000
mean      4.988276
std       9.824303
min       0.000100
25%       1.103350
50%       2.091800
75%       4.550800
max      173.072500
Name: ICU_LOS, dtype: float64
```

5.2.19 Conclusion

By analysing the distribution of ICU Length of Stay we can see clearly that most of patients only stay for 5 days or less, with a decreasing trend of the number of patients as the number of days increases.

5.3 Correlations

To finish this analysis we will attempt to find connections between the following set of features:

- Type of Admission & Hour of Admission
- Type of Admission and Diagnosis
- Type of Admission and Death
- Hour of Admission & Diagnosis
- Diagnosis & First Care Unit

- Diagnosis and Death
- First Care Unit & Death

5.3.1 Type of Admission & Hour of Admission

```
[ ]: # 1. Define the BigQuery SQL query to get counts for each combination
query = """
SELECT
    ADMISSION_TYPE,
    ADMISSION_HOUR,
    COUNT(*) AS count
FROM
    `reliable-jet-452114-s2.table.junction_table`
WHERE
    ADMISSION_TYPE IS NOT NULL
    AND ADMISSION_HOUR IS NOT NULL
GROUP BY
    ADMISSION_TYPE,
    ADMISSION_HOUR
ORDER BY
    ADMISSION_TYPE,
    ADMISSION_HOUR;
"""

# 2. Execute the query and load results into a Pandas DataFrame
adm_type_hour_df = client.query(query).to_dataframe()

adm_type_hour_df['count'] = pd.to_numeric(adm_type_hour_df['count'],
    ↪errors='coerce')
# Ensure ADMISSION_HOUR is treated as a number
adm_type_hour_df['ADMISSION_HOUR'] = pd.
    ↪to_numeric(adm_type_hour_df['ADMISSION_HOUR'], errors='coerce')
adm_type_hour_df.dropna(subset=['count', 'ADMISSION_HOUR'], inplace=True) #
    ↪Drop rows where conversion failed

# 3. Create the visualization using Matplotlib and Seaborn
if not adm_type_hour_df.empty:
    try:
        # Pivot the data to create a matrix suitable for a heatmap
        heatmap_data = adm_type_hour_df.pivot(index='ADMISSION_TYPE',
    ↪columns='ADMISSION_HOUR', values='count').fillna(0) # Fill missing combos
    ↪with 0
        heatmap_data = heatmap_data.astype(float)
```

```

plt.figure(figsize=(15, 6))

# Create the heatmap
sns.heatmap(heatmap_data, annot=False, fmt=".0f", linewidths=.5,
↪ cmap="YlGnBu")

# Add labels and title
plt.xlabel("Hour of Admission (0-23)")
plt.ylabel("Admission Type")
plt.title("Heatmap of Admission Counts by Type and Hour")
plt.yticks(rotation=0) # Ensure y-axis labels are horizontal


# Display the plot
plt.tight_layout()
plt.show()

# Display the pivoted data table
print("\nPivoted Data for Heatmap:")
print(heatmap_data)
print("\nPivoted Data Types:")
print(heatmap_data.dtypes)

except Exception as e:
    print(f"An error occurred during plotting: {e}")
    print("\nOriginal DataFrame dtypes before pivot:")
    print(adm_type_hour_df.dtypes)
    if 'heatmap_data' in locals():
        print("\nPivoted DataFrame dtypes before plotting:")
        print(heatmap_data.dtypes)

else:
    print("No data returned from the query to plot.")

```

Pivoted Data for Heatmap:

ADMISSION_HOUR	0	1	2	3	4	5	6	7	\
ADMISSION_TYPE									
ELECTIVE	70.0	20.0	12.0	9.0	2.0	4.0	6.0	3457.0	
EMERGENCY	1474.0	1348.0	1169.0	1095.0	989.0	917.0	778.0	552.0	
NEWBORN	324.0	330.0	265.0	282.0	236.0	245.0	226.0	229.0	
URGENT	23.0	14.0	12.0	10.0	4.0	8.0	13.0	23.0	

ADMISSION_HOUR	8	9	...	14	15	16	17	18	\
ADMISSION_TYPE			...						
ELECTIVE	986.0	222.0	...	461.0	84.0	87.0	51.0	42.0	
EMERGENCY	554.0	643.0	...	2184.0	2292.0	3111.0	3111.0	3077.0	
NEWBORN	347.0	320.0	...	356.0	356.0	337.0	353.0	378.0	
URGENT	19.0	30.0	...	98.0	101.0	104.0	111.0	111.0	

ADMISSION_HOUR	19	20	21	22	23
ADMISSION_TYPE					
ELECTIVE	21.0	15.0	11.0	1.0	8.0
EMERGENCY	2733.0	2907.0	2722.0	2500.0	2238.0
NEWBORN	347.0	403.0	348.0	378.0	349.0
URGENT	123.0	96.0	84.0	53.0	36.0

[4 rows x 24 columns]

Pivoted Data Types:

ADMISSION_HOUR	
0	float64
1	float64
2	float64
3	float64

```

4    float64
5    float64
6    float64
7    float64
8    float64
9    float64
10   float64
11   float64
12   float64
13   float64
14   float64
15   float64
16   float64
17   float64
18   float64
19   float64
20   float64
21   float64
22   float64
23   float64
dtype: object

```

5.3.2 Conclusion

By visualizing the heatmap, we can see that the **EMERGENCY** Type of Admission is very correlated with Hour of Admission, being much more common in the evening/night 16PM - 23PM. The Type of Admission **ELECTIVE** is also correlated with Hour of Admission, being very common at 7AM.

5.3.3 Type of Admission & Diagnosis

```

[ ]: # 1. Define the BigQuery SQL query to get counts for each combination
#     First, find the top 20 most frequent ICD-9 codes, then get the counts
#     for combinations of ADMISSION_TYPE and those top codes.
query = """
WITH TopCodes AS (
    -- Subquery to find the top 20 most frequent primary ICD-9 codes
    SELECT
        PRIMARY_ICD9_CODE
    FROM
        `reliable-jet-452114-s2.table.junction_table`
    WHERE
        PRIMARY_ICD9_CODE IS NOT NULL
    GROUP BY
        PRIMARY_ICD9_CODE
    ORDER BY
        COUNT(*) DESC
    LIMIT 20
)

```

```

-- Main query to count combinations for top codes and admission types
SELECT
    jt.ADMISSION_TYPE,
    jt.PRIMARY_ICD9_CODE,
    COUNT(*) AS count
FROM
    `reliable-jet-452114-s2.table.junction_table` jt
JOIN
    TopCodes tc ON jt.PRIMARY_ICD9_CODE = tc.PRIMARY_ICD9_CODE -- Join to
    ↪filter for top codes
WHERE
    jt.ADMISSION_TYPE IS NOT NULL
GROUP BY
    jt.ADMISSION_TYPE,
    jt.PRIMARY_ICD9_CODE
ORDER BY
    jt.ADMISSION_TYPE,
    jt.PRIMARY_ICD9_CODE;
"""

# 2. Execute the query and load results into a Pandas DataFrame
adm_type_icd9_df = client.query(query).to_dataframe()

# 3. Create the visualization using Matplotlib and Seaborn
if not adm_type_icd9_df.empty:
    heatmap_data = adm_type_icd9_df.pivot(index='ADMISSION_TYPE',
    ↪columns='PRIMARY_ICD9_CODE', values='count').fillna(0) # Fill missing combos
    ↪with 0

    plt.figure(figsize=(18, 7))

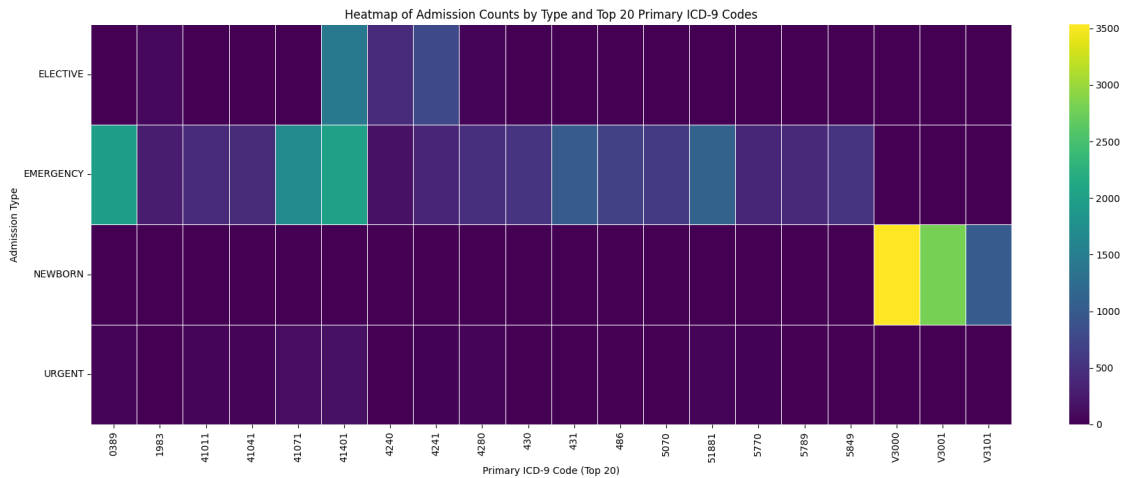
    # Create the heatmap
    heatmap_data = heatmap_data.astype(float)
    sns.heatmap(heatmap_data, annot=False, fmt=".0f", linewidths=.5,
    ↪cmap="viridis")

    plt.xlabel("Primary ICD-9 Code (Top 20)")
    plt.ylabel("Admission Type")
    plt.title("Heatmap of Admission Counts by Type and Top 20 Primary ICD-9
    ↪Codes")
    plt.xticks(rotation=90) # Rotate ICD codes for readability
    plt.yticks(rotation=0) # Ensure y-axis labels are horizontal

    # Display the plot
    plt.tight_layout()
    plt.show()

```

```
else:
    print("No data returned from the query to plot.")
```



5.3.4 Conclusion

There isn't a clear connection between Type of Admission and Diagnosis with the exceptions being V300 (Single liveborn, delivered in hospital, without cesarean section) and NEWBORN, that have a higher correlation due to their obvious connection, and some diseases like 41071(Acute myocardial infarction, subendocardial infarction, initial episode of care) with EMERGENCY, that have a low correlation.

5.3.5 Type of Admission & Death

```
[ ]: # 1. Define the BigQuery SQL query to get counts for each combination
query = """
SELECT
    ADMISSION_TYPE,
    EXPIRE_FLAG,
    COUNT(*) AS count
FROM
    `reliable-jet-452114-s2.table.junction_table`
WHERE
    ADMISSION_TYPE IS NOT NULL
    AND EXPIRE_FLAG IS NOT NULL
GROUP BY
    ADMISSION_TYPE,
    EXPIRE_FLAG
ORDER BY
    ADMISSION_TYPE,
    EXPIRE_FLAG;
"""
```

```

# 2. Execute the query and load results into a Pandas DataFrame
adm_type_expire_df = client.query(query).to_dataframe()

# 3. Create the visualization using Matplotlib and Seaborn
if not adm_type_expire_df.empty:
    # Map EXPIRE_FLAG to readable labels
    adm_type_expire_df['Outcome'] = adm_type_expire_df['EXPIRE_FLAG'].map({0: 'Survived', 1: 'Expired'})

    heatmap_data = adm_type_expire_df.pivot(index='ADMISSION_TYPE',
    ↪ columns='Outcome', values='count').fillna(0) # Fill missing combos with 0

    plt.figure(figsize=(8, 6))


    # Create the heatmap
    heatmap_data = heatmap_data.astype(float)
    sns.heatmap(heatmap_data, annot=True, fmt=".0f", linewidths=.5,
    ↪ cmap="coolwarm", cbar_kws={'label': 'Number of Patients'})

    # Add labels and title
    plt.xlabel("Patient Outcome During Hospital Stay")
    plt.ylabel("Admission Type")
    plt.title("Heatmap of Patient Outcomes by Admission Type")
    plt.yticks(rotation=0) # Ensure y-axis labels are horizontal

    # Display the plot
    plt.tight_layout()
    plt.show()

    # Display the pivoted data table
    print("\nPivoted Data for Heatmap:")
    print(heatmap_data)
else:
    print("No data returned from the query to plot.")

```



Pivoted Data for Heatmap:

Outcome	Expired	Survived
ADMISSION_TYPE		
ELECTIVE	1990.0	5236.0
EMERGENCY	18992.0	22609.0
NEWBORN	65.0	7801.0
URGENT	689.0	657.0

5.3.6 Conclusion

The clearest conclusion from this HeatMap is that patients who have **EMERGENCY** as their Type of Admission tend to **EXPIRE** much more frequently.

5.3.7 Admission Hour & Diagnosis

```
[ ]: # 1. Define the BigQuery SQL query to get counts for each combination
# First, find the top 20 most frequent ICD-9 codes, then get the counts
# for combinations of ADMISSION_HOUR and those top codes.
query = """
WITH TopCodes AS (
```

```

-- Subquery to find the top 20 most frequent primary ICD-9 codes
SELECT
    PRIMARY_ICD9_CODE
FROM
    `reliable-jet-452114-s2.table.junction_table`
WHERE
    PRIMARY_ICD9_CODE IS NOT NULL
GROUP BY
    PRIMARY_ICD9_CODE
ORDER BY
    COUNT(*) DESC
LIMIT 20
)
-- Main query to count combinations for top codes and admission hour
SELECT
    jt.ADMISSION_HOUR,
    jt.PRIMARY_ICD9_CODE,
    COUNT(*) AS count
FROM
    `reliable-jet-452114-s2.table.junction_table` jt
JOIN
    TopCodes tc ON jt.PRIMARY_ICD9_CODE = tc.PRIMARY_ICD9_CODE -- Join to
    ↪filter for top codes
WHERE
    jt.ADMISSION_HOUR IS NOT NULL
GROUP BY
    jt.ADMISSION_HOUR,
    jt.PRIMARY_ICD9_CODE
ORDER BY
    jt.ADMISSION_HOUR,
    jt.PRIMARY_ICD9_CODE;
"""

# 2. Execute the query and load results into a Pandas DataFrame
adm_hour_icd9_df = client.query(query).to_dataframe()

# 3. Create the visualization using Matplotlib and Seaborn
if not adm_hour_icd9_df.empty:
    heatmap_data = adm_hour_icd9_df.pivot(index='ADMISSION_HOUR',
    ↪columns='PRIMARY_ICD9_CODE', values='count').fillna(0) # Fill missing combos
    ↪with 0

    plt.figure(figsize=(18, 8))

    # Create the heatmap
    heatmap_data = heatmap_data.astype(float)

```

```

sns.heatmap(heatmap_data, annot=False, fmt=".0f", linewidths=.5,
cmap="rocket", cbar_kws={'label': 'Number of Patients'})

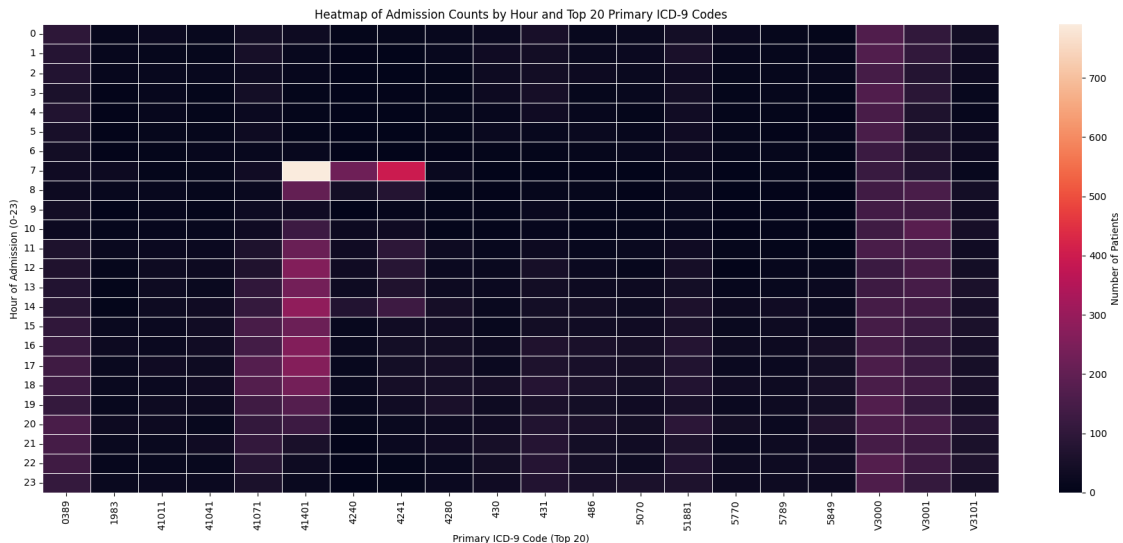
# Add labels and title
plt.xlabel("Primary ICD-9 Code (Top 20)")
plt.ylabel("Hour of Admission (0-23)")
plt.title("Heatmap of Admission Counts by Hour and Top 20 Primary ICD-9
Codes")
plt.xticks(rotation=90)
plt.yticks(rotation=0) # Ensure y-axis labels are horizontal

# Display the plot
plt.tight_layout()
plt.show()

# Display the pivoted data table
print("\nPivoted Data for Heatmap:")
print(heatmap_data)

else:
    print("No data returned from the query to plot.")

```



Pivoted Data for Heatmap:

	PRIMARY_ICD9_CODE	0389	1983	41011	41041	41071	41401	4240	4241	\
ADMISSION_HOUR										
0		95.0	13.0	22.0	17.0	42.0	29.0	9.0	10.0	
1		78.0	13.0	11.0	10.0	48.0	17.0	4.0	1.0	
2		74.0	14.0	16.0	16.0	29.0	14.0	3.0	2.0	

3	56.0	5.0	13.0	8.0	42.0	7.0	1.0	5.0
4	69.0	5.0	14.0	9.0	23.0	9.0	1.0	1.0
5	51.0	4.0	10.0	10.0	28.0	7.0	2.0	2.0
6	38.0	6.0	8.0	14.0	18.0	17.0	0.0	2.0
7	33.0	26.0	8.0	10.0	35.0	790.0	224.0	396.0
8	26.0	18.0	17.0	10.0	20.0	202.0	43.0	77.0
9	38.0	2.0	15.0	9.0	28.0	31.0	3.0	10.0
10	27.0	12.0	14.0	17.0	31.0	126.0	25.0	32.0
11	64.0	19.0	22.0	24.0	62.0	215.0	33.0	93.0
12	65.0	11.0	30.0	25.0	67.0	260.0	31.0	81.0
13	72.0	7.0	30.0	19.0	101.0	234.0	28.0	66.0
14	82.0	7.0	29.0	31.0	109.0	284.0	71.0	127.0
15	99.0	19.0	22.0	34.0	150.0	219.0	16.0	37.0
16	117.0	27.0	24.0	36.0	141.0	260.0	20.0	42.0
17	131.0	21.0	32.0	28.0	175.0	263.0	16.0	43.0
18	124.0	19.0	24.0	31.0	171.0	232.0	21.0	46.0
19	111.0	16.0	28.0	26.0	132.0	173.0	18.0	35.0
20	153.0	25.0	26.0	18.0	108.0	124.0	10.0	21.0
21	144.0	20.0	19.0	31.0	106.0	54.0	5.0	22.0
22	132.0	15.0	17.0	16.0	81.0	30.0	5.0	9.0
23	109.0	19.0	22.0	22.0	56.0	20.0	3.0	8.0

PRIMARY_ICD9_CODE	4280	430	431	486	5070	51881	5770	5789	5849	\
ADMISSION_HOUR										
0	17.0	23.0	50.0	17.0	21.0	41.0	24.0	10.0	8.0	
1	16.0	33.0	38.0	21.0	22.0	55.0	13.0	10.0	15.0	
2	9.0	29.0	40.0	26.0	19.0	33.0	11.0	10.0	16.0	
3	8.0	26.0	46.0	10.0	18.0	38.0	9.0	11.0	8.0	
4	10.0	15.0	31.0	14.0	7.0	35.0	7.0	8.0	5.0	
5	9.0	21.0	16.0	20.0	15.0	33.0	11.0	6.0	15.0	
6	6.0	8.0	8.0	10.0	17.0	25.0	10.0	10.0	8.0	
7	20.0	1.0	3.0	6.0	1.0	18.0	2.0	7.0	4.0	
8	9.0	4.0	11.0	12.0	6.0	20.0	3.0	5.0	4.0	
9	7.0	7.0	19.0	9.0	9.0	17.0	5.0	6.0	5.0	
10	8.0	8.0	15.0	8.0	11.0	18.0	6.0	9.0	5.0	
11	30.0	17.0	27.0	17.0	19.0	33.0	10.0	7.0	9.0	
12	23.0	19.0	47.0	23.0	15.0	45.0	5.0	13.0	17.0	
13	25.0	22.0	39.0	26.0	25.0	43.0	13.0	17.0	22.0	
14	45.0	21.0	41.0	38.0	30.0	60.0	27.0	37.0	22.0	
15	33.0	18.0	40.0	36.0	31.0	53.0	19.0	30.0	22.0	
16	40.0	27.0	65.0	53.0	45.0	75.0	28.0	27.0	45.0	
17	44.0	23.0	62.0	52.0	46.0	69.0	22.0	24.0	41.0	
18	49.0	44.0	79.0	56.0	43.0	72.0	35.0	30.0	48.0	
19	53.0	28.0	53.0	41.0	33.0	52.0	24.0	28.0	43.0	
20	26.0	31.0	70.0	54.0	40.0	90.0	39.0	23.0	66.0	
21	31.0	48.0	80.0	43.0	37.0	63.0	23.0	29.0	32.0	
22	20.0	38.0	78.0	42.0	30.0	73.0	28.0	28.0	37.0	
23	24.0	31.0	73.0	50.0	56.0	63.0	25.0	30.0	26.0	

PRIMARY_ICD9_CODE	V3000	V3001	V3101
ADMISSION_HOUR			
0	165.0	103.0	40.0
1	169.0	103.0	33.0
2	145.0	76.0	25.0
3	165.0	87.0	18.0
4	150.0	63.0	14.0
5	154.0	58.0	27.0
6	121.0	66.0	27.0
7	120.0	74.0	17.0
8	135.0	154.0	43.0
9	137.0	131.0	34.0
10	132.0	181.0	49.0
11	153.0	148.0	37.0
12	121.0	151.0	41.0
13	128.0	148.0	57.0
14	144.0	137.0	51.0
15	143.0	122.0	58.0
16	142.0	114.0	50.0
17	155.0	121.0	47.0
18	154.0	135.0	53.0
19	169.0	114.0	46.0
20	155.0	142.0	73.0
21	144.0	128.0	56.0
22	170.0	127.0	63.0
23	163.0	109.0	49.0

5.3.8 Conclusion

The distribution of the correlation between Diagnosis and Hour of Admission is very uniform. Some interesting outliers are 41401 (Coronary atherosclerosis of native coronary artery) and 4241 (Aortic valve disorders) with 7 AM.

5.3.9 First Care Unit & Diagnosis

```
[ ]: # 1. Define the BigQuery SQL query to get counts for each combination
#     First, find the top 20 most frequent ICD-9 codes, then get the counts
#     for combinations of FIRST_CAREUNIT and those top codes.
query = """
WITH TopCodes AS (
  -- Subquery to find the top 20 most frequent primary ICD-9 codes
  SELECT
    PRIMARY_ICD9_CODE
  FROM
    `reliable-jet-452114-s2.table.junction_table`
  WHERE
    PRIMARY_ICD9_CODE IS NOT NULL
```

```

GROUP BY
    PRIMARY_ICD9_CODE
ORDER BY
    COUNT(*) DESC
LIMIT 20
)
-- Main query to count combinations for top codes and first care unit
SELECT
    jt.FIRST_CAREUNIT,
    jt.PRIMARY_ICD9_CODE,
    COUNT(*) AS count
FROM
    `reliable-jet-452114-s2.table.junction_table` jt
JOIN
    TopCodes tc ON jt.PRIMARY_ICD9_CODE = tc.PRIMARY_ICD9_CODE -- Join to
    ↪filter for top codes
WHERE
    jt.FIRST_CAREUNIT IS NOT NULL
GROUP BY
    jt.FIRST_CAREUNIT,
    jt.PRIMARY_ICD9_CODE
ORDER BY
    jt.FIRST_CAREUNIT,
    jt.PRIMARY_ICD9_CODE;
"""

# 2. Execute the query and load results into a Pandas DataFrame
careunit_icd9_df = client.query(query).to_dataframe()

# 3. Create the visualization using Matplotlib and Seaborn
if not careunit_icd9_df.empty:
    heatmap_data = careunit_icd9_df.pivot(index='FIRST_CAREUNIT',
    ↪columns='PRIMARY_ICD9_CODE', values='count').fillna(0) # Fill missing combos
    ↪with 0

    plt.figure(figsize=(18, 8))

    # Create the heatmap
    heatmap_data = heatmap_data.astype(float)
    sns.heatmap(heatmap_data, annot=False, fmt=".0f", linewidths=.5,
    ↪cmap="Spectral", cbar_kws={'label': 'Number of Patients'})

    # Add labels and title
    plt.xlabel("Primary ICD-9 Code (Top 20)")
    plt.ylabel("First Care Unit")
    plt.title("Heatmap of Admission Counts by First Care Unit and Top 20
    ↪Primary ICD-9 Codes")


```

```
plt.xticks(rotation=90)
plt.yticks(rotation=0) # Ensure y-axis labels are horizontal

# Display the plot
plt.tight_layout()
plt.show()

# Display the pivoted data table
print("\nPivoted Data for Heatmap:")
print(heatmap_data)

else:
    print("No data returned from the query to plot.")
```



Pivoted Data for Heatmap:								
PRIMARY_ICD9_CODE	0389	1983	41011	41041	41071	41401	4240	4241 \
FIRST_CAREUNIT								
CCU	150.0	4.0	379.0	371.0	885.0	476.0	37.0	80.0
CSRU	31.0	4.0	62.0	67.0	684.0	3092.0	550.0	1075.0
MICU	1509.0	59.0	27.0	28.0	187.0	38.0	4.0	9.0
NICU	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
SICU	207.0	226.0	1.0	2.0	27.0	6.0	1.0	3.0
TSICU	91.0	50.0	4.0	3.0	20.0	5.0	0.0	1.0

PRIMARY_ICD9_CODE	4280	430	431	486	5070	51881	5770	5789 \
FIRST_CAREUNIT								
CCU	251.0	13.0	39.0	74.0	45.0	141.0	13.0	29.0
CSRU	107.0	18.0	14.0	11.0	12.0	22.0	7.0	9.0

MICU	173.0	64.0	163.0	535.0	474.0	835.0	231.0	317.0
NICU	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
SICU	23.0	351.0	627.0	50.0	43.0	86.0	99.0	43.0
TSICU	8.0	96.0	188.0	14.0	22.0	40.0	49.0	17.0

PRIMARY_ICD9_CODE	5849	V3000	V3001	V3101
FIRST_CAREUNIT				
CCU	85.0	0.0	0.0	0.0
CSRU	7.0	0.0	0.0	0.0
MICU	384.0	0.0	0.0	0.0
NICU	0.0	3534.0	2792.0	1008.0
SICU	34.0	0.0	0.0	0.0
TSICU	13.0	0.0	0.0	0.0

5.3.10 Conclusion

The strongest connections between First Care Unit and Diagnosis are V3000 (Single liveborn, delivered in hospital, without cesarean section) and NICU (Neonatal Intensive Care Unit) and 41401 (Coronary atherosclerosis of native coronary artery) and CSRU (Cardiac Surgery Recovery Unit).

5.3.11 First Care Unit & Death

```
[ ]: start_time = time.time()

# 1. Define the BigQuery SQL query to get counts for each combination
query = """
SELECT
    FIRST_CAREUNIT,
    EXPIRE_FLAG,
    COUNT(*) AS count
FROM
    `reliable-jet-452114-s2.table.junction_table`
WHERE
    FIRST_CAREUNIT IS NOT NULL
    AND EXPIRE_FLAG IS NOT NULL
GROUP BY
    FIRST_CAREUNIT,
    EXPIRE_FLAG
ORDER BY
    FIRST_CAREUNIT,
    EXPIRE_FLAG;
"""

# 2. Execute the query and load results into a Pandas DataFrame
careunit_expire_df = client.query(query).to_dataframe()

# 3. Create the visualization using Matplotlib and Seaborn
```

```

if not careunit_expire_df.empty:
    # Map EXPIRE_FLAG to readable labels
    careunit_expire_df['Outcome'] = careunit_expire_df['EXPIRE_FLAG'].map({0: 'Survived', 1: 'Expired'})

    heatmap_data = careunit_expire_df.pivot(index='FIRST_CAREUNIT',
    columns='Outcome', values='count').fillna(0) # Fill missing combos with 0

    plt.figure(figsize=(10, 7))

    # Create the heatmap
    heatmap_data = heatmap_data.astype(float)
    sns.heatmap(heatmap_data, annot=True, fmt=".0f", linewidths=.5,
    cmap="Spectral", cbar_kws={'label': 'Number of Patients'})

    # Add labels and title
    plt.xlabel("Patient Outcome During Hospital Stay")
    plt.ylabel("First Care Unit")
    plt.title("Heatmap of Patient Outcomes by First Care Unit")
    plt.yticks(rotation=0) # Ensure y-axis labels are horizontal


    # Display the plot
    plt.tight_layout()
    plt.show()

    # Display the pivoted data table
    print("\nPivoted Data for Heatmap:")
    print(heatmap_data)

else:
    print("No data returned from the query to plot.")

end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")

```



Pivoted Data for Heatmap:

Outcome	Expired	Survived
FIRST_CAREUNIT		
CCU	3432.0	3700.0
CSRU	2475.0	6612.0
MICU	10152.0	9188.0
NICU	67.0	8008.0
SICU	3619.0	4755.0
TSICU	1991.0	4040.0

Query Execution Time: 2.20 seconds

5.3.12 Conclusion

In half of the First Care Units, the ratio between patients that expired and survived is fairly equal (and close to 1). In the other half, which include the units NICU, TSICU and CSRU, the number of patients that survived is much higher than the number of patients that expired.

5.3.13 Diagnosis & Death

```
[ ]: # 1. Define the BigQuery SQL query to get counts for each combination
#     First, find the top 20 most frequent ICD-9 codes, then get the counts
#     for combinations of those top codes and EXPIRE_FLAG.
query = """
WITH TopCodes AS (
    -- Subquery to find the top 20 most frequent primary ICD-9 codes
    SELECT
        PRIMARY_ICD9_CODE
    FROM
        `reliable-jet-452114-s2.table.junction_table`
    WHERE
        PRIMARY_ICD9_CODE IS NOT NULL
    GROUP BY
        PRIMARY_ICD9_CODE
    ORDER BY
        COUNT(*) DESC
    LIMIT 20
)
-- Main query to count combinations for top codes and expire flag
SELECT
    jt.PRIMARY_ICD9_CODE,
    jt.EXPIRE_FLAG,
    COUNT(*) AS count
FROM
    `reliable-jet-452114-s2.table.junction_table` jt
JOIN
    TopCodes tc ON jt.PRIMARY_ICD9_CODE = tc.PRIMARY_ICD9_CODE -- Join to_
    ↪filter for top codes
WHERE
    jt.EXPIRE_FLAG IS NOT NULL
GROUP BY
    jt.PRIMARY_ICD9_CODE,
    jt.EXPIRE_FLAG
ORDER BY
    jt.PRIMARY_ICD9_CODE,
    jt.EXPIRE_FLAG;
"""

# 2. Execute the query and load results into a Pandas DataFrame
icd9_expire_df = client.query(query).to_dataframe()

# 3. Create the visualization using Matplotlib and Seaborn
if not icd9_expire_df.empty:
    # Map EXPIRE_FLAG to readable labels
```



```

    icd9_expire_df['Outcome'] = icd9_expire_df['EXPIRE_FLAG'].map({0: 'Survived', 1: 'Expired'})

    heatmap_data = icd9_expire_df.pivot(index='PRIMARY_ICD9_CODE',
    ↪ columns='Outcome', values='count').fillna(0) # Fill missing combos with 0

    # Reindex based on total frequency
    total_counts = icd9_expire_df.groupby('PRIMARY_ICD9_CODE')['count'].sum().
    ↪ sort_values(ascending=False)
    heatmap_data = heatmap_data.reindex(total_counts.index)

    plt.figure(figsize=(10, 10))

    # Create the heatmap
    heatmap_data = heatmap_data.astype(float)
    sns.heatmap(heatmap_data, annot=True, fmt=".0f", linewidths=.5,
    ↪ cmap="vlag", cbar_kws={'label': 'Number of Patients'})


    # Add labels and title
    plt.xlabel("Patient Outcome During Hospital Stay")
    plt.ylabel("Primary ICD-9 Code (Top 20)")
    plt.title("Heatmap of Patient Outcomes by Top 20 Primary ICD-9 Codes")
    plt.yticks(rotation=0) # Ensure y-axis labels are horizontal
    plt.xticks(rotation=0)

    # Display the plot
    plt.tight_layout()
    plt.show()

    # Display the pivoted data table
    print("\nPivoted Data for Heatmap:")
    print(heatmap_data)

else:
    print("No data returned from the query to plot.")

```



Pivoted Data for Heatmap:

Outcome	Expired	Survived
PRIMARY_ICD9_CODE		
41401	794.0	2823.0
V3000	12.0	3522.0
V3001	24.0	2768.0
0389	1314.0	674.0
41071	823.0	980.0
4241	238.0	930.0
51881	724.0	400.0
431	601.0	430.0

V3101	16.0	992.0
486	430.0	254.0
5070	398.0	198.0
4240	98.0	494.0
4280	427.0	135.0
430	164.0	378.0
5849	315.0	208.0
41011	177.0	296.0
41041	151.0	320.0
5789	220.0	195.0
5770	133.0	266.0
1983	259.0	84.0

5.3.14 Conclusion

The HeatMap that compares **Diagnosis** and **Death** is fairly uniform, with newborns counting the lowest deaths (as expected).

6 XGBoost Classifier Predictor Model

In this section we use a **XGBoost Classifier** model to predict the **ICU Length of Stay**, that we divide in bins of two days ([1-3] days, [3-5] days, ...). The **window size** we chose was **1 day** (24 hours), given that most data is collected in that window and most records have a value of **ICU Length of Stay** lower than 3 days, which means a bigger window size would exclude most of the data. We divided this section of the work in 2 parts: 1. Dividing the dataset, Categorizing the Target and Implementing Custom Sample Weights 2. Pre-Processing the features 3. Applying the model 4. Visualizing the results

The discussion of results was done in the final conclusions.

6.1 Dividing the Dataset, Categorizing the Target and Implementing Custom Sample Weights

Here we will divide the dataset in Train, Test and Validation. We will do this by: 1. Eliminating the features that are only collected after the first 24 hours (our **window size**) or are not relevant 2. Removing the records whose value for **ICU Length of Stay** is lower than 24 hours 3. Using the library **scikit-learn** to perform the division 4. Converting the target feature from **continuous values** to **bins** 5. Calculating the custom sample weights to ensure a less biased prediction towards the majority class 6. Normalizing the sample weights

```
[4]: # --- Record Start Time ---
start_time = time.time()

# --- Configuration ---
TARGET_TABLE = "reliable-jet-452114-s2.table.junction_table"
TARGET_VARIABLE = 'ICU_LOS'
WINDOW_DAYS = 1.0      # Minimum ICU stay to be included
MAX_ICU_LOS_DAYS = 25.0 # Maximum ICU stay to be included
```

```

BIN_WIDTH = 2.0          # Width of each LOS category
TEST_SIZE = 0.10
VALIDATION_SIZE = 0.10
CUSTOM_WEIGHT_ALPHA = 0.8 # Tunable parameter for custom weighting

# --- 1. Load Data from BigQuery ---
print("--- Block 1: Dividing Dataset, Target Categorization & CUSTOM Sample_
↳Weights ---")
try:
    load_query = f"SELECT * FROM `{TARGET_TABLE}`"
    full_df = client.query(load_query).to_dataframe()
    print(f"Loaded {len(full_df)} rows from {TARGET_TABLE}.")
except Exception as e:
    print(f"Error loading data from BigQuery: {e}")
    full_df = pd.DataFrame()

if not full_df.empty:
    original_row_count = len(full_df)
    print(f"Original row count: {original_row_count}")

    df_min_filtered = full_df[full_df[TARGET_VARIABLE] >= WINDOW_DAYS].copy()
    rows_after_min_los_filter = len(df_min_filtered)
    print(f"Rows remaining after MIN LOS filter (>= {WINDOW_DAYS} days):_
↳{rows_after_min_los_filter}")

    filtered_df = df_min_filtered[df_min_filtered[TARGET_VARIABLE] <=_
↳MAX_ICU_LOS_DAYS].copy()
    rows_after_max_los_filter = len(filtered_df)
    print(f"Rows remaining after MAX LOS filter (<= {MAX_ICU_LOS_DAYS} days):_
↳{rows_after_max_los_filter}")

    if rows_after_max_los_filter > 0:
        features_to_exclude = [
            'SUBJECT_ID', 'HADM_ID', 'ICUSTAY_ID', 'DOB',
            'ADMITTIME', 'DISCHTIME', 'ICU_INTIME', 'ICU_OUTTIME',
            'PRIMARY_ICD9_TITLE', 'EXPIRE_FLAG'
        ]
        features_for_training_initial = [
            col for col in filtered_df.columns
            if col not in features_to_exclude and col != TARGET_VARIABLE
        ]
        print(f"Initial features considered for training:_
↳{features_for_training_initial}")

        X = filtered_df[features_for_training_initial]
        y_continuous = filtered_df[TARGET_VARIABLE]

```

```

        if len(X) < 3:
            print("Not enough data to perform train-validation-test split_
↳after all filtering.")
            X_train, X_val, X_test = pd.DataFrame(), pd.DataFrame(), pd.
↳DataFrame()
            y_train_cat, y_val_cat, y_test_cat = pd.Series(dtype=int), pd.
↳Series(dtype=int), pd.Series(dtype=int)
            train_sample_weights = np.array([])
            num_classes = 1
        else:
            train_val_size = 1.0 - TEST_SIZE
            X_train_val, X_test, y_train_val_continuous, y_test_continuous = _
↳train_test_split(
                X, y_continuous, test_size=TEST_SIZE, random_state=42, _
↳shuffle=True
            )
            validation_split_ratio_for_temp = VALIDATION_SIZE / train_val_size_
↳if train_val_size > 0 else 0

            if len(X_train_val) < 2 or validation_split_ratio_for_temp == 0:
                print("Not enough data in train_val set to split into train and_
↳validation or invalid ratio. Using train_val as train.")
                X_train, X_val, y_train_continuous, y_val_continuous = _
↳X_train_val, pd.DataFrame(), y_train_val_continuous, pd.Series(dtype=float)
            else:
                X_train, X_val, y_train_continuous, y_val_continuous = _
↳train_test_split(
                    X_train_val, y_train_val_continuous, _
↳test_size=validation_split_ratio_for_temp, random_state=42, shuffle=True
                )

            def convert_los_to_categories(los_series, bin_width, min_los):
                if los_series.empty:
                    return pd.Series([], dtype=int)
                adjusted_los = los_series - min_los
                categories = np.floor(adjusted_los / bin_width).astype(int)
                categories = np.maximum(0, categories)
                return categories

            y_train_cat = convert_los_to_categories(y_train_continuous, _
↳bin_width=BIN_WIDTH, min_los=WINDOW_DAYS)
            y_val_cat = convert_los_to_categories(y_val_continuous, _
↳bin_width=BIN_WIDTH, min_los=WINDOW_DAYS)
            y_test_cat = convert_los_to_categories(y_test_continuous, _
↳bin_width=BIN_WIDTH, min_los=WINDOW_DAYS)

```

```

all_cats = pd.concat([y_train_cat, y_val_cat, y_test_cat]).dropna().
↳astype(int)
    if all_cats.empty:
        print("Warning: No categories found in target variable after_
↳splitting and conversion. Defaulting num_classes to 1.")
        num_classes = 1
        max_cat_overall = 0
        train_sample_weights = np.array([])
    else:
        max_cat_overall = all_cats.max()
        num_classes = max_cat_overall + 1

    if not y_train_cat.empty:
        print(f"\n--- Calculating CUSTOM Sample Weights_
↳(alpha={CUSTOM_WEIGHT_ALPHA}) ---")
        class_counts = Counter(y_train_cat)
        total_samples = len(y_train_cat)

        # Calculate custom class weights
        class_weights_map = {}
        for class_label, count in class_counts.items():
            if count > 0: # Ensure count is positive
                # Using raw count in the denominator for simplicity
                weight = 1.0 / (count ** CUSTOM_WEIGHT_ALPHA)
                class_weights_map[class_label] = weight
            else:
                class_weights_map[class_label] = 1.0

        print(f"Custom class weights map: {class_weights_map}")

        # Generate sample weights for each training instance
        train_sample_weights = np.array([class_weights_map.get(cls,
↳1.0) for cls in y_train_cat])

        # Normalize sample weights so their sum equals the number_
↳of samples
        train_sample_weights = (train_sample_weights / np.
↳sum(train_sample_weights)) * total_samples

        print(f"Computed CUSTOM sample weights for y_train_cat._
↳Shape: {train_sample_weights.shape}")
        unique_weights, counts = np.unique(train_sample_weights,
↳return_counts=True)
        print(f"Unique custom sample weights and their counts in_
↳training data: {dict(zip(unique_weights, counts))}")
    else:

```

```

        train_sample_weights = np.array([])
        print("y_train_cat is empty, no sample weights computed.")

        print(f"\nTarget variable converted to {num_classes} categories (0_
↪to {max_cat_overall}).")
        if not y_train_cat.empty: print(f"Unique categories in y_train_cat:
↪{sorted(y_train_cat.unique())}")
        if not y_val_cat.empty: print(f"Unique categories in y_val_cat:
↪{sorted(y_val_cat.unique())}")
        if not y_test_cat.empty: print(f"Unique categories in y_test_cat:
↪{sorted(y_test_cat.unique())}")

        print("\nDataset shapes:")
        print(f"X_train: {X_train.shape}, y_train_cat: {y_train_cat.shape}")
        if not X_val.empty: print(f"X_val: {X_val.shape}, y_val_cat:
↪{y_val_cat.shape}")
        else: print("X_val is empty.")
        print(f"X_test: {X_test.shape}, y_test_cat: {y_test_cat.shape}")

        if original_row_count > 0:
            percentage_kept_overall = (rows_after_max_loss_filter /
↪original_row_count) * 100
            print(f"\nPercentage of original records kept after ALL
↪filtering: {percentage_kept_overall:.2f}%")
            else:
                print("No data remaining after all filtering. Cannot proceed.")
                X_train, X_val, X_test = pd.DataFrame(), pd.DataFrame(), pd.DataFrame()
                y_train_cat, y_val_cat, y_test_cat = pd.Series(dtype=int), pd.
↪Series(dtype=int), pd.Series(dtype=int)
                features_for_training_initial = []
                train_sample_weights = np.array([])
                num_classes = 1
else:
    print("Data loading failed or table was empty. Cannot proceed.")
    X_train, X_val, X_test = pd.DataFrame(), pd.DataFrame(), pd.DataFrame()
    y_train_cat, y_val_cat, y_test_cat = pd.Series(dtype=int), pd.
↪Series(dtype=int), pd.Series(dtype=int)
    features_for_training_initial = []
    train_sample_weights = np.array([])
    num_classes = 1

# --- Record End Time ---
end_time = time.time()
execution_time = end_time - start_time
print(f"\nExecution time: {execution_time:.2f} seconds")

```

```

--- Block 1: Dividing Dataset, Target Categorization & CUSTOM Sample Weights ---
Loaded 58039 rows from reliable-jet-452114-s2.table.junction_table.
Original row count: 58039
Rows remaining after MIN LOS filter (>= 1.0 days): 46321
Rows remaining after MAX LOS filter (<= 25.0 days): 44439
Initial features considered for training: ['GENDER', 'AGE_AT_ADMISSION',
'ADMISSION_TYPE', 'ADMISSION_LOCATION', 'INSURANCE', 'ETHNICITY',
'ADMISSION_DIAGNOSIS_TEXT', 'ADMISSION_HOUR', 'FIRST_CAREUNIT',
'PRIMARY_ICD9_CODE']

--- Calculating CUSTOM Sample Weights (alpha=0.8) ---
Custom class weights map: {1: 0.0008837269036497867, 0: 0.0003628979779803216,
2: 0.0016682574871599298, 7: 0.007055908051730422, 5: 0.004504175899347122, 11:
0.012811784118365512, 4: 0.003562410621944747, 3: 0.0025666245599461337, 6:
0.0054845393635214, 8: 0.008543413297975137, 9: 0.010266510574570422, 10:
0.012552737133001474}
Computed CUSTOM sample weights for y_train_cat. Shape: (35551,)
Unique custom sample weights and their counts in training data:
{0.2585191274775793: 19965, 0.6295441747333023: 6563, 1.1884234582643785: 2966,
1.828396910593369: 1731, 2.537769129570504: 1149, 3.208658339694305: 857,
3.9070438991282415: 670, 5.026446284565101: 489, 6.086106524386989: 385,
7.3135964293557105: 306, 8.94224506055231: 238, 9.126783428621254: 232}

Target variable converted to 12 categories (0 to 11).
Unique categories in y_train_cat: [0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11]
Unique categories in y_val_cat:  [0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11]
Unique categories in y_test_cat: [0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11]

Dataset shapes:
X_train: (35551, 10), y_train_cat: (35551,)
X_val:   (4444, 10), y_val_cat:   (4444,)
X_test:  (4444, 10), y_test_cat:  (4444,)

Percentage of original records kept after ALL filtering: 76.57%

Execution time: 1.76 seconds

```

6.2 Pre-Processing the Features

Here we will pre-process the features before applying the predictor model. We will do this by: 1. Defining the features to be used in training and their respective category 2. Feature engineer the features based on their category: - Scale for Continuous Numerical features - Encode using sin and cos for Cyclical Numerical features - One-hot Encode for Low/Medium Cardinality Categorical Features - Dictifier Transformation followed by a FeatureHasher Transformation for High Cardinality Categorical features

```

[5]: # --- Record Start Time ---
start_time = time.time()

```



```

print("\n--- Block 2: Pre-Processing the Features ---")

# If X_train is empty, skip preprocessing
if 'X_train' in globals() and not X_train.empty:

    # --- 1. Define Feature Lists (based on X_train columns) ---

    _features_for_training = [
        'AGE_AT_ADMISSION',      # Numerical continuous
        'ADMISSION_TYPE',        # Categorical low cardinality
        'INSURANCE',             # Categorical low cardinality
        'PRIMARY_ICD9_CODE',     # Categorical high cardinality
        'FIRST_CAREUNIT',        # Categorical low cardinality
        'ADMISSION_HOUR'         # Numerical cyclical
    ]
    # Filter features to only those present in the training data from Block 1
    features_for_training_final = [f for f in _features_for_training if f in X_train.columns]
    print(f"Features selected for preprocessing: {features_for_training_final}")

    numerical_cont_features = [f for f in ['AGE_AT_ADMISSION'] if f in features_for_training_final]
    numerical_cycl_features = [f for f in ['ADMISSION_HOUR'] if f in features_for_training_final] # e.g., 0-23
    categorical_low_card_features = [f for f in ['ADMISSION_TYPE', 'INSURANCE', 'FIRST_CAREUNIT'] if f in features_for_training_final]
    categorical_high_card_features = [f for f in ['PRIMARY_ICD9_CODE'] if f in features_for_training_final]

    print(f"Continuous numerical features: {numerical_cont_features}")
    print(f"Cyclical numerical features: {numerical_cycl_features}")
    print(f"Low/Medium cardinality categorical features: {categorical_low_card_features}")
    print(f"High cardinality categorical features: {categorical_high_card_features}")

    # --- 2. Feature Engineering & Preprocessing Pipeline Components ---

    # a) Continuous Numerical Features: Scale
    numerical_cont_transformer = Pipeline(steps=[
        ('scaler', StandardScaler())
    ])

    # b) Cyclical Numerical Features: Encode hour using sine and cosine
    def sin_transformer(X_in):

```

```

    # Ensure input is 2D for FunctionTransformer
    X_proc = X_in.copy()
    if isinstance(X_proc, pd.Series): X_proc = X_proc.to_frame()
    elif X_proc.ndim == 1: X_proc = X_proc.reshape(-1, 1)
    return np.sin(2 * np.pi * X_proc / 24.0)

def cos_transformer(X_in):
    X_proc = X_in.copy()
    if isinstance(X_proc, pd.Series): X_proc = X_proc.to_frame()
    elif X_proc.ndim == 1: X_proc = X_proc.reshape(-1, 1)
    return np.cos(2 * np.pi * X_proc / 24.0)

# c) Low/Medium Cardinality Categorical Features: One-hot encode
categorical_low_card_transformer = Pipeline(steps=[
    ('onehot', OneHotEncoder(handle_unknown='ignore', sparse_output=False))
])

# d) High Cardinality Categorical Features: Feature Hashing
# Custom Transformer for FeatureHasher Input
class Dictifier(BaseEstimator, TransformerMixin):
    def __init__(self, col_name):
        self.col_name = col_name
    def fit(self, X, y=None):
        return self
    def transform(self, X_in):
        # X_in is expected to be a DataFrame/Series for the specific column
        if isinstance(X_in, pd.DataFrame): series = X_in.iloc[:, 0]
        elif isinstance(X_in, pd.Series): series = X_in
        else: series = pd.Series(X_in.flatten()) # Fallback for numpy array
        return [{self.col_name: str(val)} for val in series]

# --- Create the list of transformers for ColumnTransformer ---
transformers_list = []

if numerical_cont_features:
    transformers_list.append(('num_cont', numerical_cont_transformer,
↪numerical_cont_features))
    if numerical_cycl_features: # Assuming only one cyclical feature
↪'ADMISSION_HOUR' for this example
        transformers_list.append(('hour_sin',
↪FunctionTransformer(sin_transformer, validate=False),
↪numerical_cycl_features))
        transformers_list.append(('hour_cos',
↪FunctionTransformer(cos_transformer, validate=False),
↪numerical_cycl_features))
    if categorical_low_card_features:

```

```

        transformers_list.append(('cat_low', categorical_low_card_transformer,
↪categorical_low_card_features))

    if categorical_high_card_features:
        n_hash_features = 50 # Number of features for the hasher, adjust as
↪needed
        for i, col_name in enumerate(categorical_high_card_features):
            print(f" - Adding hasher for high-cardinality feature: {col_name}")
            transformer_name = f'cat_high_{col_name.replace(" ", "_").lower()}'
↪# Unique name
            high_card_pipeline = Pipeline(steps=[
                ('dictifier', Dictifier(col_name=col_name)),
                ('hasher', FeatureHasher(n_features=n_hash_features,
↪input_type='dict'))
            ])
            transformers_list.append((transformer_name, high_card_pipeline,
↪[col_name])) # Pass as list
        else:
            print("No high cardinality features specified or found for hashing.")

    # --- Define the main ColumnTransformer ---
    if transformers_list:
        preprocessor = ColumnTransformer(
            transformers=transformers_list,
            remainder='drop', # Drop any columns not specified in
↪features_for_training_final
            n_jobs=-1
        )
        print("Preprocessor defined.")
    else:
        print("No features to preprocess. Preprocessor not created.")
        preprocessor = 'passthrough' # Or handle as an error/empty pipeline

else:
    print("X_train is empty or not defined. Skipping feature preprocessing.")
    features_for_training_final = []
    preprocessor = 'passthrough' # To avoid error in next block

# --- Record End Time ---
end_time = time.time()
execution_time = end_time - start_time
print(f"\nExecution time: {execution_time:.2f} seconds")

```

--- Block 2: Pre-Processing the Features ---

Features selected for preprocessing: ['AGE_AT_ADMISSION', 'ADMISSION_TYPE', 'INSURANCE', 'PRIMARY_ICD9_CODE', 'FIRST_CAREUNIT', 'ADMISSION_HOUR']

Continuous numerical features: ['AGE_AT_ADMISSION']
 Cyclical numerical features: ['ADMISSION_HOUR']
 Low/Medium cardinality categorical features: ['ADMISSION_TYPE', 'INSURANCE', 'FIRST_CAREUNIT']
 High cardinality categorical features: ['PRIMARY_ICD9_CODE']
 - Adding hasher for high-cardinality feature: PRIMARY_ICD9_CODE
 Preprocessor defined.

Execution time: 0.00 seconds

6.3 Applying the XGBoost Classifier Model

Here we will apply the XGBoost model to the dataset. We will do this by:

1. Defining the model pipeline
2. Training the model
3. Making predictions using the trained model
4. Post-Process predictions

```
[7]: # --- Record Start Time ---
start_time_block3_weighted = time.time()
print("\n--- Block 3: Applying the XGBoost Model (Classification) - MANUAL_
↳PREDICTION FALLBACK & SAMPLE WEIGHTS ---")

# Ensure necessary variables are available
if ('X_train' not in globals() or X_train.empty or
    'y_train_cat' not in globals() or y_train_cat.empty or
    'preprocessor' not in globals() or
    'features_for_training_final' not in globals() or
    'num_classes' not in globals() or
    'train_sample_weights' not in globals()):
    print("Necessary data (X_train, y_train_cat, preprocessor, features,
↳num_classes, train_sample_weights) not available or y_train_cat is empty.
↳Skipping model training.")
    y_pred_val_final = pd.Series(dtype=int)
    y_pred_test_final = pd.Series(dtype=int)
    xgb_model_pipeline = None
    fitted_pipeline_preprocessor = None
    fitted_pipeline_classifier = None
else:
    if y_train_cat.empty:
        print(" y_train_cat is EMPTY. Cannot proceed with model training.")
        xgb_model_pipeline = None
        fitted_pipeline_preprocessor = None
        fitted_pipeline_classifier = None
    elif train_sample_weights is None or train_sample_weights.size == 0:
        print(" train_sample_weights is None or empty. Cannot proceed with
↳weighted training.")
```

```

xgb_model_pipeline = None
fitted_pipeline_preprocessor = None
fitted_pipeline_classifier = None
else:
    print(f" y_train_cat.shape: {y_train_cat.shape}, dtype: {y_train_cat.
↳dtype}")
    if y_train_cat.isnull().any():
        print(f" WARNING: y_train_cat contains NaNs! Count: {y_train_cat.
↳isnull().sum()}")
        y_train_cat = y_train_cat.astype(int) # Ensure it's int
        print(f" Type of target (y_train_cat): {type_of_target(y_train_cat)}")
        print(f" y_train_cat.min(): {y_train_cat.min()}, y_train_cat.max():
↳{y_train_cat.max()}")

    current_X_train_cols = X_train.columns.tolist()
    valid_features_for_model = [f for f in features_for_training_final if f
↳in current_X_train_cols]
    if not valid_features_for_model:
        print("WARNING: No valid features for model found in X_train.
↳Preprocessing might be incorrect or use no features.")

    X_train_processed_standalone = None
    try:
        if preprocessor == 'passthrough':
            X_train_processed_standalone =
↳X_train[valid_features_for_model].copy() if valid_features_for_model else
↳X_train.copy()
            elif hasattr(preprocessor, 'fit_transform'):
                temp_preprocessor_standalone = clone(preprocessor)
                print(" Fitting and transforming X_train with a cloned
↳preprocessor for standalone test...")
                X_train_processed_standalone = temp_preprocessor_standalone.
↳fit_transform(X_train[valid_features_for_model])
                print(f" X_train_processed_standalone shape:
↳{X_train_processed_standalone.shape}")
            else:
                print(" Preprocessor is not 'passthrough' and does not have
↳'fit_transform'. Cannot preprocess for standalone test.")
                except Exception as e_preprocess_standalone:
                    print(f" ERROR during X_train preprocessing for standalone test:
↳{e_preprocess_standalone}")

        if X_train_processed_standalone is not None:
            standalone_xgb = XGBClassifier(objective='multi:softmax',
↳num_class=num_classes, n_estimators=50, learning_rate=0.1, max_depth=3,
↳random_state=42, n_jobs=-1, eval_metric='mlogloss')

```

```

        try:
            # Pass sample_weight to standalone fit
            standalone_xgb.fit(X_train_processed_standalone, y_train_cat,
↪sample_weight=train_sample_weights)
            print(" Standalone XGBClassifier fitted successfully with
↪sample_weight!")
        except Exception as e_standalone:
            print(f" ERROR fitting standalone XGBClassifier with
↪sample_weight: {e_standalone}")
        else:
            print(" Skipping Standalone XGBClassifier test as X_train could
↪not be processed for it.")

        # preprocessor is the definition from Block 2. Pipeline will fit it.
        xgb_model_pipeline = Pipeline(steps=[
            ('preprocessor', preprocessor),
            ('classifier', XGBClassifier(objective='multi:softmax',
↪num_class=num_classes, n_estimators=100, learning_rate=0.1, max_depth=5,
↪subsample=0.8, colsample_bytree=0.8, random_state=42, n_jobs=-1,
↪eval_metric='mlogloss'))
        ])

        fitted_pipeline_preprocessor = None
        fitted_pipeline_classifier = None
        pipeline_fit_successful = False

        try:
            print(f" Attempting to fit the pipeline on
↪X_train[valid_features_for_model] (shape: {X_train[valid_features_for_model]
↪shape} if valid_features_for_model else X_train.shape) and y_train_cat...")
            fit_params = {'classifier__sample_weight': train_sample_weights}
            xgb_model_pipeline.fit(X_train, y_train_cat, **fit_params)
            print(" Pipeline fitted successfully with sample_weight!")
            pipeline_fit_successful = True
            fitted_pipeline_preprocessor = xgb_model_pipeline.
↪named_steps['preprocessor']
            fitted_pipeline_classifier = xgb_model_pipeline.
↪named_steps['classifier']

        except Exception as e_pipeline_fit:
            print(f" ERROR during pipeline.fit() with sample_weight:
↪{e_pipeline_fit}")
            xgb_model_pipeline = None # Ensure pipeline is None if fit fails

        # Initialize prediction variables
        y_pred_val_final = pd.Series(dtype=int)

```

```

y_pred_test_final = pd.Series(dtype=int)

if pipeline_fit_successful and fitted_pipeline_preprocessor and
↳fitted_pipeline_classifier:
    # Validation set predictions
    if 'X_val' in globals() and not X_val.empty:
        if 'y_val_cat' in globals() and not y_val_cat.empty: # Check if
↳there's a target to compare
            try:
                # Use the same features for transform as were used for
↳training the preprocessor
                X_val_to_transform = X_val # Pass the X_val with all
↳original features
                X_val_processed = fitted_pipeline_preprocessor.
↳transform(X_val_to_transform)
                print(f"    X_val_processed shape: {X_val_processed.
↳shape}")
                y_pred_val_cat_manual = fitted_pipeline_classifier.
↳predict(X_val_processed)
                y_pred_val_final = pd.Series(y_pred_val_cat_manual,
↳index=X_val.index)
                print(f"    Manual predictions on validation set
↳successful. Shape: {y_pred_val_final.shape}")
                except Exception as e_manual_pred_val:
                    print(f"    ERROR during manual prediction on
↳validation set: {e_manual_pred_val}")
                else:
                    print("    Validation set target (y_val_cat) is empty.
↳Skipping manual validation predictions.")
                else:
                    print("    Validation set (X_val) is empty or not available.
↳Skipping manual validation predictions.")

    # Test set predictions
    if 'X_test' in globals() and not X_test.empty:
        if 'y_test_cat' in globals() and not y_test_cat.empty: # Check
↳if there's a target to compare
            try:
                X_test_to_transform = X_test # Pass the X_test with all
↳original features
                X_test_processed = fitted_pipeline_preprocessor.
↳transform(X_test_to_transform)
                print(f"    X_test_processed shape: {X_test_processed.
↳shape}")
                y_pred_test_cat_manual = fitted_pipeline_classifier.
↳predict(X_test_processed)

```

```

        y_pred_test_final = pd.Series(y_pred_test_cat_manual,
↪index=X_test.index)
        print(f"    Manual predictions on test set successful.
↪Shape: {y_pred_test_final.shape}")
        except Exception as e_manual_pred_test:
            print(f"    ERROR during manual prediction on test set:
↪{e_manual_pred_test}")
        else:
            print("    Test set target (y_test_cat) is empty. Skipping
↪manual test predictions.")
        else:
            print("    Test set (X_test) is empty or not available. Skipping
↪manual test predictions.")

    elif xgb_model_pipeline is not None:
        try:
            if 'X_val' in globals() and not X_val.empty and 'y_val_cat' in
↪globals() and not y_val_cat.empty:
                y_pred_val_cat_pipeline = xgb_model_pipeline.predict(X_val)
↪# Use full X_val
                y_pred_val_final = pd.Series(y_pred_val_cat_pipeline,
↪index=X_val.index)
                print(f"    Pipeline predictions on validation set successful.
↪ Shape: {y_pred_val_final.shape}")
                if 'X_test' in globals() and not X_test.empty and 'y_test_cat'
↪in globals() and not y_test_cat.empty:
                    y_pred_test_cat_pipeline = xgb_model_pipeline.
↪predict(X_test) # Use full X_test
                    y_pred_test_final = pd.Series(y_pred_test_cat_pipeline,
↪index=X_test.index)
                    print(f"    Pipeline predictions on test set successful.
↪Shape: {y_pred_test_final.shape}")
                    except Exception as e_pipeline_predict:
                        print(f"    ERROR during pipeline.predict() fallback:
↪{e_pipeline_predict}")
                    else:
                        print("\nPipeline did not fit successfully. No predictions will be
↪made.")

# --- Record End Time ---
end_time = time.time()
execution_time = end_time - start_time
print(f"\nXGBoost training & prediction execution time (MANUAL PREDICTION
↪FALLBACK & SAMPLE WEIGHTS block): {execution_time:.2f} seconds")

```



```

--- Block 3: Applying the XGBoost Model (Classification) - MANUAL PREDICTION
FALLBACK & SAMPLE WEIGHTS ---
y_train_cat.shape: (35551,), dtype: int64
Type of target (y_train_cat): multiclass
y_train_cat.min(): 0, y_train_cat.max(): 11
Fitting and transforming X_train with a cloned preprocessor for standalone
test...
X_train_processed_standalone shape: (35551, 68)
Standalone XGBClassifier fitted successfully with sample_weight!
Attempting to fit the pipeline on X_train[valid_features_for_model] (shape:
(35551, 6)) and y_train_cat...
Pipeline fitted successfully with sample_weight!
X_val_processed shape: (4444, 68)
Manual predictions on validation set successful. Shape: (4444,)
X_test_processed shape: (4444, 68)
Manual predictions on test set successful. Shape: (4444,)

XGBoost training & prediction execution time (MANUAL PREDICTION FALLBACK &
SAMPLE WEIGHTS block): 31.53 seconds

```

6.4 Vizualizing the Results

To end we use statistical measures and adequate plots to visualize the prediction results.

```

[8]: # --- Record Start Time ---
start_time = time.time()
print("\n--- Block 4: Evaluating and Visualizing Results ---")

# Check if necessary data for evaluation is available
if ('y_val_cat' not in globals() or
    'y_test_cat' not in globals() or
    'y_pred_val_final' not in globals() or
    'y_pred_test_final' not in globals() or
    'num_classes' not in globals()):
    print("Necessary data for evaluation/visualization is missing. Skipping.")
else:
    # Define descriptive labels for your bins for plotting
    bin_labels_for_plots = []
    if 'WINDOW_DAYS' in globals() and 'BIN_WIDTH' in globals() and num_classes_
    ↪ > 0 :
        for i in range(num_classes):
            lower_bound = WINDOW_DAYS + i * BIN_WIDTH
            upper_bound = WINDOW_DAYS + (i + 1) * BIN_WIDTH
            bin_labels_for_plots.append(f"[{lower_bound:.1f}-{upper_bound:.
            ↪ 1f})")
    else: # Fallback labels
        bin_labels_for_plots = [f"Cat {i}" for i in range(num_classes)]

```

```

print(f"Using bin labels for plots: {bin_labels_for_plots}")

# --- 1. Evaluate the Model ---
# Validation Set Metrics (if y_val_cat and y_pred_val_final are not empty)
if not y_val_cat.empty and not y_pred_val_final.empty:
    # Ensure labels in classification report and confusion matrix cover all
    ↪ possible classes
    unique_labels_present_val = np.union1d(y_val_cat.unique(),
    ↪ y_pred_val_final.unique())
    # Ensure all labels from 0 to num_classes-1 are considered for
    ↪ consistency if some classes have 0 instances
    report_labels_val = list(range(num_classes))

    accuracy_val = accuracy_score(y_val_cat, y_pred_val_final)
    kappa_val = cohen_kappa_score(y_val_cat, y_pred_val_final,
    ↪ labels=report_labels_val if report_labels_val else None)
    report_val = classification_report(y_val_cat, y_pred_val_final,
    ↪ labels=report_labels_val, target_names=bin_labels_for_plots, zero_division=0)

    print("\n--- Validation Set Metrics (XGBoost - Categorical) ---")
    print(f"Accuracy: {accuracy_val:.3f}")
    print(f"Cohen's Kappa: {kappa_val:.3f}")
    print("Classification Report (Validation):\n", report_val)
else:
    print("\nValidation data (actual or predicted) is empty. Skipping
    ↪ validation metrics.")

# Test Set Metrics (if y_test_cat and y_pred_test_final are not empty)
if not y_test_cat.empty and not y_pred_test_final.empty:
    unique_labels_present_test = np.union1d(y_test_cat.unique(),
    ↪ y_pred_test_final.unique())
    report_labels_test = list(range(num_classes))

    accuracy_test = accuracy_score(y_test_cat, y_pred_test_final)
    kappa_test = cohen_kappa_score(y_test_cat, y_pred_test_final,
    ↪ labels=report_labels_test if report_labels_test else None)
    report_test = classification_report(y_test_cat, y_pred_test_final,
    ↪ labels=report_labels_test, target_names=bin_labels_for_plots,
    ↪ zero_division=0)
    cm_test = confusion_matrix(y_test_cat, y_pred_test_final,
    ↪ labels=report_labels_test)

    print("\n--- Test Set Metrics (XGBoost - Categorical) ---")
    print(f"Accuracy: {accuracy_test:.3f}")
    print(f"Cohen's Kappa: {kappa_test:.3f}")

```

```

print("Classification Report (Test):\n", report_test)
# print("Confusion Matrix (Test):\n", cm_test)

# --- 2. Visualize Results (Test Set) ---
# a) Confusion Matrix Heatmap
if num_classes > 0:
    plt.figure(figsize=(min(10, num_classes + 2), min(8, num_classes + 1)))
    sns.heatmap(cm_test, annot=True, fmt='d', cmap='Blues',
                xticklabels=bin_labels_for_plots,
                yticklabels=bin_labels_for_plots)
    plt.xlabel("Predicted Category")
    plt.ylabel("Actual Category")
    plt.title("Confusion Matrix (Test Set - XGBoost)")
    plt.tight_layout()
    plt.show()

    # b) Actual vs. Predicted Class Distribution
    plt.figure(figsize=(max(8, num_classes * 0.8), 6))
    actual_counts = y_test_cat.value_counts()
    ↪reindex(range(num_classes), fill_value=0).sort_index()
    predicted_counts = y_pred_test_final.value_counts()
    ↪reindex(range(num_classes), fill_value=0).sort_index()

    df_counts = pd.DataFrame({'Actual': actual_counts, 'Predicted': ↪
    ↪predicted_counts})
    if not df_counts.empty:
        # Use the generated bin_labels_for_plots for the x-axis ticks
        df_counts.index = [bin_labels_for_plots[i] for i in df_counts.
    ↪index if i < len(bin_labels_for_plots)]

        df_counts.plot(kind='bar', width=0.8)
        plt.title("Actual vs. Predicted Class Distribution (Test Set)")
        plt.xlabel("ICU LOS Category")
        plt.ylabel("Number of Cases")
        plt.xticks(rotation=45, ha="right")
        plt.legend()
        plt.grid(axis='y', linestyle='--')
        plt.tight_layout()
        plt.show()
    else:
        print("Cannot plot class distribution: Counts data is empty.")
else:
    print("Number of classes is 0, cannot generate plots.")
else:
    print("\nTest data (actual or predicted) is empty. Skipping test_
    ↪metrics and visualization.")

```

```
# --- Record End Time ---
end_time = time.time()
execution_time = end_time - start_time
print(f"\nExecution time: {execution_time:.2f} seconds")
```

--- Block 4: Evaluating and Visualizing Results ---

Using bin labels for plots: ['[1.0-3.0)', '[3.0-5.0)', '[5.0-7.0)', '[7.0-9.0)', '[9.0-11.0)', '[11.0-13.0)', '[13.0-15.0)', '[15.0-17.0)', '[17.0-19.0)', '[19.0-21.0)', '[21.0-23.0)', '[23.0-25.0)']

--- Validation Set Metrics (XGBoost - Categorical) ---

Accuracy: 0.429

Cohen's Kappa: 0.071

Classification Report (Validation):

	precision	recall	f1-score	support
[1.0-3.0)	0.62	0.71	0.66	2473
[3.0-5.0)	0.21	0.13	0.16	828
[5.0-7.0)	0.07	0.03	0.04	387
[7.0-9.0)	0.05	0.03	0.04	216
[9.0-11.0)	0.04	0.03	0.03	148
[11.0-13.0)	0.04	0.03	0.03	105
[13.0-15.0)	0.05	0.09	0.06	70
[15.0-17.0)	0.06	0.07	0.06	75
[17.0-19.0)	0.03	0.07	0.04	44
[19.0-21.0)	0.02	0.08	0.03	38
[21.0-23.0)	0.03	0.06	0.04	34
[23.0-25.0)	0.02	0.08	0.03	26
accuracy			0.43	4444
macro avg	0.10	0.12	0.10	4444
weighted avg	0.40	0.43	0.41	4444

--- Test Set Metrics (XGBoost - Categorical) ---


Accuracy: 0.420

Cohen's Kappa: 0.073


Classification Report (Test):

	precision	recall	f1-score	support
[1.0-3.0)	0.60	0.70	0.64	2414
[3.0-5.0)	0.20	0.14	0.16	843
[5.0-7.0)	0.13	0.05	0.07	402
[7.0-9.0)	0.11	0.06	0.08	231

[9.0-11.0)	0.02	0.01	0.02	147
[11.0-13.0)	0.06	0.05	0.05	97
[13.0-15.0)	0.08	0.11	0.09	83
[15.0-17.0)	0.06	0.08	0.07	63
[17.0-19.0)	0.05	0.08	0.06	62
[19.0-21.0)	0.01	0.02	0.01	43
[21.0-23.0)	0.01	0.03	0.02	31
[23.0-25.0)	0.05	0.18	0.08	28
accuracy			0.42	4444
macro avg	0.11	0.13	0.11	4444
weighted avg	0.39	0.42	0.40	4444



<Figure size 960x600 with 0 Axes>



Execution time: 0.74 seconds

7 Final Conclusions

7.0.1 Features Distribution

We were able to get some interesting conclusions from the distribution analysis for each feature, namely: - Most of the Admissions in ICU Stay occur within the first hour in the hospital - Most of patients only stay in ICU for **5 days or less**, with a decreasing trend of the number of patients as the number of days increases. - There is clearly an outlier (7 AM) in the distribution of patients by Hour of Admission

7.0.2 Features Correlation

We were also able to draw interesting conclusions from the correlation between different features, namely: - The correlation between **Diagnosis** and **ICU Length of Stay** is evident, with Twin mate, liveborn, delivered in hospital, with cesarean section having a much higher mean value compared to other types of Diagnosis - Regarding the connection between **Age** and **Length of ICU Stay**, we can conclude that the vast majority of ICU stays are relatively short across all adult age groups, although the spread of the values of **Length of ICU Stay** appears to widen slightly for older patients - All **Insurance** systems have a similar distribution and mean

for ICU Length of Stay, with a clear exception of Self Pay, whose average value is significantly lower. - The EMERGENCY Type of Admission is very correlated with Hour of Admission, being much more common in the evening/night 16PM - 23PM. The Type of Admission Elective is also correlated with Hour of Admission, being very common at 7AM. - Patients who have EMERGENCY as their Type of Admission tend to EXPIRE(die) much more frequently. - Some interesting outliers in the distribution of the correlation between Diagnosis and Hour of Admission are 41401 (Coronary atherosclerosis of native coronary artery) and 4241 (Aortic valve disorders) with 7 AM. - NICU (Neonatal Intensive Care Unit) has the lowest number of deaths and MICU (Medical Intensive Care Unit) has the highest number of deaths.

7.0.3 Predictor Model

- The accuracy value in the Validation and in Testing set are similar and have medium-low values, which mean the model wasn't very successful at predicting the patients ICU Length of Stay during the validation phase
- The Cohen's Kappa value for both Validation and Testing is very low, which indicates the model is not much better than a random chance model at predicting the correct classes
- The heatmap shows that the model is very good at predicting correctly the majority class [1-3] days, but mostly fails at predicting other classes correctly
- The bar plot shows that the model predictions distribution is pretty similar to the actual data distribution, which is a good sign for possible future improvements
- In summary, the predictions for the majority class are accurate but for minority classes are not, which is a very common problem in Machine Learning prediction models.
- In future works, we could try to improve these bad results by performing better feature engineering, performing fine-tuning in the parameters of the predictor model, choose more wisely the features to use and possibly choose a more adequate model for the dataset.

7.0.4 BigQuery and Execution Time

- The BigQuery proved to be a reliable and intuitive platform to perform Machine Learning projects
- The queries that were timed (not all were because all the results would be very similar) executed almost instantly (range from < 1 second to ~3/4 seconds)
- All the steps for executing the XGBoost Classifier model were also (and surprisingly!) very fast, with all executions running almost instantly

7.0.5 Project Developed by:

- Francisco Macieira → up202207166@edu.fc.up.pt
- Manuel Silva → up202108874@edu.fe.up.pt
- Nuno Gomes → up202206195@edu.fc.up.pt