

Dask_Pipeline

July 1, 2025

1 Configuration and Imports

```
[2]: from google.colab import auth
auth.authenticate_user()
```

```
[62]: import dask.dataframe as dd
from dask_ml.preprocessing import StandardScaler, OneHotEncoder
from dask_ml.model_selection import train_test_split
from google.cloud import bigquery
import bigframes as bf
import bigframes.pandas as bpd
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import numpy as np
import time
from sklearn.utils.class_weight import compute_sample_weight
from sklearn.utils.multiclass import type_of_target
from sklearn.preprocessing import StandardScaler, OneHotEncoder,
    ↳FunctionTransformer
from sklearn.feature_extraction import FeatureHasher
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.metrics import accuracy_score, classification_report,
    ↳confusion_matrix, cohen_kappa_score
from sklearn.base import BaseEstimator, TransformerMixin, clone
from collections import Counter
from xgboost import XGBClassifier
```

```
[4]: project_id = "my-first-gcp-project-452814"
client = bigquery.Client(project=project_id)

print("Project ID:", client.project)
```

Project ID: my-first-gcp-project-452814

2 Datasets

The datasets used in these project were the ones we considered most relevant for the task at hand. They were uploaded to the BigQuery project `my-first-gcp-project-452814` via Google Cloud Storage.

Below we perform some queries to vizualize and ensure data quality in all of these the datasets.

2.1 Chartevents Reduced

Contains time-stamped clinical data and measurements recorded for patients during their hospital stay.

- `ROW_ID`: Unique identifier for the row.
- `SUBJECT_ID`: Foreign key to the `PATIENTS` table.
- `HADM_ID`: Foreign key to the `ADMISSIONS` table.
- `ITEMID`: Foreign key to the `D_ITEMS` table, indicating the type of measurement.
- `CHARTTIME`: Timestamp when the measurement was recorded.
- `STORETIME`: Timestamp when the measurement was stored.
- `VALUE`: The value of the measurement (can be numeric or text).
- `VALUENUM`: Numeric value of the measurement, if applicable.
- `VALUEUOM`: Unit of measurement.
- `WARNING`: Flags indicating potential issues with the data.

2.1.1 Reducing the original Chartevents dataset

We decided to create a more small and acessible table using the original `chartevents`. This new table, `chartevents_reduced`, that will be used from now on, only contains the measurements of the most common disease in the original table.

```
[ ]: start_time = time.time()

query = """
-- Step 1: Create a new table for disease-related data
CREATE OR REPLACE TABLE `my-first-gcp-project-452814.cdle_project_dataset.
↳chartevents_reduced` AS

-- Step 2: Identify and select only disease-related measurements
WITH disease_measurements AS (
  SELECT *
  FROM `my-first-gcp-project-452814.cdle_project_dataset.chartevents`
  WHERE ITEMID IN (
    -- Cardiovascular
    220045, -- Heart Rate
    220050, -- Blood Pressure Systolic
    220051, -- Blood Pressure Diastolic

    -- Metabolic/Endocrine
    220179, -- Glucose
```

```

50912,    -- Creatinine
50809,    -- Glucose (serum)

-- Respiratory
220277,   -- SpO2
224690,   -- Respiratory Rate

-- Infection/Inflammation
50813,    -- Lactate (sepsis marker)

-- Liver
50821     -- Bilirubin
)
),

-- Step 3: Find the top 3 most common disease measurements
top_disease_measurements AS (
    SELECT
        ITEMID,
        COUNT(*) AS measurement_count
    FROM disease_measurements
    GROUP BY ITEMID
    ORDER BY measurement_count DESC
    LIMIT 3
)

-- Step 4: Create final table with only top disease measurements
SELECT d.*
FROM disease_measurements d
JOIN top_disease_measurements t ON d.ITEMID = t.ITEMID;
"""

query_job = client.query(query)
print("Dataset reduced sucessfully")

end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")

```

Dataset reduced sucessfully
Query Execution Time: 0.70 seconds

2.1.2 Visualization

```

[ ]: start_time = time.time()

dataset_id = "cdle_project_dataset"

```

```

table_id = "chartevents_reduced"

table_ref = client.dataset(dataset_id).table(table_id)

chartevents = client.list_rows(table_ref).to_dataframe()

end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")

chartevents.head()

```

Query Execution Time: 7.32 seconds

```

[ ]:
  ROW_ID  SUBJECT_ID  HADM_ID  ICUSTAY_ID  ITEMID  \
0  19184847      66298   152072    200105  220045
1  24904858      79894   106711    201234  220045
2    759709       9002   120994    203667  220045
3  30895546      91558   119749    204020  220045
4  34071445      99469   179324    204842  220179

      CHARTTIME      STORETIME  CGID  VALUE  VALUENUM  \
0  2104-10-24 00:05:00+00:00      NaT  <NA>      0      0.0
1  2104-05-01 02:00:00+00:00  2104-05-01 02:10:00+00:00  16037      0      0.0
2  2177-05-05 01:40:00+00:00  2177-05-05 01:48:00+00:00  18784      0      0.0
3  2183-09-01 09:10:00+00:00  2183-09-01 09:18:00+00:00  16526      0      0.0
4  2183-05-13 16:20:00+00:00  2183-05-13 17:10:00+00:00  19589      0      0.0

  VALUEUOM  WARNING  ERROR  RESULTSTATUS  STOPPED
0      bpm         0      0          None      None
1      bpm         0      0          None      None
2      bpm         0      0          None      None
3      bpm         0      0          None      None
4     mmHg         0      0          None      None

```

2.1.3 Data quality check

```

[ ]: start_time = time.time()

query = """
-- Data Quality Assessment for chartevents_reduced
WITH stats AS (
  SELECT
    COUNT(*) AS total_records,
    COUNT(DISTINCT SUBJECT_ID) AS unique_patients,
    COUNT(DISTINCT HADM_ID) AS unique_admissions,
    COUNT(DISTINCT ITEMID) AS unique_measurement_types

```

```

FROM `my-first-gcp-project-452814.cdle_project_dataset.chartevents_reduced`
),

measurement_analysis AS (
  SELECT
    ITEMID,
    COUNT(*) AS record_count,
    ROUND(COUNT(*)*100/(SELECT total_records FROM stats), 2) AS percentage_of_total,
    MIN(VALUENUM) AS min_value,
    MAX(VALUENUM) AS max_value,
    AVG(VALUENUM) AS avg_value,
    COUNT(CASE WHEN VALUENUM IS NULL THEN 1 END) AS null_value_counts,
    COUNT(CASE WHEN VALUE = '' THEN 1 END) AS empty_string_counts,
    MIN(CHARTTIME) AS earliest_measurement,
    MAX(CHARTTIME) AS latest_measurement
  FROM `my-first-gcp-project-452814.cdle_project_dataset.chartevents_reduced`
  GROUP BY ITEMID
),

temporal_analysis AS (
  SELECT
    EXTRACT(YEAR FROM CHARTTIME) AS year,
    EXTRACT(MONTH FROM CHARTTIME) AS month,
    COUNT(*) AS measurements_count
  FROM `my-first-gcp-project-452814.cdle_project_dataset.chartevents_reduced`
  GROUP BY year, month
  ORDER BY year, month
)

SELECT
  -- Basic Statistics
  s.total_records,
  s.unique_patients,
  s.unique_admissions,
  s.unique_measurement_types,

  -- Measurement-specific quality metrics
  ARRAY(
    SELECT AS STRUCT * FROM measurement_analysis
    ORDER BY record_count DESC
  ) AS measurement_quality,

  -- Temporal distribution
  (SELECT COUNT(*) FROM temporal_analysis) AS months_with_data,
  (SELECT MIN(year) FROM temporal_analysis) AS first_year,
  (SELECT MAX(year) FROM temporal_analysis) AS last_year,

```

```

-- Data completeness
(SELECT COUNT(*) FROM `my-first-gcp-project-452814.cdle_project_dataset.
↳chartevents_reduced` WHERE CHARTTIME IS NULL) AS null_timestamps,
(SELECT COUNT(*) FROM `my-first-gcp-project-452814.cdle_project_dataset.
↳chartevents_reduced` WHERE SUBJECT_ID IS NULL) AS null_patient_ids,

-- Clinical validity checks
(SELECT COUNT(*) FROM `my-first-gcp-project-452814.cdle_project_dataset.
↳chartevents_reduced`
WHERE ITEMID = 220045 AND (VALUENUM < 20 OR VALUENUM > 250)) AS
↳abnormal_heart_rates,

(SELECT COUNT(*) FROM `my-first-gcp-project-452814.cdle_project_dataset.
↳chartevents_reduced`
WHERE ITEMID = 220050 AND (VALUENUM < 50 OR VALUENUM > 300)) AS
↳abnormal_bp_readings
FROM stats s
"""

query_job = client.query(query)
data_quality = query_job.to_dataframe()

end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")

data_quality

```

Query Execution Time: 3.49 seconds

```

[ ]:  total_records  unique_patients  unique_admissions  \
0          6724529          17717          21927

      unique_measurement_types  \
0                          3

                                measurement_quality  months_with_data  \
0  [{'ITEMID': 220045, 'record_count': 2762225, '...'          1289

      first_year  last_year  null_timestamps  null_patient_ids  \
0          2100          2209              0              0

      abnormal_heart_rates  abnormal_bp_readings
0              1177              0

```

2.2 Admissions

Contains information about patient admissions to the hospital, such as admission and discharge times, admission type, and insurance information.

- ROW_ID: Unique identifier for each row
- SUBJECT_ID: Foreign key to the PATIENTS table.
- HADM_ID: Unique identifier for the hospital admission
- ADMITTIME: Timestamp for hospital admission.
- DISCHTIME: Timestamp for hospital discharge.
- DEATHTIME: Timestamp for patient death, if applicable.
- ADMISSION_TYPE: Type of admission, e.g., EMERGENCY, ELECTIVE, URGENT.
- ADMISSION_LOCATION: Location of the patient prior to admission.
- DISCHARGE_LOCATION: Location to which the patient was discharged.
- INSURANCE: The patient's insurance provider.
- LANGUAGE: The patient's primary language.
- RELIGION: The patient's religious affiliation.
- MARITAL_STATUS: The patient's marital status.
- ETHNICITY: The patient's ethnicity.
- EDREGTIME: Emergency Department registration time
- EDOUTTIME: Emergency Department departure time
- DIAGNOSIS: The patient's primary diagnosis.
- HOSPITAL_EXPIRE_FLAG: Indicates if the patient died in the hospital.
- HAS_CHARTEVENTS_DATA: Flag indicating if there is chart event data.

2.2.1 Visualization

```
[ ]: dataset_id = "cdle_project_dataset"
    table_id = "Admissions"

    table_ref = client.dataset(dataset_id).table(table_id)

    admissions = client.list_rows(table_ref).to_dataframe()

    admissions.head()
```

```
[ ]:  ROW_ID  SUBJECT_ID  HADM_ID  ADMITTIME  \
0      4060         3369   126808  2111-01-24  11:53:00+00:00
1      50952        74869   123152  2150-09-05  17:49:00+00:00
2      12812        10484   113233  2190-09-18  22:39:00+00:00
3      13573        11091   164694  2101-06-07  13:57:00+00:00
4      33654        27527   155091  2131-08-27  18:01:00+00:00

      DISCHTIME  DEATHTIME  ADMISSION_TYPE  \
0  2111-01-25  22:40:00+00:00  2111-01-25  22:40:00+00:00  EMERGENCY
1  2150-09-12  18:30:00+00:00  2150-09-12  18:30:00+00:00  EMERGENCY
2  2190-09-24  20:40:00+00:00  2190-09-24  20:40:00+00:00  EMERGENCY
3  2101-09-18  07:20:00+00:00  2101-09-18  07:20:00+00:00  EMERGENCY
```

4 2131-10-03 05:30:00+00:00 2131-10-03 05:30:00+00:00 EMERGENCY

	ADMISSION_LOCATION	DISCHARGE_LOCATION	INSURANCE	LANGUAGE	\
0	PHYS REFERRAL/NORMAL DELI	DEAD/EXPIRED	Private	None	
1	TRANSFER FROM HOSP/EXTRAM	DEAD/EXPIRED	Medicare	ENGL	
2	TRANSFER FROM HOSP/EXTRAM	DEAD/EXPIRED	Medicaid	None	
3	CLINIC REFERRAL/PREMATURE	DEAD/EXPIRED	Private	None	
4	CLINIC REFERRAL/PREMATURE	DEAD/EXPIRED	Private	PTUN	

	RELIGION	MARITAL_STATUS	ETHNICITY	EDREGTIME	EDOUTTIME	\
0	UNOBTAINABLE	SINGLE	WHITE	NaT	NaT	
1	CATHOLIC	WIDOWED	ASIAN - JAPANESE	NaT	NaT	
2	CATHOLIC	MARRIED	WHITE	NaT	NaT	
3	CATHOLIC	SEPARATED	WHITE	NaT	NaT	
4	NOT SPECIFIED	MARRIED	WHITE	NaT	NaT	

	DIAGNOSIS	HOSPITAL_EXPIRE_FLAG	\
0	None	1	
1	? SEROTONIN SYNDROME	1	
2	(AML) ACUTE MYELOGENOUS LEUKEMIA	1	
3	(AML) ACUTE MYELOGENOUS LEUKEMIA	1	
4	(AML) ACUTE MYELOGENOUS LEUKEMIA	1	

	HAS_CHARTEVENTS_DATA
0	1
1	1
2	1
3	1
4	1

2.2.2 Data quality check

```
[ ]: query = ""
-- Data Quality Assessment for admissions
SELECT
    COUNT(*) AS total_records,
    COUNT(DISTINCT HADM_ID) AS unique_admissions,
    COUNT(DISTINCT SUBJECT_ID) AS unique_patients,
    COUNT(CASE WHEN ADMITTIME IS NULL THEN 1 END) AS null_admit_times,
    COUNT(CASE WHEN DISCHTIME IS NULL THEN 1 END) AS null_discharge_times,
    COUNT(CASE WHEN ADMISSION_TYPE IS NULL THEN 1 END) AS null_admission_types,
    COUNT(CASE WHEN ADMISSION_LOCATION IS NULL THEN 1 END) AS_
↪null_admission_locations,
    COUNT(CASE WHEN INSURANCE IS NULL THEN 1 END) AS null_insurance_info,
    MIN(ADMITTIME) AS first_admission,
    MAX(ADMITTIME) AS last_admission
FROM `my-first-gcp-project-452814.cdle_project_dataset.Admissions`;
```



```
"""
```

```
query_job = client.query(query)
data_quality = query_job.to_dataframe()
data_quality
```

```
[ ]:  total_records  unique_admissions  unique_patients  null_admit_times  \
0          58976          58976          46520          0

      null_discharge_times  null_admission_types  null_admission_locations  \
0              0              0              0              0

      null_insurance_info          first_admission          last_admission
0              0  2100-06-07 19:59:00+00:00  2210-08-17 17:13:00+00:00
```

2.3 Callout

Contains information about requests for services or consultations for patients.

- ROW_ID: Unique identifier for the row.
- SUBJECT_ID: Foreign key to the PATIENTS table.
- HADM_ID: Foreign key to the ADMISSIONS table.
- CALLOUT_ID: Unique identifier for the callout request.
- CALLOUTTIME: Timestamp for the callout request.
- SERVICE_ID: ID of the service requested.
- LOCATION: Location of the patient when the callout was placed.
- STATUS: Status of the callout request.
- OUTCOME: Outcome of the callout request.
- ACKNOWLEDGE_TIME: Timestamp when the callout was acknowledged.
- OUTCOMETIME: Timestamp when the callout outcome was recorded.
- FIRSTRESERVATIONTIME: of the first reservation.
- CURRENTRESERVATIONTIME: Timestamp of the current reservation.
- CREATETIME: Timestamp when the row was created.
- UPDATETIME: Timestamp when the row was updated.
- CALLOUT_WARDID: Ward ID of the callout.
- CALLOUT_SERVICEREQUEST: Service requested.
- CALLOUT_TELEPHONE: Telephone number for the callout.
- REQUEST_TELE: Telephone request.
- REQUEST_RESP: Respiratory reqTimestamp.
- REQUEST_CDIF: C. difficile request.
- REQUEST_MRSA: MRSA request.
- REQUEST_VRE: VRE request.
- DISCHARGE_WARDID: Discharge ward ID.
- ACKNOWLEDGE_STATUS: Acknowledge status.

2.3.1 Visualization

```
[ ]: dataset_id = "cdle_project_dataset"
table_id = "Callout"

table_ref = client.dataset(dataset_id).table(table_id)

callout = client.list_rows(table_ref).to_dataframe()

callout.head()
```

```
[ ]:  ROW_ID  SUBJECT_ID  HADM_ID  SUBMIT_WARDID  SUBMIT_CAREUNIT  CURR_WARDID  \
0    15115      31974   144780          <NA>          None          <NA>
1      161        309   162308           7          None           2
2      169        333   160548           7          None           2
3      197        383   173723           7          None           2
4      136        253   176189           7          None           2
```

```
  CURR_CAREUNIT  CALLOUT_WARDID  CALLOUT_SERVICE  REQUEST_TELE  ...  \
0          None                1             MED            1  ...
1          CCU                 2             CCU            1  ...
2          CCU                 2             CCU            1  ...
3          CCU                 2             CCU            1  ...
4          CCU                 2             CCU            1  ...
```

```
  CALLOUT_STATUS  CALLOUT_OUTCOME  DISCHARGE_WARDID  ACKNOWLEDGE_STATUS  \
0      Inactive      Discharged           0      Unacknowledged
1      Inactive      Discharged           2      Acknowledged
2      Inactive      Discharged           2      Acknowledged
3      Inactive      Discharged           2      Unacknowledged
4      Inactive      Discharged           2      Acknowledged
```

```
  CREATETIME  UPDATETIME  \
0 2191-01-26 13:55:10+00:00 2191-01-26 13:55:10+00:00
1 2160-06-05 10:22:04+00:00 2160-06-05 10:22:04+00:00
2 2137-09-30 09:42:12+00:00 2137-09-30 09:42:12+00:00
3 2143-09-08 10:53:04+00:00 2143-09-08 10:53:04+00:00
4 2174-01-23 09:57:24+00:00 2174-01-23 10:44:12+00:00
```

```
  ACKNOWLEDGETIME  OUTCOMETIME  FIRSTRESERVATIONTIME  \
0              NaT 2191-01-26 14:10:04+00:00          NaT
1 2160-06-05 11:20:06+00:00 2160-06-05 19:25:01+00:00          NaT
2 2137-09-30 09:45:08+00:00 2137-10-01 14:40:02+00:00          NaT
3              NaT 2143-09-08 11:55:02+00:00          NaT
4 2174-01-23 11:10:50+00:00 2174-01-23 13:40:02+00:00          NaT
```

```
CURRENTRESERVATIONTIME
```

0	NaT
1	NaT
2	NaT
3	NaT
4	NaT

[5 rows x 24 columns]

2.3.2 Data quality check

```
[ ]: query = """
-- Data Quality Assessment for callout
WITH basic_stats AS (
    SELECT
        COUNT(*) AS total_records,
        COUNT(DISTINCT SUBJECT_ID) AS unique_patients,
        COUNT(DISTINCT HADM_ID) AS unique_admissions,
        COUNT(DISTINCT ROW_ID) AS unique_row_ids,
        COUNT(DISTINCT SUBMIT_WARDID) AS unique_submit_ward_ids,
        COUNT(DISTINCT SUBMIT_CAREUNIT) AS unique_submit_care_units,
        COUNT(DISTINCT CURR_WARDID) AS unique_current_ward_ids,
        COUNT(DISTINCT CURR_CAREUNIT) AS unique_current_care_units,
        COUNT(DISTINCT CALLOUT_WARDID) AS unique_callout_ward_ids,
        COUNT(DISTINCT CALLOUT_SERVICE) AS unique_callout_service,
        COUNT(DISTINCT REQUEST_TELE) AS unique_request_telephones,
        COUNT(DISTINCT REQUEST_RESP) AS unique_request_resp,
        COUNT(DISTINCT REQUEST_CDIFF) AS unique_request_cdif,
        COUNT(DISTINCT REQUEST_MRSA) AS unique_request_mrsa,
        COUNT(DISTINCT REQUEST_VRE) AS unique_request_vre,
        COUNT(DISTINCT CALLOUT_STATUS) AS unique_callout_statuses,
        COUNT(DISTINCT CALLOUT_OUTCOME) AS unique_callout_outcomes,
        COUNT(DISTINCT DISCHARGE_WARDID) AS unique_discharge_ward_ids,
        COUNT(DISTINCT ACKNOWLEDGE_STATUS) AS unique_acknowledge_statuses
    FROM `my-first-gcp-project-452814.cdle_project_dataset.Callout`
),
completeness AS (
    SELECT
        COUNT(CASE WHEN SUBJECT_ID IS NULL THEN 1 END) AS null_subject_ids,
        COUNT(CASE WHEN HADM_ID IS NULL THEN 1 END) AS null_hadm_ids,
        COUNT(CASE WHEN SUBMIT_WARDID IS NULL THEN 1 END) AS_
↪null_submit_ward_ids,
        COUNT(CASE WHEN SUBMIT_CAREUNIT IS NULL THEN 1 END) AS_
↪null_submit_care_units,
        COUNT(CASE WHEN SUBMIT_CAREUNIT = '' THEN 1 END) AS_
↪empty_submit_care_units,
        COUNT(CASE WHEN CURR_WARDID IS NULL THEN 1 END) AS null_curr_ward_ids,
```

```

        COUNT(CASE WHEN CURR_CAREUNIT IS NULL THEN 1 END) AS_
↪null_curr_care_units,
        COUNT(CASE WHEN CURR_CAREUNIT = '' THEN 1 END) AS empty_curr_care_units,
        COUNT(CASE WHEN CALLOUT_WARDID IS NULL THEN 1 END) AS_
↪null_callout_ward_ids,
        COUNT(CASE WHEN CALLOUT_SERVICE IS NULL THEN 1 END) AS_
↪null_callout_service,
        COUNT(CASE WHEN CALLOUT_SERVICE = '' THEN 1 END) AS_
↪empty_callout_service,
        COUNT(CASE WHEN REQUEST_TELE IS NULL THEN 1 END) AS_
↪null_request_telephones,
        COUNT(CASE WHEN REQUEST_RESP IS NULL THEN 1 END) AS null_request_resp,
        COUNT(CASE WHEN REQUEST_CDIF IS NULL THEN 1 END) AS null_request_cdif,
        COUNT(CASE WHEN REQUEST_MRSA IS NULL THEN 1 END) AS null_request_mrsa,
        COUNT(CASE WHEN REQUEST_VRE IS NULL THEN 1 END) AS null_request_vre,
        COUNT(CASE WHEN CALLOUT_STATUS IS NULL THEN 1 END) AS_
↪null_callout_statuses,
        COUNT(CASE WHEN CALLOUT_STATUS = '' THEN 1 END) AS_
↪empty_callout_statuses,
        COUNT(CASE WHEN CALLOUT_OUTCOME IS NULL THEN 1 END) AS_
↪null_callout_outcomes,
        COUNT(CASE WHEN CALLOUT_OUTCOME = '' THEN 1 END) AS_
↪empty_callout_outcomes,
        COUNT(CASE WHEN DISCHARGE_WARDID IS NULL THEN 1 END) AS_
↪null_discharge_ward_ids,
        COUNT(CASE WHEN ACKNOWLEDGE_STATUS IS NULL THEN 1 END) AS_
↪null_acknowledge_statuses,
        COUNT(CASE WHEN ACKNOWLEDGE_STATUS = '' THEN 1 END) AS_
↪empty_acknowledge_statuses,
        COUNT(CASE WHEN CREATETIME IS NULL THEN 1 END) AS null_createtimes,
        COUNT(CASE WHEN UPDATETIME IS NULL THEN 1 END) AS null_updatetimes,
        COUNT(CASE WHEN ACKNOWLEDGETIME IS NULL THEN 1 END) AS_
↪null_acknowledgetimes,
        COUNT(CASE WHEN OUTCOMETIME IS NULL THEN 1 END) AS null_outcometimes,
        COUNT(CASE WHEN FIRSTRESERVATIONTIME IS NULL THEN 1 END) AS_
↪null_firstreservationtimes,
        COUNT(CASE WHEN CURRENTRESERVATIONTIME IS NULL THEN 1 END) AS_
↪null_currentreservationtimes
    FROM `my-first-gcp-project-452814.cdle_project_dataset.Callout`
),
temporal_analysis AS (
    SELECT
        MIN(CREATETIME) AS first_creation_time,
        MAX(CREATETIME) AS last_creation_time,
        MIN(UPDATETIME) AS first_update_time,
        MAX(UPDATETIME) AS last_update_time,

```

```

        MIN(ACKNOWLEDGETIME) AS first_acknowledge_time,
        MAX(ACKNOWLEDGETIME) AS last_acknowledge_time,
        MIN(OUTCOMETIME) AS first_outcome_time,
        MAX(OUTCOMETIME) AS last_outcome_time,
        MIN(FIRSTRESERVATIONTIME) AS first_reservation_time,
        MAX(FIRSTRESERVATIONTIME) AS last_reservation_time,
        MIN(CURRENTRESERVATIONTIME) AS first_current_reservation_time,
        MAX(CURRENTRESERVATIONTIME) AS last_current_reservation_time
    FROM `my-first-gcp-project-452814.cdle_project_dataset.Callout`
)
SELECT
    bs.*,
    c.*,
    t.*
FROM basic_stats bs
CROSS JOIN completeness c
CROSS JOIN temporal_analysis t;
"""

query_job = client.query(query)
data_quality = query_job.to_dataframe()
data_quality

```

```

[ ]:   total_records  unique_patients  unique_admissions  unique_row_ids  \
0      34499          22871          28732          34499

      unique_submit_ward_ids  unique_submit_care_units  unique_current_ward_ids  \
0              9              5              35

      unique_current_care_units  unique_callout_ward_ids  unique_callout_service  \
0              5              37              21

...      first_update_time      last_update_time  \
0  ... 2100-06-08 12:58:29+00:00 2210-08-20 16:05:16+00:00

      first_acknowledge_time      last_acknowledge_time  \
0 2100-06-08 12:58:32+00:00 2210-08-20 16:05:27+00:00

      first_outcome_time      last_outcome_time  \
0 2100-06-08 15:10:26+00:00 2210-08-20 18:55:15+00:00

      first_reservation_time      last_reservation_time  \
0 2100-06-08 11:55:26+00:00 2210-08-20 16:25:16+00:00

      first_current_reservation_time      last_current_reservation_time
0      2100-08-09 14:42:25+00:00      2209-08-05 14:01:23+00:00

```

```
[1 rows x 60 columns]
```

2.4 ICD Diagnoses Description

Contains descriptions for ICD-9 diagnosis codes.

- ROW_ID: Unique identifier for the row.
- ICD9_CODE: The ICD-9 diagnosis code.
- SHORT_TITLE: Short description of the diagnosis.
- LONG_TITLE: Long description of the diagnosis.

2.4.1 Visualization

```
[ ]: dataset_id = "cdle_project_dataset"
      table_id = "D_ICD_Diagnoses"

      table_ref = client.dataset(dataset_id).table(table_id)

      d_icd_diagnoses = client.list_rows(table_ref).to_dataframe()

      d_icd_diagnoses.head()
```

```
[ ]:  ROW_ID  ICD9_CODE          SHORT_TITLE  \
0      5120      4957  "ventilation" pneumonit
1     11159     94416   1 deg burn back of hand
2     11157     94414  1 deg burn fingr w thumb
3      3658     36911   1 eye-sev/oth-blind NOS
4     12505     94811  10-19% bdy brn/10-19% 3d

                                LONG_TITLE
0                                "Ventilation" pneumonitis
1                Erythema [first degree] of back of hand
2  Erythema [first degree] of two or more digits ...
3  Better eye: severe vision impairment; lesser e...
4  Burn [any degree] involving 10-19 percent of b...
```

2.4.2 Data quality check

```
[ ]: query = """
      -- Data Quality Assessment for d_icd_diagnoses
      SELECT
          COUNT(*) AS total_records,
          COUNT(DISTINCT ROW_ID) AS unique_row_ids,
          COUNT(DISTINCT ICD9_CODE) AS unique_icd9_codes,
          COUNT(DISTINCT SHORT_TITLE) AS unique_short_titles,
          COUNT(DISTINCT LONG_TITLE) AS unique_long_titles,
```

```

-- Completeness checks
COUNT(CASE WHEN ROW_ID IS NULL THEN 1 END) AS null_row_ids,
COUNT(CASE WHEN ICD9_CODE IS NULL THEN 1 END) AS null_icd9_codes,
COUNT(CASE WHEN SHORT_TITLE IS NULL THEN 1 END) AS null_short_titles,
COUNT(CASE WHEN LONG_TITLE IS NULL THEN 1 END) AS null_long_titles,

COUNT(CASE WHEN ICD9_CODE = '' THEN 1 END) AS empty_icd9_codes,
COUNT(CASE WHEN SHORT_TITLE = '' THEN 1 END) AS empty_short_titles,
COUNT(CASE WHEN LONG_TITLE = '' THEN 1 END) AS empty_long_titles,

-- Basic Analysis of ICD9 Code Length
AVG(LENGTH(ICD9_CODE)) AS avg_icd9_code_length,
MIN(LENGTH(ICD9_CODE)) AS min_icd9_code_length,
MAX(LENGTH(ICD9_CODE)) AS max_icd9_code_length,
COUNT(CASE WHEN LENGTH(ICD9_CODE) NOT BETWEEN 3 AND 5 THEN 1 END) AS
potential_invalid_icd9_length_count,

-- Potential data inconsistencies
COUNT(CASE WHEN SHORT_TITLE LIKE '%NOS%' AND LONG_TITLE NOT LIKE '%not_
otherwise specified%' THEN 1 END) AS short_title_nos_long_title_mismatch,
COUNT(CASE WHEN SHORT_TITLE LIKE '%NEC%' AND LONG_TITLE NOT LIKE '%not_
elsewhere classified%' THEN 1 END) AS short_title_nec_long_title_mismatch

FROM
`my-first-gcp-project-452814.cdle_project_dataset.D_ICD_Diagnoses`;
"""
query_job = client.query(query)
data_quality = query_job.to_dataframe()
data_quality

```

```

[ ]:  total_records  unique_row_ids  unique_icd9_codes  unique_short_titles  \
0          14567          14567          14567          14328

      unique_long_titles  null_row_ids  null_icd9_codes  null_short_titles  \
0          14562          0          0          0

      null_long_titles  empty_icd9_codes  empty_short_titles  empty_long_titles  \
0          0          0          0          0

      avg_icd9_code_length  min_icd9_code_length  max_icd9_code_length  \
0          4.686483          3          5

      potential_invalid_icd9_length_count  short_title_nos_long_title_mismatch  \
0          0          2195

      short_title_nec_long_title_mismatch
0          2138

```

2.5 D Items

Contains metadata about different medical items (measurements, procedures, medications) recorded in the database.

- ROW_ID: Unique identifier for the row.
- ITEMID: Unique identifier for the item.
- LABEL: Label or name of the item.
- ABBREVIATION: Abbreviation for the item.
- DBSOURCE: Source database for the item.
- LINKSTO: Table linked to.
- CATEGORY: Category of the item.
- UNITNAME: Unit of measurement for the item.
- PARAM_TYPE: Parameter type.
- CONCEPTID: Identifier for the concept.

2.5.1 Visualization

```
[ ]: dataset_id = "cdle_project_dataset"
table_id = "D_Items"

table_ref = client.dataset(dataset_id).table(table_id)

d_items = client.list_rows(table_ref).to_dataframe()

d_items.head()
```

```
[ ]:  ROW_ID  ITEMID                                LABEL ABBREVIATION \
0      457    497  Patient controlled analgesia (PCA) [Inject]      None
1      458    498                                PCA Lockout (Min)      None
2      459    499                                PCA Medication      None
3      460    500                                PCA Total Dose      None
4      461    501                                PCV Exh Vt (Obser)      None

      DBSOURCE      LINKSTO  CATEGORY  UNITNAME  PARAM_TYPE  CONCEPTID
0  carevue  chartevents      None      None      None      None
1  carevue  chartevents      None      None      None      None
2  carevue  chartevents      None      None      None      None
3  carevue  chartevents      None      None      None      None
4  carevue  chartevents      None      None      None      None
```

2.5.2 Data quality check

```
[ ]: query = ""
-- Data Quality Assessment for d_items
SELECT
    COUNT(*) AS total_records,
    COUNT(DISTINCT ROW_ID) AS unique_row_ids,
```



```

COUNT(DISTINCT ITEMID) AS unique_itemids,
COUNT(DISTINCT LABEL) AS unique_labels,
COUNT(DISTINCT ABBREVIATION) AS unique_abbreviations,
COUNT(DISTINCT DBSOURCE) AS unique_dbsources,
COUNT(DISTINCT LINKSTO) AS unique_linkstos,
COUNT(DISTINCT CATEGORY) AS unique_categories,
COUNT(DISTINCT UNITNAME) AS unique_unitnames,
COUNT(DISTINCT PARAM_TYPE) AS unique_param_types,
COUNT(DISTINCT CONCEPTID) AS unique_conceptids,

-- Completeness checks (NULL values)
COUNT(CASE WHEN ROW_ID IS NULL THEN 1 END) AS null_row_ids,
COUNT(CASE WHEN ITEMID IS NULL THEN 1 END) AS null_itemids,
COUNT(CASE WHEN LABEL IS NULL THEN 1 END) AS null_labels,
COUNT(CASE WHEN ABBREVIATION IS NULL THEN 1 END) AS null_abbreviations,
COUNT(CASE WHEN DBSOURCE IS NULL THEN 1 END) AS null_dbsources,
COUNT(CASE WHEN LINKSTO IS NULL THEN 1 END) AS null_linkstos,
COUNT(CASE WHEN CATEGORY IS NULL THEN 1 END) AS null_categories,
COUNT(CASE WHEN UNITNAME IS NULL THEN 1 END) AS null_unitnames,
COUNT(CASE WHEN PARAM_TYPE IS NULL THEN 1 END) AS null_param_types,
COUNT(CASE WHEN CONCEPTID IS NULL THEN 1 END) AS null_conceptids,

-- Completeness checks (Empty strings)
COUNT(CASE WHEN LABEL = '' THEN 1 END) AS empty_labels,
COUNT(CASE WHEN ABBREVIATION = '' THEN 1 END) AS empty_abbreviations,
COUNT(CASE WHEN DBSOURCE = '' THEN 1 END) AS empty_dbsources,
COUNT(CASE WHEN LINKSTO = '' THEN 1 END) AS empty_linkstos,
COUNT(CASE WHEN CATEGORY = '' THEN 1 END) AS empty_categories,
COUNT(CASE WHEN UNITNAME = '' THEN 1 END) AS empty_unitnames,
COUNT(CASE WHEN PARAM_TYPE = '' THEN 1 END) AS empty_param_types,
COUNT(CASE WHEN CONCEPTID = '' THEN 1 END) AS empty_conceptids

FROM
    `my-first-gcp-project-452814.cdle_project_dataset.D_Items`;
"""
query_job = client.query(query)
data_quality = query_job.to_dataframe()
data_quality

```

```

[ ]:  total_records  unique_row_ids  unique_itemids  unique_labels  \
0          12487          12487          12487          11847

      unique_abbreviations  unique_dbsources  unique_linkstos  unique_categories  \
0                   2907                3                7                94

      unique_unitnames  unique_param_types  ...  null_param_types  \
0                   53                7  ...                9495

```

```

null_conceptids  empty_labels  empty_abbreviations  empty_dbsources  \
0                12487          0                0                0

empty_linkstos  empty_categories  empty_unitnames  empty_param_types  \
0                0                0                0                0

empty_conceptids
0                0

[1 rows x 29 columns]

```

2.6 ICD Diagnoses

Contains ICD-9 diagnosis codes assigned to patients during their hospital admissions.

- ROW_ID: Unique identifier for the row.
- SUBJECT_ID: Foreign key to the PATIENTS table.
- HADM_ID: Foreign key to the ADMISSIONS table.
- SEQ_NUM: Sequence number for the diagnosis within the admission.
- ICD9_CODE: The ICD-9 diagnosis code.

2.6.1 Visualization

```

[ ]: dataset_id = "cdle_project_dataset"
    table_id = "Diagnoses_ICD"

    table_ref = client.dataset(dataset_id).table(table_id)

    diagnoses_icd = client.list_rows(table_ref).to_dataframe()

    diagnoses_icd.head()

```

```

[ ]:
  ROW_ID  SUBJECT_ID  HADM_ID  SEQ_NUM  ICD9_CODE
0    1297         109   172335         1    40301
1    1298         109   172335         2     486
2    1299         109   172335         3    58281
3    1300         109   172335         4    5855
4    1301         109   172335         5    4254

```

2.6.2 Data quality check

```

[ ]: query = """
    -- Data Quality Assessment for diagnoses_icd
    SELECT
        COUNT(*) AS total_records,
        COUNT(DISTINCT ROW_ID) AS unique_row_ids,

```

```

COUNT(DISTINCT SUBJECT_ID) AS unique_patients,
COUNT(DISTINCT HADM_ID) AS unique_admissions,
COUNT(DISTINCT SEQ_NUM) AS unique_sequence_numbers,
COUNT(DISTINCT ICD9_CODE) AS unique_icd9_codes,

-- Completeness checks (NULL values)
COUNT(CASE WHEN ROW_ID IS NULL THEN 1 END) AS null_row_ids,
COUNT(CASE WHEN SUBJECT_ID IS NULL THEN 1 END) AS null_subject_ids,
COUNT(CASE WHEN HADM_ID IS NULL THEN 1 END) AS null_hadm_ids,
COUNT(CASE WHEN SEQ_NUM IS NULL THEN 1 END) AS null_sequence_numbers,
COUNT(CASE WHEN ICD9_CODE IS NULL THEN 1 END) AS null_icd9_codes,

-- Completeness checks (Empty strings)
COUNT(CASE WHEN ICD9_CODE = '' THEN 1 END) AS empty_icd9_codes,

-- Analysis of ICD9 Code Length
AVG(LENGTH(ICD9_CODE)) AS avg_icd9_code_length,
MIN(LENGTH(ICD9_CODE)) AS min_icd9_code_length,
MAX(LENGTH(ICD9_CODE)) AS max_icd9_code_length,
COUNT(CASE WHEN LENGTH(ICD9_CODE) NOT BETWEEN 3 AND 5 THEN 1 END) AS
potential_invalid_icd9_length_count

FROM
    `my-first-gcp-project-452814.cdle_project_dataset.Diagnoses_ICD`;
"""
query_job = client.query(query)
data_quality = query_job.to_dataframe()
data_quality

```

```

[ ]:  total_records  unique_row_ids  unique_patients  unique_admissions  \
0          651047          651047          46520          58976

      unique_sequence_numbers  unique_icd9_codes  null_row_ids  null_subject_ids  \
0                   39          6984          0          0

      null_hadm_ids  null_sequence_numbers  null_icd9_codes  empty_icd9_codes  \
0              0          47          47          0

      avg_icd9_code_length  min_icd9_code_length  max_icd9_code_length  \
0          4.448883          3          5

      potential_invalid_icd9_length_count
0          0

```

2.7 Icustays

Contains information about patient stays in the intensive care unit (ICU).

- ROW_ID: Unique identifier for the row.
- SUBJECT_ID: Foreign key to the PATIENTS table.
- HADM_ID: Foreign key to the ADMISSIONS table.
- ICUSTAY_ID: Unique identifier for the ICU stay.
- DBSOURCE: Source database.
- FIRST_CAREUNIT: First care unit the patient was in.
- LAST_CAREUNIT: Last care unit the patient was in.
- FIRST_WARDID: First ward ID.
- LAST_WARDID: Last ward ID.
- INTIME: Timestamp for ICU admission.
- OUTTIME: Timestamp for ICU discharge.
- LOS: Length of ICU stay.

2.7.1 Visualization

```
[ ]: dataset_id = "cdle_project_dataset"
table_id = "ICUSTays"

table_ref = client.dataset(dataset_id).table(table_id)

icustays = client.list_rows(table_ref).to_dataframe()

icustays.head()
```

```
[ ]:  ROW_ID  SUBJECT_ID  HADM_ID  ICUSTAY_ID  DBSOURCE  FIRST_CAREUNIT  \
0      372         275   129886    219649    carevue          CCU
1      389         291   113649    256641    carevue          CCU
2      390         291   125726    275109  metavision          CCU
3      394         294   152578    222074    carevue          CCU
4      401         301   160332    288401    carevue          CCU

      LAST_CAREUNIT  FIRST_WARDID  LAST_WARDID          INTIME  \
0              CCU              7              7  2170-10-07 11:28:53+00:00
1              CCU              7              7  2102-04-08 23:05:28+00:00
2              CCU              7              7  2106-04-17 12:26:17+00:00
3              CCU              7              7  2118-01-17 21:45:05+00:00
4              CCU              7              7  2189-11-11 12:12:33+00:00

          OUTTIME      LOS
0  2170-10-14 14:38:07+00:00  7.1314
1  2102-04-09 11:20:11+00:00  0.5102
2  2106-04-18 22:05:39+00:00  1.4023
3  2118-01-20 11:12:45+00:00  2.5609
4  2189-11-13 22:11:28+00:00  2.4159
```

2.7.2 Data quality check

```
[ ]: query = """
-- Data Quality Assessment for icustays
SELECT
    COUNT(*) AS total_records,
    COUNT(DISTINCT ROW_ID) AS unique_row_ids,
    COUNT(DISTINCT SUBJECT_ID) AS unique_patients,
    COUNT(DISTINCT HADM_ID) AS unique_admissions,
    COUNT(DISTINCT ICUSTAY_ID) AS unique_icustay_ids,
    COUNT(DISTINCT DBSOURCE) AS unique_dbsources,
    COUNT(DISTINCT FIRST_CAREUNIT) AS unique_first_careunits,
    COUNT(DISTINCT LAST_CAREUNIT) AS unique_last_careunits,
    COUNT(DISTINCT FIRST_WARDID) AS unique_first_wardids,
    COUNT(DISTINCT LAST_WARDID) AS unique_last_wardids,

    -- Completeness checks (NULL values)
    COUNT(CASE WHEN ROW_ID IS NULL THEN 1 END) AS null_row_ids,
    COUNT(CASE WHEN SUBJECT_ID IS NULL THEN 1 END) AS null_subject_ids,
    COUNT(CASE WHEN HADM_ID IS NULL THEN 1 END) AS null_hadm_ids,
    COUNT(CASE WHEN ICUSTAY_ID IS NULL THEN 1 END) AS null_icustay_ids,
    COUNT(CASE WHEN DBSOURCE IS NULL THEN 1 END) AS null_dbsources,
    COUNT(CASE WHEN FIRST_CAREUNIT IS NULL THEN 1 END) AS null_first_careunits,
    COUNT(CASE WHEN LAST_CAREUNIT IS NULL THEN 1 END) AS null_last_careunits,
    COUNT(CASE WHEN FIRST_WARDID IS NULL THEN 1 END) AS null_first_wardids,
    COUNT(CASE WHEN LAST_WARDID IS NULL THEN 1 END) AS null_last_wardids,
    COUNT(CASE WHEN INTIME IS NULL THEN 1 END) AS null_intimes,
    COUNT(CASE WHEN OUTTIME IS NULL THEN 1 END) AS null_outtimes,
    COUNT(CASE WHEN LOS IS NULL THEN 1 END) AS null_los,

    -- Completeness checks (Empty strings)
    COUNT(CASE WHEN FIRST_CAREUNIT = '' THEN 1 END) AS empty_first_careunits,
    COUNT(CASE WHEN LAST_CAREUNIT = '' THEN 1 END) AS empty_last_careunits,
    COUNT(CASE WHEN DBSOURCE = '' THEN 1 END) AS empty_dbsources,

    -- Basic Analysis of LOS (Length of ICU Stay)
    AVG(LOS) AS avg_los,
    MIN(LOS) AS min_los,
    MAX(LOS) AS max_los,
    COUNT(CASE WHEN LOS < 0 THEN 1 END) AS negative_los_count -- Potential data_
↪issue

FROM
    `my-first-gcp-project-452814.cdle_project_dataset.ICUStays`;
"""
query_job = client.query(query)
data_quality = query_job.to_dataframe()
```

```
data_quality
```

```
[ ]: total_records unique_row_ids unique_patients unique_admissions \
0          61532          61532          46476          57786

    unique_icustay_ids unique_dbsources unique_first_careunits \
0          61532              3              6

    unique_last_careunits unique_first_wardids unique_last_wardids ... \
0              6              16              17 ...

    null_intimes null_outtimes null_los empty_first_careunits \
0              0              10              10              0

    empty_last_careunits empty_dbsources avg_los min_los max_los \
0              0              0  4.917972  0.0001  173.0725

    negative_los_count
0              0

[1 rows x 29 columns]
```

2.8 Patients

Contains demographic information about the patients in the database.

- ROW_ID: Unique identifier for the row.
- SUBJECT_ID: Unique identifier for the patient.
- GENDER: Patient's gender.
- DOB: Patient's date of birth.
- DOD: Patient's date of death, if applicable.
- DOD_HOSP: Date of death in hospital.
- DOD_SSN: Date of death according to Social Security records.
- EXPIRE_FLAG: Flag indicating if the patient expired.

2.8.1 Visualization

```
[ ]: dataset_id = "cdle_project_dataset"
    table_id = "Patients"

    table_ref = client.dataset(dataset_id).table(table_id)

    patients = client.list_rows(table_ref).to_dataframe()

    patients.head()
```

```
[ ]:  ROW_ID  SUBJECT_ID  GENDER          DOB  DOD  DOD_HOSP  DOD_SSN  \
0      234      249      F  2075-03-13  00:00:00+00:00  NaT      NaT      NaT
1      238      253      F  2089-11-26  00:00:00+00:00  NaT      NaT      NaT
2      242      258      F  2124-09-19  00:00:00+00:00  NaT      NaT      NaT
3      243      260      F  2105-03-23  00:00:00+00:00  NaT      NaT      NaT
4      247      264      F  2162-11-30  00:00:00+00:00  NaT      NaT      NaT

      EXPIRE_FLAG
0          0
1          0
2          0
3          0
4          0
```

2.8.2 Data quality check

```
[ ]: query = ""
-- Data Quality Assessment for patients
SELECT
    COUNT(*) AS total_records,
    COUNT(DISTINCT ROW_ID) AS unique_row_ids,
    COUNT(DISTINCT SUBJECT_ID) AS unique_patients,
    COUNT(DISTINCT GENDER) AS unique_genders,
    COUNT(DISTINCT DOB) AS unique_dates_of_birth,
    COUNT(DISTINCT DOD) AS unique_dates_of_death,
    COUNT(DISTINCT DOD_HOSP) AS unique_dates_of_death_hospital,
    COUNT(DISTINCT DOD_SSN) AS unique_dates_of_death_ssn,
    COUNT(DISTINCT EXPIRE_FLAG) AS unique_expire_flags,

    -- Completeness checks (NULL values)
    COUNT(CASE WHEN ROW_ID IS NULL THEN 1 END) AS null_row_ids,
    COUNT(CASE WHEN SUBJECT_ID IS NULL THEN 1 END) AS null_subject_ids,
    COUNT(CASE WHEN GENDER IS NULL THEN 1 END) AS null_genders,
    COUNT(CASE WHEN DOB IS NULL THEN 1 END) AS null_dates_of_birth,
    COUNT(CASE WHEN DOD IS NULL THEN 1 END) AS null_dates_of_death,
    COUNT(CASE WHEN DOD_HOSP IS NULL THEN 1 END) AS null_dates_of_death_hospital,
    COUNT(CASE WHEN DOD_SSN IS NULL THEN 1 END) AS null_dates_of_death_ssn,
    COUNT(CASE WHEN EXPIRE_FLAG IS NULL THEN 1 END) AS null_expire_flags,

    -- Completeness checks (Empty strings)
    COUNT(CASE WHEN GENDER = '' THEN 1 END) AS empty_genders,

    -- Basic Analysis of Dates
    MIN(DOB) AS first_dob,
    MAX(DOB) AS last_dob,
    MIN(DOD) AS first_dod,
```

```

MAX(DOD) AS last_dod,

-- Potential Data Inconsistencies
COUNT(CASE WHEN DOD IS NOT NULL AND EXPIRE_FLAG = 0 THEN 1 END) AS
↳died_but_not_expired_flag,
COUNT(CASE WHEN DOD IS NULL AND EXPIRE_FLAG = 1 THEN 1 END) AS
↳expired_flag_but_no_dod

FROM
`my-first-gcp-project-452814.cdle_project_dataset.Patients`;
"""
query_job = client.query(query)
data_quality = query_job.to_dataframe()
data_quality

```

```

[ ]:   total_records  unique_row_ids  unique_patients  unique_genders  \
0           46520           46520           46520           2

      unique_dates_of_birth  unique_dates_of_death  \
0           32540           12911

      unique_dates_of_death_hospital  unique_dates_of_death_ssn  \
0           8747           11301

      unique_expire_flags  null_row_ids  ...  null_dates_of_death_hospital  \
0           2           0  ...           36546

      null_dates_of_death_ssn  null_expire_flags  empty_genders  \
0           33142           0           0

           first_dob           last_dob  \
0 1800-07-02 00:00:00+00:00 2201-07-24 00:00:00+00:00

           first_dod           last_dod  \
0 2100-06-19 00:00:00+00:00 2211-06-10 00:00:00+00:00

      died_but_not_expired_flag  expired_flag_but_no_dod
0           0           0

[1 rows x 24 columns]

```

3 Junction of tables

To make the pre-processing and Length of ICU Stay prediction more efficient we decided to create a table, using a query, that contain the most relevant columns of each of the datasets.

The most relevant columns chosen were: - Patient Info: SUBJECT_ID, GENDER, DOB, EXPIRE_FLAG,

AGE_AT_ADMISSION - Admission Info: HADM_ID, ADMITTIME, DISCHTIME, ADMISSION_TYPE, ADMISSION_LOCATION, INSURANCE, ETHNICITY, ADMISSION_DIAGNOSIS_TEXT, ADMISSION_HOUR - ICU Stay Info: ICUSTAY_ID, ICU_INTIME, ICU_OUTTIME, ICU_LOS (Target), FIRST_CAREUNIT - Diagnosis Info: PRIMARY_ICD9_CODE, PRIMARY_ICD9_TITLE - Callout Info: NUM_CALLOUTS

```
[ ]: query = ""
-- Creating a junction table with aggregated diagnoses, callouts and engineered
    ↳ features

CREATE OR REPLACE TABLE `my-first-gcp-project-452814.cdle_project_dataset.
    ↳ junction_table` AS

-- Define CTEs within the AS clause
WITH
  -- CTE to select only the primary diagnosis for each hospital admission
  PrimaryDiagnosis AS (
    SELECT
      di.HADM_ID,
      di.ICD9_CODE,
      dd.SHORT_TITLE AS ICD9_SHORT_TITLE
    FROM
      `my-first-gcp-project-452814.cdle_project_dataset.Diagnoses_ICD` AS di
    LEFT JOIN
      `my-first-gcp-project-452814.cdle_project_dataset.D_ICD_Diagnoses` AS dd
    ↳ ON di.ICD9_CODE = dd.ICD9_CODE
    WHERE
      di.SEQ_NUM = 1 -- Filter for primary diagnosis only
  ),

  -- CTE to count the number of callouts for each hospital admission
  CalloutCount AS (
    SELECT
      HADM_ID,
      COUNT(*) AS NUM_CALLOUTS -- Count callouts per admission
    FROM
      `my-first-gcp-project-452814.cdle_project_dataset.Callout`
    GROUP BY
      HADM_ID
  )

-- Main SELECT statement that uses the CTEs
SELECT
  -- Patient demographics & calculated age
  p.SUBJECT_ID,
  p.GENDER,
  p.DOB,
  p.EXPIRE_FLAG,
```

```

    DATE_DIFF(DATE(a.ADMITTIME), DATE(p.DOB), YEAR) AS AGE_AT_ADMISSION, --□
↪Calculate age

-- Admission details & extracted features
a.HADM_ID,
a.ADMITTIME,
a.DISCHTIME,
a.ADMISSION_TYPE,
a.ADMISSION_LOCATION,
a.INSURANCE,
a.ETHNICITY,
a.DIAGNOSIS AS ADMISSION_DIAGNOSIS_TEXT,
EXTRACT(HOUR FROM a.ADMITTIME) AS ADMISSION_HOUR,

-- ICU stay details (TARGET VARIABLE HERE)
icu.ICUSTAY_ID,
icu.INTIME AS ICU_INTIME,
icu.OUTTIME AS ICU_OUTTIME,
icu.LOS AS ICU_LOS,          -- TARGET VARIABLE
icu.FIRST_CAREUNIT,

-- Primary Diagnosis details (from CTE)
pd.ICD9_CODE AS PRIMARY_ICD9_CODE,
pd.ICD9_SHORT_TITLE AS PRIMARY_ICD9_TITLE,

-- Aggregated Callout details (from CTE)
cc.NUM_CALLOUTS

FROM
-- Start with patients table
`my-first-gcp-project-452814.cdle_project_dataset.Patients` AS p
LEFT JOIN
-- Join with admissions using SUBJECT_ID
`my-first-gcp-project-452814.cdle_project_dataset.Admissions` AS a ON p.
↪SUBJECT_ID = a.SUBJECT_ID
LEFT JOIN
-- Join with ICU stays using HADM_ID
`my-first-gcp-project-452814.cdle_project_dataset.ICUSTays` AS icu ON a.
↪HADM_ID = icu.HADM_ID
LEFT JOIN
-- Join with pre-filtered primary diagnosis using HADM_ID
PrimaryDiagnosis AS pd ON a.HADM_ID = pd.HADM_ID
LEFT JOIN
-- Join with pre-aggregated callout counts using HADM_ID
CalloutCount AS cc ON a.HADM_ID = cc.HADM_ID

```

```
"""
```

```
query_job = client.query(query)
query_job.result()
```

```
[ ]: <google.cloud.bigquery.table._EmptyRowIterator at 0x79c9620a6dd0>
```

3.0.1 Visualization

```
[ ]: dataset_id = "cdle_project_dataset"
      table_id = "junction_table"

      table_ref = client.dataset(dataset_id).table(table_id)

      junction_table = client.list_rows(table_ref).to_dataframe()

      junction_table.head()
```

```
[ ]:  SUBJECT_ID  GENDER  DOB  EXPIRE_FLAG  AGE_AT_ADMISSION  \
0      21852      M  2079-10-15  00:00:00+00:00      1      54
1      19082      F  1869-10-30  00:00:00+00:00      1     300
2      23493      M  2103-04-03  00:00:00+00:00      0      0
3      20778      F  2105-09-02  00:00:00+00:00      1     75
4      13647      M  2081-12-22  00:00:00+00:00      0     27
```

```
      HADM_ID  ADMITTIME  DISCHTIME  ADMISSION_TYPE  \
0  144317  2133-06-02  05:09:00+00:00  2133-06-02  03:14:00+00:00  EMERGENCY
1  187326  2169-10-30  04:28:00+00:00  2169-11-06  19:45:00+00:00  EMERGENCY
2  127281  2103-04-03  09:25:00+00:00  2103-04-08  12:17:00+00:00  NEWBORN
3  113683  2180-12-02  17:46:00+00:00  2180-12-11  19:21:00+00:00  EMERGENCY
4  143439  2108-12-17  01:53:00+00:00  2108-12-23  13:00:00+00:00  EMERGENCY
```

```
      ADMISSION_LOCATION  ...  ADMISSION_DIAGNOSIS_TEXT  \
0  ** INFO NOT AVAILABLE **  ...  CARDIAC ARREST
1  ** INFO NOT AVAILABLE **  ...  RESPIRATORY FAILURE,UROSEPSIS
2  ** INFO NOT AVAILABLE **  ...  NEWBORN
3  ** INFO NOT AVAILABLE **  ...  ANKLE FRACTURE
4  ** INFO NOT AVAILABLE **  ...  GUN SHOT WOUND
```

```
      ADMISSION_HOUR  ICUSTAY_ID  ICU_INTIME  \
0      5      251228  2133-06-02  05:10:11+00:00
1      4      263042  2169-10-30  04:29:22+00:00
2      9      252466  2103-04-03  09:36:42+00:00
3     17      260515  2180-12-05  16:40:35+00:00
4      1      292926  2108-12-17  02:40:14+00:00
```

```
      ICU_OUTTIME  ICU_LOS  FIRST_CAREUNIT  PRIMARY_ICD9_CODE  \
```

0	2133-06-02 05:10:27+00:00	0.0002	CSRU	4271
1	2169-11-03 15:29:54+00:00	4.4587	MICU	51881
2	2103-04-03 22:29:06+00:00	0.5364	NICU	76519
3	2180-12-06 17:48:37+00:00	1.0472	MICU	8244
4	2108-12-18 18:06:51+00:00	1.6435	TSICU	86349

	PRIMARY_ICD9_TITLE	NUM_CALLOUTS
0	Parox ventric tachycard	<NA>
1	Acute respiratry failure	1
2	Preterm NEC 2500+g	<NA>
3	Fx bimalleolar-closed	1
4	Colon injury NEC-closed	<NA>

[5 rows x 22 columns]

4 Data Pre-Processing - Dask

In this section of the project, we will use **Dask** to perform some pre-processing on the `junction_table`, to deal with repeated rows and null values

4.1 Duplicate Rows Analysis

```
[5]: dtype = {
    "SUBJECT_ID": "float64",
    "GENDER": "object",
    "DOB": "object",
    "EXPIRE_FLAG": "float64",
    "AGE_AT_ADMISSION": "float64",
    "HADM_ID": "float64",
    "ADMITTIME": "object",
    "DISCHTIME": "object",
    "ADMISSION_TYPE": "object",
    "ADMISSION_LOCATION": "object",
    "INSURANCE": "object",
    "ETHNICITY": "object",
    "ADMISSION_DIAGNOSIS_TEXT": "object",
    "ADMISSION_HOUR": "float64",
    "ICUSTAY_ID": "float64",
    "ICU_INTIME": "object",
    "ICU_OUTTIME": "object",
    "ICU_LOS": "float64",
    "FIRST_CAREUNIT": "object",
    "PRIMARY_ICD9_CODE": "object",
    "PRIMARY_ICD9_TITLE": "object",
    "NUM_CALLOUTS": "float64"
}
```

```
ddf = dd.read_csv('gs://n_cdle_project/junction_table', dtype=dtype)
```

```
[6]: print(ddf.dtypes)
```

```
SUBJECT_ID          float64
GENDER              string[pyarrow]
DOB                 string[pyarrow]
EXPIRE_FLAG         float64
AGE_AT_ADMISSION    float64
HADM_ID             float64
ADMITTIME           string[pyarrow]
DISCHTIME           string[pyarrow]
ADMISSION_TYPE      string[pyarrow]
ADMISSION_LOCATION  string[pyarrow]
INSURANCE           string[pyarrow]
ETHNICITY           string[pyarrow]
ADMISSION_DIAGNOSIS_TEXT string[pyarrow]
ADMISSION_HOUR      float64
ICUSTAY_ID          float64
ICU_INTIME          string[pyarrow]
ICU_OUTTIME         string[pyarrow]
ICU_LOS             float64
FIRST_CAREUNIT      string[pyarrow]
PRIMARY_ICD9_CODE   string[pyarrow]
PRIMARY_ICD9_TITLE  string[pyarrow]
NUM_CALLOUTS        float64
dtype: object
```

```
[8]: start_time = time.time()

# Group by ICUSTAY_ID and count number of rows
result_ddf = (
    ddf.groupby("ICUSTAY_ID")
        .size()
        .to_frame("number_of_rows")      # convert to DataFrame and name the
    ↪count column
        .reset_index()                  # reset index to make ICUSTAY_ID a
    ↪column again
        .query("number_of_rows > 1")
        .sort_values("number_of_rows", ascending=False)
)

# Compute the result
results = result_ddf.compute()

# Print results
```

```

print("ICUSTAY_ID | number_of_rows")
print("-----")
for _, row in results.iterrows():
    print(f"{row['ICUSTAY_ID']} | {row['number_of_rows']}")

end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")

```

```

ICUSTAY_ID | number_of_rows
-----
Query Execution Time: 0.90 seconds

```

4.1.1 Conclusion

Based on the results of the query, we can conclude that besides the rows that contain the value None that are not any duplicate rows.

4.2 Null Values Analysis

```

[10]: start_time = time.time()

# Total number of rows
total_rows = ddf.shape[0].compute()

# Count of nulls per column
null_counts = ddf.isna().sum().compute()

# Percentage of nulls per column
null_percentages = (null_counts / total_rows * 100).round(2)

# Combine into a single transposed DataFrame for display
null_analysis_df = pd.DataFrame({
    'null_count': null_counts,
    'null_percentage': null_percentages
}).T

print(f"Total number of rows: {total_rows}")
print(" Null Value Analysis (Transposed View):")
print(null_analysis_df)

end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")

```

```

Total number of rows: 62722
Null Value Analysis (Transposed View):
SUBJECT_ID  GENDER  DOB  EXPIRE_FLAG  AGE_AT_ADMISSION  \

```

null_count	0.0	0.0	0.0	0.0	0.0
null_percentage	0.0	0.0	0.0	0.0	0.0

	HADM_ID	ADMITTIME	DISCHTIME	ADMISSION_TYPE	\
null_count	0.0	0.0	0.0	0.0	
null_percentage	0.0	0.0	0.0	0.0	

	ADMISSION_LOCATION	...	ADMISSION_DIAGNOSIS_TEXT	\
null_count		0.0	...	25.00
null_percentage		0.0	...	0.04

	ADMISSION_HOUR	ICUSTAY_ID	ICU_INTIME	ICU_OUTTIME	ICU_LOS	\
null_count	0.0	1190.0	1190.0	1200.00	1200.00	
null_percentage	0.0	1.9	1.9	1.91	1.91	

	FIRST_CAREUNIT	PRIMARY_ICD9_CODE	PRIMARY_ICD9_TITLE	\
null_count	1190.0	47.00	839.00	
null_percentage	1.9	0.07	1.34	

	NUM_CALLOUTS
null_count	31366.00
null_percentage	50.01

[2 rows x 22 columns]

Query Execution Time: 2.34 seconds

4.2.1 Conclusion

There are 62722 rows. The core patient and admission data is complete. However, approximately 1.9% of rows have nulls for ICU-related features, including the target variable ICU_LOS, primarily representing non-ICU admissions or missing discharge times. The primary diagnosis title has minor missingness (~1.3%). The biggest problem is the NUM_CALLOUTS feature, that has a high null rate (50%), requiring specific handling.

4.2.2 Approach to Missing Values in NUM_CALLOUTS

Knowing that the feature NUM_CALLOUTS is not too essential to predict the length of stay, we will remove this column from the table to avoid future processing problems.

```
[11]: ddf = ddf.drop(columns=["NUM_CALLOUTS"])
```

4.2.3 Approach to other Missing Values

Knowing that the rows with missing values in other columns, including the target column ICU_LOS represent a very small percentage of the dataset, we decided the simplest and most effective solution would be to remove these rows.

```
[12]: ddf = ddf.dropna()
```

4.2.4 Running the Null Values Query again

In order to assess if our changes were successful, we decided to run the Null Values query again.

```
[14]: start_time = time.time()

# Total number of rows
total_rows = ddf.shape[0].compute()

# Count of nulls per column
null_counts = ddf.isna().sum().compute()

# Percentage of nulls per column
null_percentages = (null_counts / total_rows * 100).round(2)

# Combine into a single transposed DataFrame for display
null_analysis_df = pd.DataFrame({
    'null_count': null_counts,
    'null_percentage': null_percentages
}).T

print(f"Total number of rows: {total_rows}")
print(" Null Value Analysis (Transposed View):")
print(null_analysis_df)

end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")
```

Total number of rows: 60748

Null Value Analysis (Transposed View):

	SUBJECT_ID	GENDER	DOB	EXPIRE_FLAG	AGE_AT_ADMISSION	\
null_count	0.0	0.0	0.0	0.0	0.0	
null_percentage	0.0	0.0	0.0	0.0	0.0	

	HADM_ID	ADMITTIME	DISCHTIME	ADMISSION_TYPE	\
null_count	0.0	0.0	0.0	0.0	
null_percentage	0.0	0.0	0.0	0.0	

	ADMISSION_LOCATION	...	ETHNICITY	ADMISSION_DIAGNOSIS_TEXT	\
null_count	0.0	...	0.0	0.0	
null_percentage	0.0	...	0.0	0.0	

	ADMISSION_HOUR	ICUSTAY_ID	ICU_INTIME	ICU_OUTTIME	ICU_LOS	\
null_count	0.0	0.0	0.0	0.0	0.0	
null_percentage	0.0	0.0	0.0	0.0	0.0	

	FIRST_CAREUNIT	PRIMARY_ICD9_CODE	PRIMARY_ICD9_TITLE
--	----------------	-------------------	--------------------

null_count	0.0	0.0	0.0
null_percentage	0.0	0.0	0.0

[2 rows x 21 columns]

Query Execution Time: 2.56 seconds

4.2.5 Conclusion

There are 60748 rows (~96.8% of the original table) and there are no missing values, as expected.

Fixing Column Types

```
[15]: timestamp_cols = ["DOB", "ADMITTIME", "DISCHTIME", "ICU_INTIME", "ICU_OUTTIME"]

for col in timestamp_cols:
    ddf[col] = dd.to_datetime(ddf[col], errors="coerce")

int_cols = ["SUBJECT_ID", "EXPIRE_FLAG", "AGE_AT_ADMISSION", "HADM_ID",
            ↪ "ADMISSION_HOUR", "ICUSTAY_ID"]

for col in int_cols:
    ddf[col] = ddf[col].astype("int64")

cat_cols = ["GENDER", "ADMISSION_TYPE", "ADMISSION_LOCATION", "INSURANCE",
            ↪ "ETHNICITY", "ADMISSION_DIAGNOSIS_TEXT", "FIRST_CAREUNIT",
            ↪ "PRIMARY_ICD9_CODE", "PRIMARY_ICD9_TITLE"]

for col in cat_cols:
    ddf[col] = ddf[col].astype("category")
```

5 Dataset Analysis - Dask

In this section, we use **Dask** to perform a dataset analysis on three different topics, **Patients**, **Medical Data** and **Correlations**.

In **Patients** and **Medical Data** we will analyse the distribution of several features and try to find connections between them and the target feature (**Length of ICU Stay**), to get a better understanding of what features will be more relevant to our ML model.

In **Correlations** we will try to find connections between different features of the dataset, to gain an even better insight of the data.

5.1 Patients

We will start this analysis with an in-depth analysis of the features regarding patients characteristics, which include:

- Age
- Insurance

5.1.1 Age Analysis

```
[19]: start_time = time.time()


# 1. Filter Dask DataFrame for non-null AGE_AT_ADMISSION
age_ddf = ddf["AGE_AT_ADMISSION"]

# 2. Compute to Pandas
age_series = age_ddf.compute()
age_df = age_series.to_frame(name="AGE_AT_ADMISSION") # <- Convert to DataFrame

# 3. Plot histogram + KDE
if not age_df.empty:
    plt.figure(figsize=(10, 6))
    sns.histplot(data=age_df, x='AGE_AT_ADMISSION', kde=True, bins=30)
    plt.xlabel("Age at Admission (Years)")
    plt.ylabel("Number of Patients")
    plt.title("Distribution of Patient Age at Admission")
    plt.tight_layout()
    plt.show()
else:
    print("No non-null AGE_AT_ADMISSION values to plot.")

# 4. Print summary statistics
print("\nAge at Admission Statistics:")
print(age_df["AGE_AT_ADMISSION"].describe())

end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")
```



Age at Admission Statistics:

```
count    60748.000000
mean      64.939998
std       57.138132
min        0.000000
25%       44.000000
50%       62.000000
75%       76.000000
max       311.000000
Name: AGE_AT_ADMISSION, dtype: float64
Query Execution Time: 2.38 seconds
```

5.1.2 Eliminating rows with Age bigger then 120 years

Given the histogram results, we decided to eliminate the rows with outliers that most definitely represent input errors, given that there are no people near 300 years old.

```
[21]: start_time = time.time()

age_df_filtered = age_df[age_df["AGE_AT_ADMISSION"] <= 120]

# 3. Plot histogram + KDE
if not age_df_filtered.empty:
    plt.figure(figsize=(10, 6))
    sns.histplot(data=age_df_filtered, x='AGE_AT_ADMISSION', kde=True, bins=30)
```

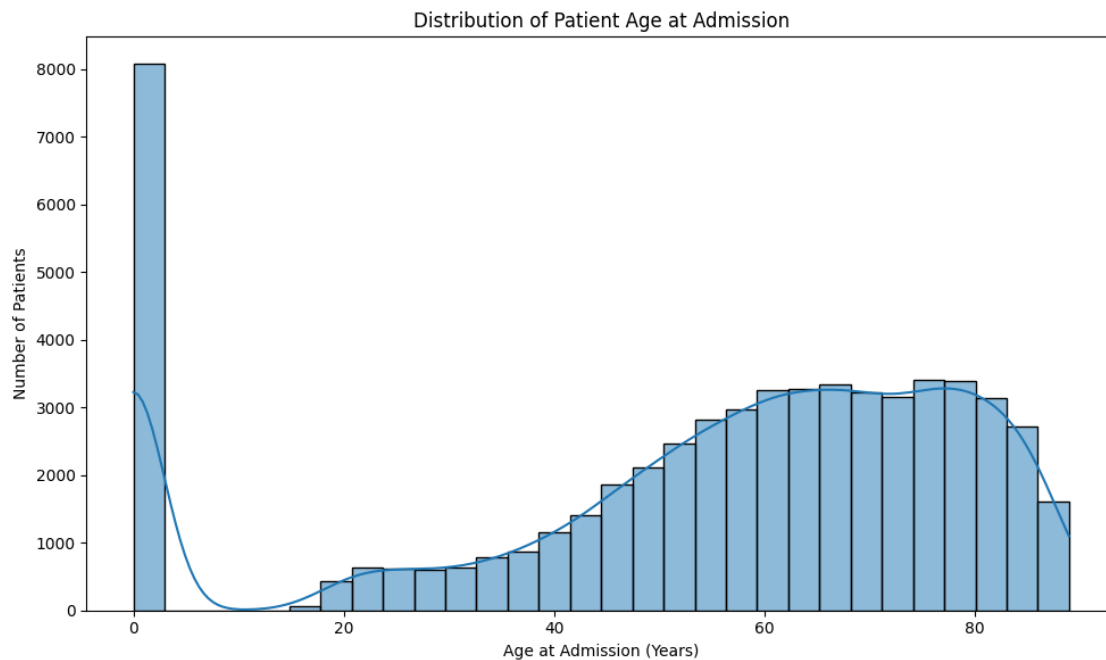
```

plt.xlabel("Age at Admission (Years)")
plt.ylabel("Number of Patients")
plt.title("Distribution of Patient Age at Admission")
plt.tight_layout()
plt.show()
else:
    print("No non-null AGE_AT_ADMISSION values to plot.")

# 4. Print summary statistics
print("\nAge at Admission Statistics:")
print(age_df_filtered["AGE_AT_ADMISSION"].describe())

end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")

```



```

Age at Admission Statistics:
count    58039.000000
mean      53.943693
std       26.560400
min        0.000000
25%       43.000000
50%       61.000000
75%       74.000000
max       89.000000

```

Name: AGE_AT_ADMISSION, dtype: float64

Query Execution Time: 0.78 seconds

After eliminating these input error mistakes, we ended up with a table containing 58039 rows (~95.5% of the previous table).

5.1.3 Age & Length of ICU Stay

```
[28]: start_time = time.time()


# 1. Filter AGE_AT_ADMISSION to remove outliers
filtered_ddf = ddf[ddf["AGE_AT_ADMISSION"] <= 100]

# 2. Select relevant columns and bring to memory
age_los_df = filtered_ddf[["AGE_AT_ADMISSION", "ICU_LOS"]].compute()

# 3. Plot and correlation
if not age_los_df.empty:
    plt.figure(figsize=(10, 7))
    sns.scatterplot(x="AGE_AT_ADMISSION", y="ICU_LOS", data=age_los_df, alpha=0.
↪3, s=15)
    plt.xlabel("Age at Admission (Years)")
    plt.ylabel("ICU Length of Stay (Days)")
    plt.title("Relationship between Age at Admission and ICU Length of Stay")
    plt.ylim(bottom=0)
    plt.xlim(left=0)
    plt.tight_layout()
    plt.show()

    # Correlation
    correlation = age_los_df["AGE_AT_ADMISSION"].corr(age_los_df["ICU_LOS"])
    print(f"\nCorrelation between Age at Admission and ICU LOS: {correlation:.
↪3f}")
else:
    print("No data to plot.")

end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")
```



Correlation between Age at Admission and ICU LOS: -0.161
 Query Execution Time: 6.84 seconds

5.1.4 Conclusion

By analysing the histogram of **Age**, we can see that the most common age intervals in the dataset are 0 and 50–80. The dataset purposely omits data from minor patients with the exception of newborns.

Regarding the connection between **Age** and **Length of ICU Stay**, we can conclude that: 1. The vast majority of ICU stays are relatively short (concentrated roughly below 25 days) across all adult age groups. 2. While short stays are common for all ages, the spread of the values of **Length of ICU Stay** appears to widen slightly for older patients, which means that very long stays seem slightly more prevalent among older age groups.

5.1.5 Insurance Analysis

```
[34]: start_time = time.time()

# 1. Group by insurance and count occurrences
insurance_counts = (
    filtered_ddf.groupby("INSURANCE")
```

```

        .size()
        .compute()
        .sort_values(ascending=False)
        .reset_index(name="count")
        .rename(columns={0: "INSURANCE"})
    )

# 2. Plot horizontal bar chart
if not insurance_counts.empty:
    plt.figure(figsize=(10, 5))
    sns.barplot(x='count', y='INSURANCE', data=insurance_counts,
        palette='magma', orient='h')
    plt.xlabel("Number of Patients")
    plt.ylabel("Insurance Type")
    plt.title("Distribution of Patients by Insurance Type")
    plt.tight_layout()
    plt.show()
else:
    print("No data to plot.")

end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")

```

```

/usr/local/lib/python3.11/dist-
packages/dask/dataframe/dask_expr/_groupby.py:1562: FutureWarning: The default
of observed=False is deprecated and will be changed to True in a future version
of pandas. Pass observed=False to retain current behavior or observed=True to
adopt the future default and silence this warning.

```

```

    self._meta = self.obj._meta.groupby(
<ipython-input-34-7ff799173bd5>:16: FutureWarning:

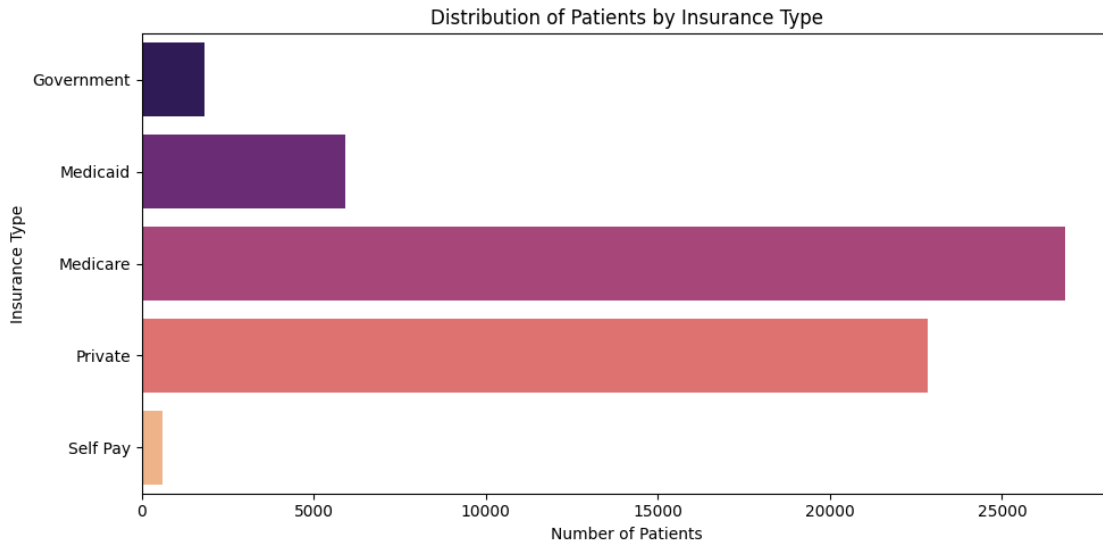
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```

    sns.barplot(x='count', y='INSURANCE', data=insurance_counts, palette='magma',
orient='h')

```



Query Execution Time: 2.77 seconds

5.1.6 Insurance & Length of ICU Stay

```
[36]: start_time = time.time()

# 1. Select relevant columns
insurance_los_ddf = filtered_ddf[["INSURANCE", "ICU_LOS"]]

# 2. Bring to memory for plotting/stats
insurance_los_df = insurance_los_ddf.compute()

# 3. Plot and summarize
if not insurance_los_df.empty:
    plt.figure(figsize=(12, 7))
    sns.boxplot(x='INSURANCE', y='ICU_LOS', data=insurance_los_df,
                palette='magma', showfliers=False)
    plt.xlabel("Insurance Type")
    plt.ylabel("Length of ICU Stay (Days)")
    plt.title("Distribution of Length of ICU Stay by Insurance Type")
    plt.xticks(rotation=45, ha='right')
    plt.ylim(bottom=0)
    plt.tight_layout()
    plt.show()

    # Summary statistics
    print("\nSummary Statistics for ICU LOS by Insurance Type:")
    print(insurance_los_df.groupby("INSURANCE")["ICU_LOS"].describe())
else:
```



```

print("No data returned for insurance vs. LOS analysis.")

end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")

```

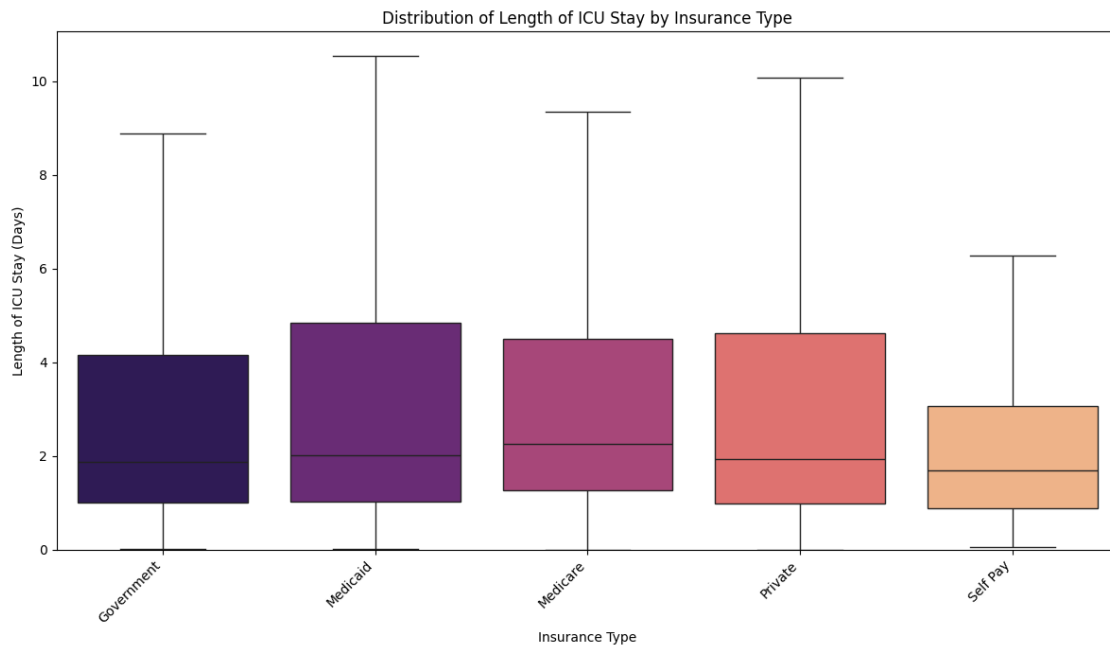
<ipython-input-36-e7db03b38274>:12: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```

sns.boxplot(x='INSURANCE', y='ICU_LOS', data=insurance_los_df,
palette='magma', showfliers=False)

```



Summary Statistics for ICU LOS by Insurance Type:

	count	mean	std	min	25%	50%	75%	\
INSURANCE								
Government	1822.0	4.895602	9.793321	0.0057	0.998850	1.87805	4.167150	
Medicaid	5919.0	5.900825	12.800130	0.0036	1.025900	2.02050	4.835200	
Medicare	26843.0	4.306665	6.252719	0.0001	1.260800	2.26430	4.501300	
Private	22845.0	5.611555	12.140707	0.0002	0.982100	1.92950	4.625200	
Self Pay	610.0	3.062356	4.486948	0.0614	0.888425	1.68900	3.057675	

max

```

INSURANCE
Government    101.8397
Medicaid     169.4202
Medicare      173.0725
Private       171.6227
Self Pay      43.1465
Query Execution Time: 7.77 seconds

```

<ipython-input-36-e7db03b38274>:23: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```
print(insurance_los_df.groupby("INSURANCE")["ICU_LOS"].describe())
```

5.1.7 Conclusion

By analysing the bar chart of **Insurance**, we can see that the most common insurance systems are **Medicare** followed by **Private**, and that there is a significantly lower number of patients using other systems.

Regarding the connection between **Insurance** and **Length of ICU Stay**, we can conclude that all systems have a similar distribution and mean in the comparison of the two metrics, with the exception of **Self Pay**, whose average value for **Length of ICU Stay** is significantly lower.

5.2 Medical Data

We will continue the dataset analysis with an in-depth analysis of the features regarding patients medical records while they were hospitalized, which include:

- Type of Admission
- Diagnostic Code
- First Care Unit
- Length of ICU Stay (Target)

5.2.1 Type of Admission Analysis

```

[33]: start_time = time.time()

# 1. Group by ADMISSION_TYPE and count
adm_type_counts = (
    filtered_ddf.groupby("ADMISSION_TYPE")
        .size()
        .compute()
        .sort_values(ascending=False)
        .reset_index(name="count")
)

# 2. Plot bar chart
if not adm_type_counts.empty:
    plt.figure(figsize=(8, 5))

```

```

    ax = sns.barplot(x='ADMISSION_TYPE', y='count', data=adm_type_counts,
↪palette='cubehelix')
    plt.xlabel("Admission Type")
    plt.ylabel("Number of Patients")
    plt.title("Distribution of Patients by Admission Type")
    plt.xticks(rotation=45, ha='right')

    # Add value labels to bars
    for container in ax.containers:
        ax.bar_label(container)

    plt.tight_layout()
    plt.show()
else:
    print("No data returned for admission type distribution.")

end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")

```

/usr/local/lib/python3.11/dist-packages/dask/dataframe/dask_expr/_groupby.py:1562: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```

    self._meta = self.obj._meta.groupby(
<ipython-input-33-6ba51112c5e1>:15: FutureWarning:

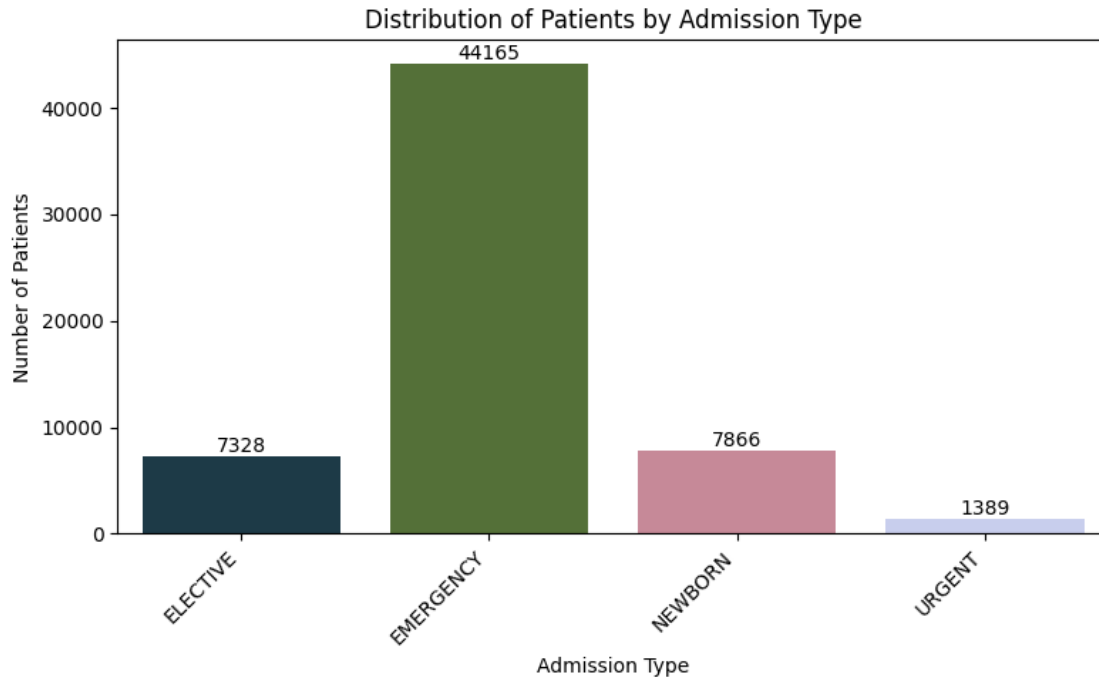
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```

    ax = sns.barplot(x='ADMISSION_TYPE', y='count', data=adm_type_counts,
palette='cubehelix')

```



Query Execution Time: 2.66 seconds

5.2.2 Type of Admission & Length of ICU Stay

```
[37]: start_time = time.time()

# 1. Select relevant columns from the cleaned DataFrame
adm_type_los_ddf = filtered_ddf[["ADMISSION_TYPE", "ICU_LOS"]]

# 2. Compute to Pandas
adm_type_los_df = adm_type_los_ddf.compute()

# 3. Plot and describe
if not adm_type_los_df.empty:
    plt.figure(figsize=(10, 6))
    sns.boxplot(x='ADMISSION_TYPE', y='ICU_LOS', data=adm_type_los_df,
                palette='cubehelix', showfliers=False)
    plt.xlabel("Admission Type")
    plt.ylabel("ICU Length of Stay (Days)")
    plt.title("Distribution of ICU Length of Stay by Admission Type")
    plt.xticks(rotation=45, ha='right')
    plt.ylim(bottom=0)
    plt.tight_layout()
    plt.show()
```

```

# Summary statistics
print("\nSummary Statistics for ICU LOS by Admission Type:")
print(adm_type_los_df.groupby("ADMISSION_TYPE")["ICU_LOS"].describe())
else:
    print("No data returned for admission type vs. LOS analysis.")

end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")

```

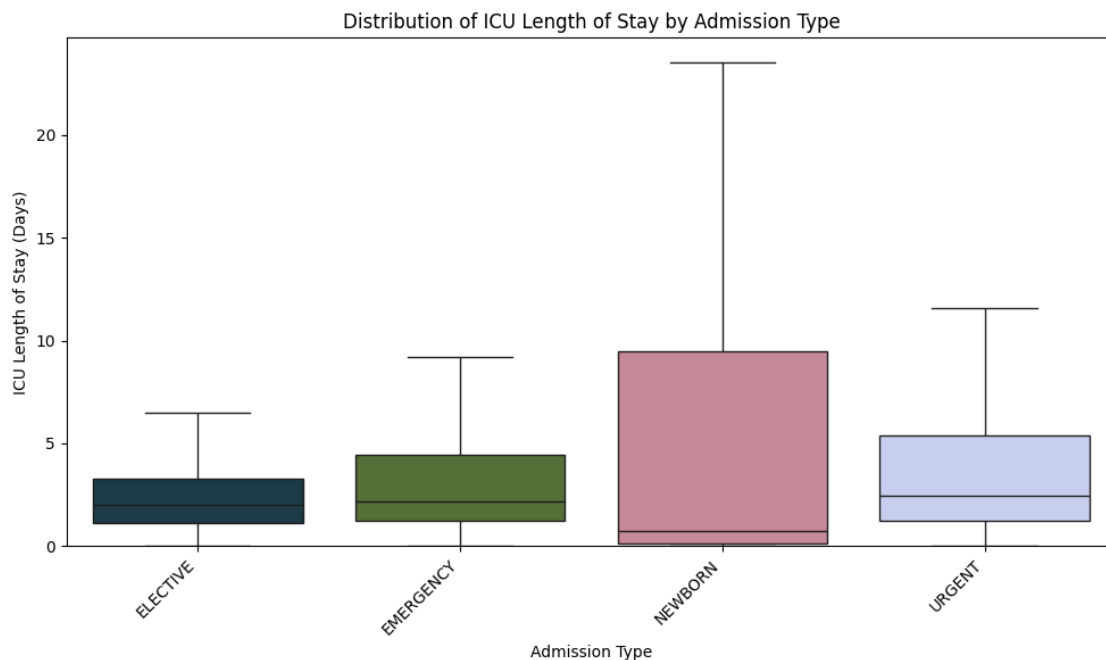
<ipython-input-37-1121dcee2c5f>:12: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```

sns.boxplot(x='ADMISSION_TYPE', y='ICU_LOS', data=adm_type_los_df,
palette='cubehelix', showfliers=False)

```



Summary Statistics for ICU LOS by Admission Type:

	count	mean	std	min	25%	50%	\
ADMISSION_TYPE							
ELECTIVE	7226.0	3.503513	6.008552	0.0004	1.140675	1.9895	
EMERGENCY	41601.0	4.309662	6.436865	0.0001	1.229700	2.1897	
NEWBORN	7866.0	9.878190	20.483139	0.0008	0.137675	0.7253	

URGENT	1346.0	5.356694	8.359788	0.0025	1.241075	2.4351
--------	--------	----------	----------	--------	----------	--------

	75%	max
ADMISSION_TYPE		
ELECTIVE	3.283750	173.0725
EMERGENCY	4.420000	169.4202
NEWBORN	9.500975	171.6227
URGENT	5.380250	97.4897

Query Execution Time: 8.09 seconds

<ipython-input-37-1121dcee2c5f>:23: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```
print(adm_type_los_df.groupby("ADMISSION_TYPE")["ICU_LOS"].describe())
```

5.2.3 Conclusion

The distributions of patients by Type of Admission is uneven, with most admissions being classified as EMERGENCY. The least recurrent type of admissions is Urgent.

There is an evident connection between Type of Admission and Length of ICU Stay, with NEWBORN having the highest mean values (close to 10 days) while the other types of admission have mean values between 3 to 5 days.

5.2.4 Diagnosis Analysis

```
[51]: start_time = time.time()

# Ensure PRIMARY_ICD9_CODE is string
filtered_ddf = filtered_ddf.
    ↪ assign(PRIMARY_ICD9_CODE=filtered_ddf["PRIMARY_ICD9_CODE"].astype(str))

# 1. Count top 20 ICD-9 codes (and fix column names properly)
icd_counts = (
    filtered_ddf["PRIMARY_ICD9_CODE"]
    .value_counts()
    .compute()
    .nlargest(20)
    .reset_index()
    .rename(columns={"index": "PRIMARY_ICD9_CODE", 0: "count"})
)

# 2. Get titles
icd_titles = (
    filtered_ddf[["PRIMARY_ICD9_CODE", "PRIMARY_ICD9_TITLE"]]
    .drop_duplicates(subset=["PRIMARY_ICD9_CODE"])
    .compute()
    .dropna(subset=["PRIMARY_ICD9_TITLE"])
)
```

```

)

# 3. Merge counts with titles
icd9_distribution_df = icd_counts.merge(icd_titles, on="PRIMARY_ICD9_CODE",
    ↪how="left")

# 4. Plot
if not icd9_distribution_df.empty:
    plt.figure(figsize=(10, 10))
    icd9_distribution_df["PRIMARY_ICD9_TITLE"] =
    ↪icd9_distribution_df["PRIMARY_ICD9_TITLE"].astype(str)
    icd9_distribution_df["CODE_TITLE"] = (
        icd9_distribution_df["PRIMARY_ICD9_CODE"] + ": " +
        icd9_distribution_df["PRIMARY_ICD9_TITLE"].replace("nan", "N/A")
    )

    sns.barplot(
        x="count",
        y="CODE_TITLE",
        data=icd9_distribution_df,
        palette="viridis",
        orient="h"
    )

    plt.xlabel("Number of Patients")
    plt.ylabel("Primary ICD-9 Code")
    plt.title("Distribution of Top 20 Primary ICD-9 Codes at Admission")
    plt.tight_layout()
    plt.show()

    print("\nTop ICD-9 Code Distribution:")
    print(icd9_distribution_df[["PRIMARY_ICD9_CODE", "count",
    ↪"PRIMARY_ICD9_TITLE"]])
else:
    print("No data returned for primary ICD-9 code distribution.")


end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")

```

<ipython-input-51-bd3e9e7b3720>:36: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(
```



Top ICD-9 Code Distribution:

	PRIMARY_ICD9_CODE	count	PRIMARY_ICD9_TITLE
0	41401	3617	Crnry athrscl natve vssl
1	V3000	3534	Single lb in-hosp w/o cs
2	V3001	2792	Single lb in-hosp w cs
3	0389	1988	Septicemia NOS
4	41071	1803	Subendo infarct, initial
5	4241	1168	Aortic valve disorder
6	51881	1124	Acute respiratry failure
7	431	1031	Intracerebral hemorrhage
8	V3101	1008	Twin-mate lb-in hos w cs
9	486	684	Pneumonia, organism NOS

10	5070	596	Food/vomit pneumonitis
11	4240	592	Mitral valve disorder
12	4280	562	CHF NOS
13	430	542	Subarachnoid hemorrhage
14	5849	523	Acute kidney failure NOS
15	41011	473	AMI anterior wall, init
16	41041	471	AMI inferior wall, init
17	5789	415	Gastrointest hemorr NOS
18	5770	399	Acute pancreatitis
19	1983	343	Sec mal neo brain/spine

Query Execution Time: 12.88 seconds

5.2.5 Diagnosis & Length of ICU Stay

```
[52]: start_time = time.time()

# Ensure PRIMARY_ICD9_CODE is as string
filtered_ddf = filtered_ddf.
    ↪ assign(PRIMARY_ICD9_CODE=filtered_ddf["PRIMARY_ICD9_CODE"].astype(str))

# 1. Get top 20 most frequent ICD-9 codes (as strings)
top_codes_list = (
    filtered_ddf["PRIMARY_ICD9_CODE"]
    .value_counts()
    .compute()
    .nlargest(20)
    .index
    .tolist()
)

# 2. Filter and select necessary columns
if top_codes_list:
    icd9_los_ddf = filtered_ddf[
        filtered_ddf["PRIMARY_ICD9_CODE"].isin(top_codes_list)
    ][["PRIMARY_ICD9_CODE", "ICU_LOS"]]

# 3. Compute to pandas
icd9_los_df = icd9_los_ddf.compute()

# 4. Plot
if not icd9_los_df.empty:
    plt.figure(figsize=(15, 8))
    sns.boxplot(
        x="PRIMARY_ICD9_CODE",
        y="ICU_LOS",
        data=icd9_los_df,
        palette="viridis",
```

```

        order=top_codes_list,
        showfliers=False
    )
    plt.xlabel("Primary ICD-9 Code")
    plt.ylabel("ICU Length of Stay (Days)")
    plt.title("Distribution of ICU Length of Stay by Top 20 Primary ICD-9_
↳Codes")
    plt.xticks(rotation=90)
    plt.ylim(bottom=0)
    plt.tight_layout()
    plt.show()

    print("\nSummary Statistics for ICU LOS by Top 20 Primary ICD-9 Codes:")
    print(icd9_los_df.groupby("PRIMARY_ICD9_CODE")["ICU_LOS"].describe())

    else:
        print("No data returned for top ICD-9 codes vs. LOS analysis.")
else:
    print("Cannot proceed with LOS analysis as no top ICD-9 codes were_
↳identified.")

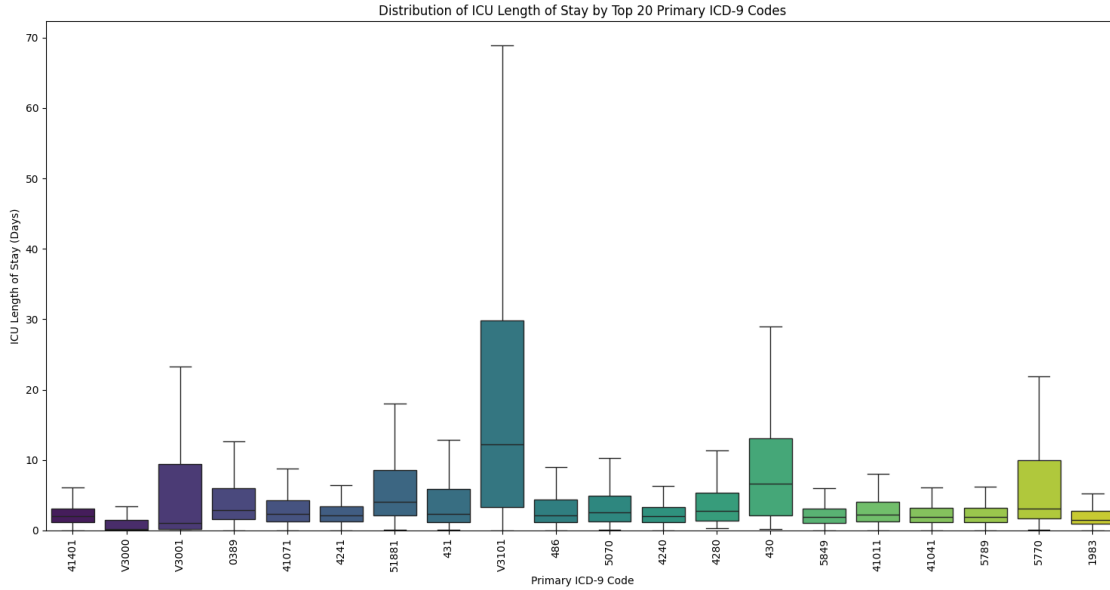
end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")

```

<ipython-input-52-3bf58f3a6526>:28: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(
```



Summary Statistics for ICU LOS by Top 20 Primary ICD-9 Codes:

	count	mean	std	min	25%	50%	\
PRIMARY_ICD9_CODE							
0389	1988.0	5.329782	7.084884	0.0079	1.614200	2.87470	
1983	343.0	2.509505	3.203271	0.0079	1.005550	1.51880	
41011	473.0	3.895759	5.718157	0.0435	1.314000	2.20950	
41041	471.0	3.449264	4.951962	0.0033	1.230300	1.94030	
41071	1803.0	4.026161	5.369919	0.0016	1.288900	2.35840	
41401	3617.0	2.876750	3.885769	0.0048	1.154600	1.99070	
4240	592.0	3.466497	5.780785	0.0077	1.159450	2.02135	
4241	1168.0	3.624112	5.788327	0.0014	1.233175	2.14675	
4280	562.0	4.949970	6.628354	0.3021	1.387300	2.73745	
430	542.0	9.060081	8.904613	0.1733	2.170775	6.62515	
431	1031.0	4.795295	5.835963	0.1118	1.213400	2.30710	
486	684.0	3.930527	4.905842	0.0280	1.196175	2.15545	
5070	596.0	4.541065	5.972763	0.1101	1.334350	2.55865	
51881	1124.0	6.957914	7.858505	0.1524	2.096675	4.12695	
5770	399.0	8.341084	12.501480	0.0566	1.757150	3.08880	
5789	415.0	2.812222	3.300042	0.0214	1.149050	1.96770	
5849	523.0	3.024175	4.131413	0.0012	1.110750	1.90300	
V3000	3534.0	4.589990	14.277310	0.0037	0.104725	0.20955	
V3001	2792.0	10.340239	21.146993	0.0008	0.162275	1.02360	
V3101	1008.0	21.719032	26.430318	0.0098	3.314175	12.19845	

	75%	max
PRIMARY_ICD9_CODE		
0389	6.028650	91.5726

1983	2.743500	30.7172
41011	4.038900	76.9211
41041	3.198900	40.9936
41071	4.309700	53.0333
41401	3.134600	68.2052
4240	3.360900	75.9919
4241	3.437825	98.6446
4280	5.372050	68.9517
430	13.097175	54.5119
431	5.869350	51.7955
486	4.354575	41.5576
5070	4.938750	59.4319
51881	8.548500	71.0056
5770	9.929350	101.7390
5789	3.202750	36.3250
5849	3.082300	43.2606
V3000	1.457850	167.5077
V3001	9.443825	171.6227
V3101	29.786825	142.3605

Query Execution Time: 15.99 seconds

5.2.6 Conclusion

There is a very unsymmetric distribution of Diagnosis. In the top-20 most common diagnosis, with the most common being 41401 (Coronary atherosclerosis of native vessel). The correlation between some Diagnosis and ICU Length of Stay is evident, with V3101 (Twin mate, liveborn, delivered in hospital, with cesarean section) having a mean value very high compared to other types of Diagnosis. There are also diagnosis who have an average of ICU Length of Stay very low in comparison, like V3000 (Single liveborn, delivered in hospital, without cesarean section).

5.2.7 First Care Unit Analysis

```
[53]: start_time = time.time()

# 1. Group by FIRST_CAREUNIT and count
careunit_counts = (
    filtered_ddf.groupby("FIRST_CAREUNIT")
    .size()
    .compute()
    .sort_values(ascending=False)
    .reset_index(name="count")
)

# 2. Plot
if not careunit_counts.empty:
    plt.figure(figsize=(10, 6))
```

```

    ax = sns.barplot(x="FIRST_CAREUNIT", y="count", data=careunit_counts,
↪palette="Spectral")
    plt.xlabel("First Care Unit")
    plt.ylabel("Number of Patients")
    plt.title("Distribution of Patients by First Care Unit")
    plt.xticks(rotation=45, ha="right")

    # Add count labels
    for container in ax.containers:
        ax.bar_label(container)

    plt.tight_layout()
    plt.show()
else:
    print("No data returned for first care unit distribution.")

end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")

```

/usr/local/lib/python3.11/dist-packages/dask/dataframe/dask_expr/_groupby.py:1562: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```

    self._meta = self.obj._meta.groupby(
<ipython-input-53-9eb78ae8727a>:15: FutureWarning:

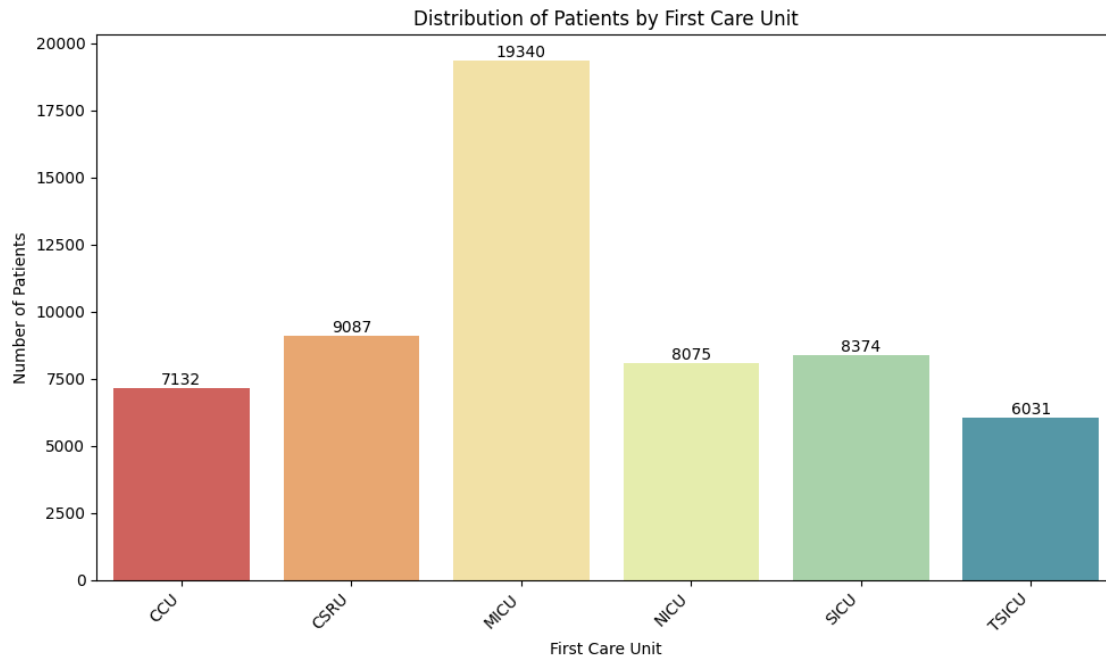
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```

    ax = sns.barplot(x="FIRST_CAREUNIT", y="count", data=careunit_counts,
palette="Spectral")

```



Query Execution Time: 3.80 seconds

5.2.8 First Care of Unit & Length of ICU Stay

```
[54]: start_time = time.time()

# 1. Select relevant columns from the cleaned DataFrame
careunit_los_ddf = filtered_ddf[["FIRST_CAREUNIT", "ICU_LOS"]]

# 2. Compute to Pandas
careunit_los_df = careunit_los_ddf.compute()

# 3. Plot and describe
if not careunit_los_df.empty:
    plt.figure(figsize=(12, 7))
    sns.boxplot(
        x="FIRST_CAREUNIT",
        y="ICU_LOS",
        data=careunit_los_df,
        palette="Spectral",
        showfliers=False
    )
    plt.xlabel("First Care Unit")
    plt.ylabel("ICU Length of Stay (Days)")
    plt.title("Distribution of ICU Length of Stay by First Care Unit")
    plt.xticks(rotation=45, ha="right")
```

```

plt.ylim(bottom=0)
plt.tight_layout()
plt.show()

# Summary statistics
print("\nSummary Statistics for ICU LOS by First Care Unit:")
print(careunit_los_df.groupby("FIRST_CAREUNIT")["ICU_LOS"].describe())
else:
    print("No data returned for first care unit vs. LOS analysis.")

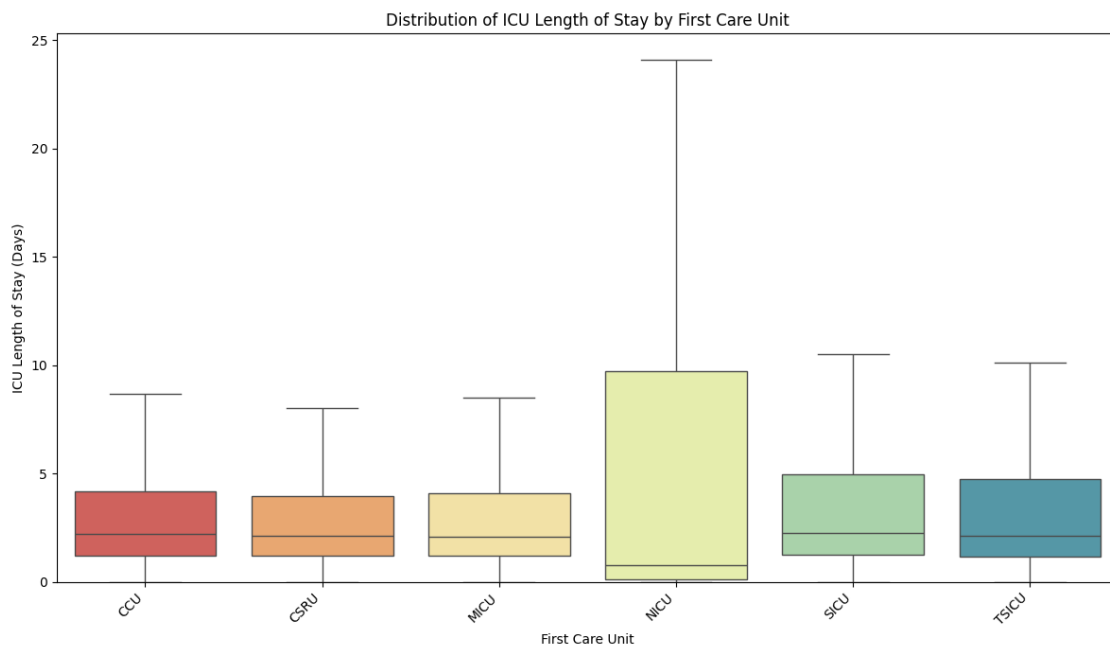
end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")

```

<ipython-input-54-e2d327078eaf>:12: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(
```



Summary Statistics for ICU LOS by First Care Unit:

	count	mean	std	min	25%	50%	\
FIRST_CAREUNIT							

CCU	7132.0	3.953257	5.619626	0.0012	1.213125	2.19915
CSRU	9087.0	3.888229	6.097731	0.0001	1.213850	2.14770
MICU	19340.0	4.042175	5.890149	0.0004	1.189600	2.09700
NICU	8075.0	10.004395	20.644582	0.0008	0.140500	0.78470
SICU	8374.0	4.749462	6.997957	0.0003	1.258775	2.26470
TSICU	6031.0	4.519062	6.760158	0.0016	1.174100	2.12800

	75%	max
FIRST_CAREUNIT		
CCU	4.200700	100.1225
CSRU	3.943200	153.9280
MICU	4.109775	116.8327
NICU	9.723350	171.6227
SICU	4.964775	101.7390
TSICU	4.755900	173.0725

Query Execution Time: 8.19 seconds

<ipython-input-54-e2d327078eaf>:29: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```
print(careunit_los_df.groupby("FIRST_CAREUNIT")["ICU_LOS"].describe())
```

5.2.9 Conclusion

There is a relative uniform distribution in the number of patients by First Care Unit with the exception being MICU (Medical Intensive Care Unit), that has a count of patients that is more than double of each other type of First Care Unit.

The correlation between Type of Care Unit and ICU Length of Stay is also relatively uniform with an exception, NICU (Neonatal Intensive Care Unit).

5.2.10 Length of ICU Stay Analysis

```
[55]: start_time = time.time()

# 1. Select ICU_LOS column from the filtered dataframe
icu_los_series = filtered_ddf["ICU_LOS"]

# 2. Compute to Pandas for plotting
icu_los_df = icu_los_series.compute().to_frame(name="ICU_LOS")

# 3. Filter for LOS within 0-25 days
plot_data = icu_los_df[(icu_los_df["ICU_LOS"] >= 0) & (icu_los_df["ICU_LOS"] <= 25)]

# 4. Plot
if not plot_data.empty:
    plt.figure(figsize=(10, 6))
```




```

sns.histplot(data=plot_data, x="ICU_LOS", kde=False, bins=5)
plt.xlabel("ICU Length of Stay (Days)")
plt.ylabel("Number of Patients")
plt.title("Distribution of ICU Length of Stay (0-25 Days)")
plt.xlim(left=0, right=25)
plt.tight_layout()
plt.show()

# Descriptive stats
print("\nICU Length of Stay Statistics:")
print(icu_los_df["ICU_LOS"].describe())
else:
    print("No data returned for ICU LOS distribution (0-25 days).")

end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")

```



ICU Length of Stay Statistics:

count	58039.000000
mean	4.988276
std	9.824303
min	0.000100
25%	1.103350
50%	2.091800

```
75%          4.550800
max          173.072500
Name: ICU_LOS, dtype: float64
Query Execution Time: 6.30 seconds
```

5.2.11 Conclusion

By analysing the distribution of ICU Length of Stay we can see clearly that most of patients only stay for 5 days or less, with a decreasing trend of the number of patients as the number of days increases.

5.3 Correlations

To finish this analysis we will attempt to find connections between the following set of features:

- Type of Admission & Hour of Admission
- Type of Admission and Death
- Diagnosis & First Care Unit
- First Care Unit & Death

5.3.1 Type of Admission & Hour of Admission

```
[60]: start_time = time.time()

# 1. Group by ADMISSION_TYPE and ADMISSION_HOUR and count occurrences
adm_type_hour_df = (
    filtered_ddf.groupby(["ADMISSION_TYPE", "ADMISSION_HOUR"])
    .size()
    .compute()
    .reset_index(name="count")
)

# 2. Ensure numeric types
adm_type_hour_df["count"] = pd.to_numeric(adm_type_hour_df["count"],
    ↪errors="coerce")
adm_type_hour_df["ADMISSION_HOUR"] = pd.
    ↪to_numeric(adm_type_hour_df["ADMISSION_HOUR"], errors="coerce")
adm_type_hour_df.dropna(subset=["count", "ADMISSION_HOUR"], inplace=True)

# 3. Plot heatmap
if not adm_type_hour_df.empty:
    try:
        heatmap_data = (
            adm_type_hour_df
            .pivot(index="ADMISSION_TYPE", columns="ADMISSION_HOUR",
    ↪values="count")
            .fillna(0)
            .astype(float)
```

```

    )

    plt.figure(figsize=(15, 6))
    sns.heatmap(heatmap_data, annot=False, fmt=".0f", linewidths=.5,
    cmap="YlGnBu")
    plt.xlabel("Hour of Admission (0-23)")
    plt.ylabel("Admission Type")
    plt.title("Heatmap of Admission Counts by Type and Hour")
    plt.yticks(rotation=0)
    plt.tight_layout()
    plt.show()

    # Print the table
    print("\nPivoted Data for Heatmap:")
    print(heatmap_data)
    print("\nPivoted Data Types:")
    print(heatmap_data.dtypes)

except Exception as e:
    print(f"An error occurred during plotting: {e}")
    print("\nOriginal DataFrame dtypes before pivot:")
    print(adm_type_hour_df.dtypes)
    if "heatmap_data" in locals():
        print("\nPivoted DataFrame dtypes before plotting:")
        print(heatmap_data.dtypes)

else:
    print("No data returned from the grouped admission type/hour analysis.")

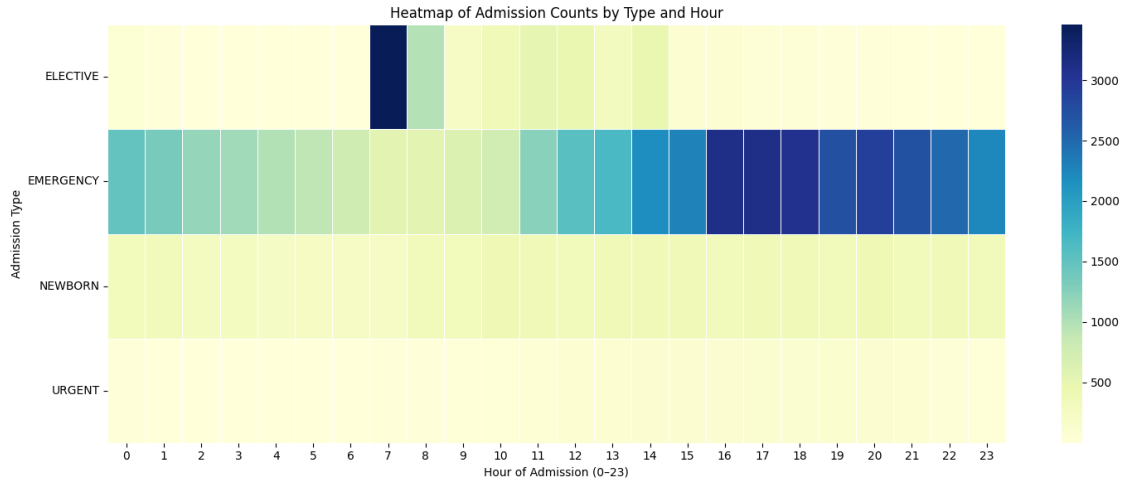
end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")

```

```

/usr/local/lib/python3.11/dist-
packages/dask/dataframe/dask_expr/_groupby.py:1562: FutureWarning: The default
of observed=False is deprecated and will be changed to True in a future version
of pandas. Pass observed=False to retain current behavior or observed=True to
adopt the future default and silence this warning.
    self._meta = self.obj._meta.groupby(

```



Pivoted Data for Heatmap:

ADMISSION_HOUR	0	1	2	3	4	5	6	7	\
ADMISSION_TYPE									
ELECTIVE	70.0	20.0	12.0	9.0	2.0	4.0	6.0	3457.0	
EMERGENCY	1474.0	1348.0	1169.0	1095.0	989.0	917.0	778.0	552.0	
NEWBORN	324.0	330.0	265.0	282.0	236.0	245.0	226.0	229.0	
URGENT	23.0	14.0	12.0	10.0	4.0	8.0	13.0	23.0	

ADMISSION_HOUR	8	9	...	14	15	16	17	18	\
ADMISSION_TYPE			...						
ELECTIVE	986.0	222.0	...	461.0	84.0	87.0	51.0	42.0	
EMERGENCY	554.0	643.0	...	2184.0	2292.0	3111.0	3111.0	3077.0	
NEWBORN	347.0	320.0	...	356.0	356.0	337.0	353.0	378.0	
URGENT	19.0	30.0	...	98.0	101.0	104.0	111.0	111.0	

ADMISSION_HOUR	19	20	21	22	23
ADMISSION_TYPE					
ELECTIVE	21.0	15.0	11.0	1.0	8.0
EMERGENCY	2733.0	2907.0	2722.0	2500.0	2238.0
NEWBORN	347.0	403.0	348.0	378.0	349.0
URGENT	123.0	96.0	84.0	53.0	36.0

[4 rows x 24 columns]

Pivoted Data Types:

ADMISSION_HOUR	
0	float64
1	float64
2	float64
3	float64

```

4    float64
5    float64
6    float64
7    float64
8    float64
9    float64
10   float64
11   float64
12   float64
13   float64
14   float64
15   float64
16   float64
17   float64
18   float64
19   float64
20   float64
21   float64
22   float64
23   float64
dtype: object
Query Execution Time: 3.44 seconds

```

5.3.2 Conclusion

By visualizing the heatmap, we can see that the EMERGENCY Type of Admission is very correlated with Hour of Admission, being much more common in the evening/night 16PM - 23PM. The Type of Admission ELECTIVE is also correlated with Hour of Admission, being very common at 7AM.

5.3.3 Type of Admission & Death

```

[57]: start_time = time.time()

# 1. Group by ADMISSION_TYPE and EXPIRE_FLAG
adm_type_expire_df = (
    filtered_ddf.groupby(["ADMISSION_TYPE", "EXPIRE_FLAG"])
    .size()
    .compute()
    .reset_index(name="count")
)

# 2. Map EXPIRE_FLAG to readable outcome labels
adm_type_expire_df["Outcome"] = adm_type_expire_df["EXPIRE_FLAG"].map({0: "Survived", 1: "Expired"})

# 3. Pivot for heatmap
if not adm_type_expire_df.empty:
    heatmap_data = (

```

```

        adm_type_expire_df
        .pivot(index="ADMISSION_TYPE", columns="Outcome", values="count")
        .fillna(0)
        .astype(float)
    )

    # 4. Plot
    plt.figure(figsize=(8, 6))
    sns.heatmap(
        heatmap_data,
        annot=True,
        fmt=".0f",
        linewidths=.5,
        cmap="coolwarm",
        cbar_kws={"label": "Number of Patients"}
    )
    plt.xlabel("Patient Outcome During Hospital Stay")
    plt.ylabel("Admission Type")
    plt.title("Heatmap of Patient Outcomes by Admission Type")
    plt.yticks(rotation=0)
    plt.tight_layout()
    plt.show()


    # 5. Print pivoted table
    print("\nPivoted Data for Heatmap:")
    print(heatmap_data)
else:
    print("No data returned for outcome heatmap analysis.")

end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")

```

/usr/local/lib/python3.11/dist-packages/dask/dataframe/dask_expr/_groupby.py:1562: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```
self._meta = self.obj._meta.groupby(
```



Pivoted Data for Heatmap:

Outcome	Expired	Survived
ADMISSION_TYPE		
ELECTIVE	1990.0	5236.0
EMERGENCY	18992.0	22609.0
NEWBORN	65.0	7801.0
URGENT	689.0	657.0

Query Execution Time: 3.20 seconds

5.3.4 Conclusion

The clearest conclusion from this HeatMap is that patients who have EMERGENCY as their Type of Admission tend to EXPIRE much more frequently.

5.3.5 First Care Unit & Diagnosis

```
[58]: start_time = time.time()

# 1. Force PRIMARY_ICD9_CODE to string
```

```

filtered_ddf = filtered_ddf.
    ↪assign(PRIMARY_ICD9_CODE=filtered_ddf["PRIMARY_ICD9_CODE"].astype(str))

# 2. Get top 20 ICD-9 codes as strings
top_icd9_codes = (
    filtered_ddf["PRIMARY_ICD9_CODE"]
    .value_counts()
    .compute()
    .nlargest(20)
    .index
    .tolist()
)

# 3. Filter rows where ICD-9 is in top list and FIRST_CAREUNIT is not null
careunit_icd9_ddf = filtered_ddf[
    (filtered_ddf["PRIMARY_ICD9_CODE"].isin(top_icd9_codes)) &
    (filtered_ddf["FIRST_CAREUNIT"].notnull())
][["FIRST_CAREUNIT", "PRIMARY_ICD9_CODE"]]

# 4. Group and count
careunit_icd9_df = (
    careunit_icd9_ddf
    .groupby(["FIRST_CAREUNIT", "PRIMARY_ICD9_CODE"])
    .size()
    .compute()
    .reset_index(name="count")
)

# 5. Pivot and plot
if not careunit_icd9_df.empty:
    heatmap_data = (
        careunit_icd9_df
        .pivot(index="FIRST_CAREUNIT", columns="PRIMARY_ICD9_CODE",
    ↪values="count")
        .fillna(0)
        .astype(float)
    )

    plt.figure(figsize=(18, 8))
    sns.heatmap(
        heatmap_data,
        annot=False,
        fmt=".0f",
        linewidths=.5,
        cmap="Spectral",
        cbar_kws={'label': 'Number of Patients'}
    )

```



```

plt.xlabel("Primary ICD-9 Code (Top 20)")
plt.ylabel("First Care Unit")
plt.title("Heatmap of Admission Counts by First Care Unit and Top 20_
↳Primary ICD-9 Codes")
plt.xticks(rotation=90)
plt.yticks(rotation=0)
plt.tight_layout()
plt.show()

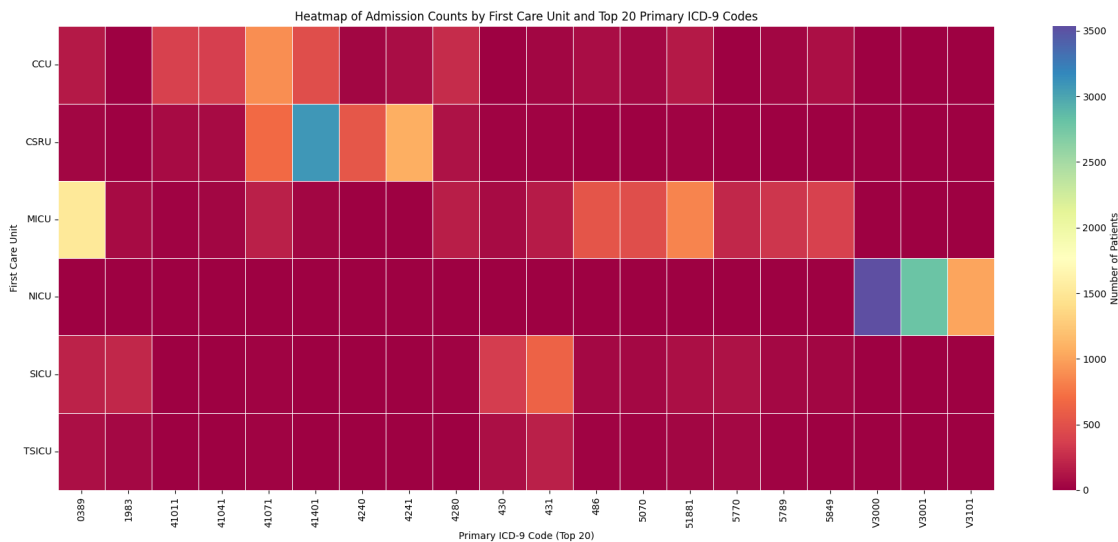
print("\nPivoted Data for Heatmap:")
print(heatmap_data)
else:
    print("No data returned for First Care Unit vs Top ICD-9 Code heatmap.")

end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")

```

/usr/local/lib/python3.11/dist-packages/dask/dataframe/dask_expr/_groupby.py:1562: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```
self._meta = self.obj._meta.groupby(
```



Pivoted Data for Heatmap:

```

PRIMARY_ICD9_CODE    0389    1983    41011    41041    41071    41401    4240    4241  \
FIRST_CAREUNIT

```

CCU	150.0	4.0	379.0	371.0	885.0	476.0	37.0	80.0
CSRU	31.0	4.0	62.0	67.0	684.0	3092.0	550.0	1075.0
MICU	1509.0	59.0	27.0	28.0	187.0	38.0	4.0	9.0
NICU	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
SICU	207.0	226.0	1.0	2.0	27.0	6.0	1.0	3.0
TSICU	91.0	50.0	4.0	3.0	20.0	5.0	0.0	1.0

PRIMARY_ICD9_CODE	4280	430	431	486	5070	51881	5770	5789	\
FIRST_CAREUNIT									
CCU	251.0	13.0	39.0	74.0	45.0	141.0	13.0	29.0	
CSRU	107.0	18.0	14.0	11.0	12.0	22.0	7.0	9.0	
MICU	173.0	64.0	163.0	535.0	474.0	835.0	231.0	317.0	
NICU	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
SICU	23.0	351.0	627.0	50.0	43.0	86.0	99.0	43.0	
TSICU	8.0	96.0	188.0	14.0	22.0	40.0	49.0	17.0	

PRIMARY_ICD9_CODE	5849	V3000	V3001	V3101
FIRST_CAREUNIT				
CCU	85.0	0.0	0.0	0.0
CSRU	7.0	0.0	0.0	0.0
MICU	384.0	0.0	0.0	0.0
NICU	0.0	3534.0	2792.0	1008.0
SICU	34.0	0.0	0.0	0.0
TSICU	13.0	0.0	0.0	0.0

Query Execution Time: 6.32 seconds

5.3.6 Conclusion

The strongest connections between First Care Unit and Diagnosis are V3000 (Single liveborn, delivered in hospital, without cesarean section) and NICU (Neonatal Intensive Care Unit) and 41401 (Coronary atherosclerosis of native coronary artery) and CSRU (Cardiac Surgery Recovery Unit).

5.3.7 First Care Unit & Death

```
[61]: start_time = time.time()

# 1. Group by FIRST_CAREUNIT and EXPIRE_FLAG
careunit_expire_df = (
    filtered_ddf.groupby(["FIRST_CAREUNIT", "EXPIRE_FLAG"])
    .size()
    .compute()
    .reset_index(name="count")
)

# 2. Map EXPIRE_FLAG to readable labels
careunit_expire_df["Outcome"] = careunit_expire_df["EXPIRE_FLAG"].map({0: "Survived", 1: "Expired"})
```

```

# 3. Pivot for heatmap
if not careunit_expire_df.empty:
    heatmap_data = (
        careunit_expire_df
        .pivot(index="FIRST_CAREUNIT", columns="Outcome", values="count")
        .fillna(0)
        .astype(float)
    )

# 4. Plot
plt.figure(figsize=(10, 7))
sns.heatmap(
    heatmap_data,
    annot=True,
    fmt=".0f",
    linewidths=.5,
    cmap="Spectral",
    cbar_kws={'label': 'Number of Patients'}
)
plt.xlabel("Patient Outcome During Hospital Stay")
plt.ylabel("First Care Unit")
plt.title("Heatmap of Patient Outcomes by First Care Unit")
plt.yticks(rotation=0)
plt.tight_layout()
plt.show()


# 5. Print pivot table
print("\nPivoted Data for Heatmap:")
print(heatmap_data)
else:
    print("No data returned for First Care Unit vs Outcome heatmap.")

end_time = time.time()
execution_time = end_time - start_time
print(f"Query Execution Time: {execution_time:.2f} seconds")

```

/usr/local/lib/python3.11/dist-packages/dask/dataframe/dask_expr/_groupby.py:1562: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```
self._meta = self.obj._meta.groupby(
```



Pivoted Data for Heatmap:

Outcome	Expired	Survived
FIRST_CAREUNIT		
CCU	3432.0	3700.0
CSRU	2475.0	6612.0
MICU	10152.0	9188.0
NICU	67.0	8008.0
SICU	3619.0	4755.0
TSICU	1991.0	4040.0

Query Execution Time: 4.98 seconds

5.3.8 Conclusion

In half of the First Care Units, the ratio between patients that expired and survived is fairly equal (and close to 1). In the other half, which include the units NICU, TSICU and CSRU, the number of patients that survived is much higher than the number of patients that expired.

6 XGBoost Classifier Predictor Model

In this section we use a XGBoost Classifier model to predict the ICU Length of Stay, that we divide in bins of two days ([1-3] days, [3-5] days, ...). The window size we chose was 1 day (24

hours), given that most data is collected in that window and most records have a value of **ICU Length of Stay** lower than 3 days, which means a bigger window size would exclude most of the data. We divided this section of the work in 2 parts: 1. Dividing the dataset, Categorizing the Target and Implementing Custom Sample Weights 2. Pre-Processing the features 3. Applying the model 4. Visualizing the results

The discussion of results was done in the final conclusions.

6.1 Dividing the Dataset, Categorizing the Target and Implementing Custom Sample Weights

Here we will divide the dataset in Train, Test and Validation. We will do this by: 1. Eliminating the features that are only collected after the first 24 hours (our **window size**) or are not relevant 2. Removing the records whose value for **ICU Length of Stay** is lower than 24 hours 3. Using the library **scikit-learn** to perform the division 4. Converting the target feature from **continuous** values to bins 5. Calculating the custom sample weights to ensure a less biased prediction towards the majority class 6. Normalizing the sample weights

```
[64]: # --- Configuration ---
TARGET_VARIABLE = 'ICU_LOS'
WINDOW_DAYS = 1.0
MAX_ICU_LOS_DAYS = 25.0
BIN_WIDTH = 2.0
TEST_SIZE = 0.10
VALIDATION_SIZE = 0.10
CUSTOM_WEIGHT_ALPHA = 0.8

print("--- Block 1: Dividing Dataset, Target Categorization & CUSTOM Sample_
↳Weights ---")
start_time = time.time()

try:
    ddf = filtered_ddf
    original_row_count = ddf.shape[0].compute()
    print(f"Original row count: {original_row_count}")

    ddf = ddf[ddf[TARGET_VARIABLE] >= WINDOW_DAYS]
    rows_after_min_los_filter = ddf.shape[0].compute()
    print(f"Rows remaining after MIN LOS filter (>= {WINDOW_DAYS} days):_
↳{rows_after_min_los_filter}")

    ddf = ddf[ddf[TARGET_VARIABLE] <= MAX_ICU_LOS_DAYS]
    rows_after_max_los_filter = ddf.shape[0].compute()
    print(f"Rows remaining after MAX LOS filter (<= {MAX_ICU_LOS_DAYS} days):_
↳{rows_after_max_los_filter}")

    if rows_after_max_los_filter > 0:
        features_to_exclude = [
```

```

        'SUBJECT_ID', 'HADM_ID', 'ICUSTAY_ID', 'DOB',
        'ADMITTIME', 'DISCHTIME', 'ICU_INTIME', 'ICU_OUTTIME',
        'PRIMARY_ICD9_TITLE', 'EXPIRE_FLAG'
    ]

    all_columns = ddf.columns
    features_for_training_initial = [
        col for col in all_columns
        if col not in features_to_exclude and col != TARGET_VARIABLE
    ]

    print(f"Initial features considered for training:␣
↪{features_for_training_initial}")

    X = ddf[features_for_training_initial]
    y_continuous = ddf[TARGET_VARIABLE]

    X_pd = X.compute()
    y_continuous_pd = y_continuous.compute()

    if len(X_pd) < 3:
        raise ValueError("Not enough data to split.")

    train_val_size = 1.0 - TEST_SIZE
    X_train_val, X_test, y_train_val_cont, y_test_cont = train_test_split(
        X_pd, y_continuous_pd, test_size=TEST_SIZE, random_state=42,␣
↪shuffle=True
    )

    val_split_ratio = VALIDATION_SIZE / train_val_size if train_val_size >␣
↪0 else 0

    if len(X_train_val) < 2 or val_split_ratio == 0:
        X_train, X_val = X_train_val, pd.DataFrame()
        y_train_cont, y_val_cont = y_train_val_cont, pd.Series(dtype=float)
    else:
        X_train, X_val, y_train_cont, y_val_cont = train_test_split(
            X_train_val, y_train_val_cont, test_size=val_split_ratio,␣
↪random_state=42, shuffle=True
        )

    def convert_los_to_categories(los_series, bin_width, min_los):
        adjusted_los = los_series - min_los
        categories = np.floor(adjusted_los / bin_width).astype(int)
        return np.maximum(0, categories)

```

```

    y_train_cat = convert_los_to_categories(y_train_cont, BIN_WIDTH,
↪WINDOW_DAYS)
    y_val_cat = convert_los_to_categories(y_val_cont, BIN_WIDTH,
↪WINDOW_DAYS)
    y_test_cat = convert_los_to_categories(y_test_cont, BIN_WIDTH,
↪WINDOW_DAYS)

    all_cats = pd.concat([
        pd.Series(y_train_cat),
        pd.Series(y_val_cat),
        pd.Series(y_test_cat)
    ]).dropna().astype(int)

    if all_cats.empty:
        num_classes = 1
        max_cat_overall = 0
        train_sample_weights = np.array([])
    else:
        max_cat_overall = all_cats.max()
        num_classes = max_cat_overall + 1

        if len(y_train_cat) > 0:
            print(f"\n--- Calculating CUSTOM Sample Weights,
↪(alpha={CUSTOM_WEIGHT_ALPHA}) ---")
            class_counts = Counter(y_train_cat)
            total_samples = len(y_train_cat)

            class_weights_map = {
                cls: 1.0 / (count ** CUSTOM_WEIGHT_ALPHA) if count > 0 else
↪1.0
                for cls, count in class_counts.items()
            }

            train_sample_weights = np.array([class_weights_map.get(cls, 1.
↪0) for cls in y_train_cat])
            train_sample_weights = (train_sample_weights / np.
↪sum(train_sample_weights)) * total_samples

            print(f"Custom class weights map: {class_weights_map}")
            print(f"Computed CUSTOM sample weights. Shape:
↪{train_sample_weights.shape}")
        else:
            train_sample_weights = np.array([])

        print(f"\nTarget variable converted to {num_classes} categories (0 to
↪{max_cat_overall}).")

```

```

print("\nDataset shapes:")
print(f"X_train: {X_train.shape}, y_train_cat: {y_train_cat.shape}")
if not X_val.empty:
    print(f"X_val: {X_val.shape}, y_val_cat: {y_val_cat.shape}")
else:
    print("X_val is empty.")
print(f"X_test: {X_test.shape}, y_test_cat: {y_test_cat.shape}")

if original_row_count > 0:
    percentage_kept = (rows_after_max_los_filter / original_row_count)
    ↪* 100
    print(f"\nPercentage of original records kept after ALL filtering:
    ↪{percentage_kept:.2f}%")
else:
    raise ValueError("No data left after filtering.")

except Exception as e:
    print(f"Error: {e}")
    X_train, X_val, X_test = pd.DataFrame(), pd.DataFrame(), pd.DataFrame()
    y_train_cat = y_val_cat = y_test_cat = pd.Series(dtype=int)
    train_sample_weights = np.array([])
    num_classes = 1
    features_for_training_initial = []

end_time = time.time()
print(f"\nExecution time: {end_time - start_time:.2f} seconds")

```

--- Block 1: Dividing Dataset, Target Categorization & CUSTOM Sample Weights ---

Original row count: 58039

Rows remaining after MIN LOS filter (≥ 1.0 days): 46321

Rows remaining after MAX LOS filter (≤ 25.0 days): 44439

Initial features considered for training: ['GENDER', 'AGE_AT_ADMISSION', 'ADMISSION_TYPE', 'ADMISSION_LOCATION', 'INSURANCE', 'ETHNICITY', 'ADMISSION_DIAGNOSIS_TEXT', 'ADMISSION_HOUR', 'FIRST_CAREUNIT', 'PRIMARY_ICD9_CODE']

--- Calculating CUSTOM Sample Weights (alpha=0.8) ---

Custom class weights map: {0: 0.0003634369074874825, 4: 0.003567379116988738, 1: 0.0008840502035876041, 3: 0.0025809482616361177, 2: 0.0016579799257901674, 8: 0.008045657161370626, 6: 0.005605339521480853, 9: 0.00995539059449496, 10: 0.012264988137738082, 5: 0.004555275341602907, 7: 0.0068874041958582035, 11: 0.012945881104681681}

Computed CUSTOM sample weights. Shape: (35551,)

Target variable converted to 12 categories (0 to 11).

Dataset shapes:


```
X_train: (35551, 10), y_train_cat: (35551,)
X_val: (4444, 10), y_val_cat: (4444,)
X_test: (4444, 10), y_test_cat: (4444,)
```

Percentage of original records kept after ALL filtering: 76.57%

Execution time: 36.40 seconds

6.2 Pre-Processing the Features

Here we will pre-process the features before applying the predictor model. We will do this by: 1. Defining the features to be used in training and their respective category 2. Feature engineer the features based on their category: - Scale for Continuous Numerical features - Encode using sin and cos for Cyclical Numerical features - One-hot Encode for Low/Medium Cardinality Categorical Features - Dictifier Transformation followed by a FeatureHasher Transformation for High Cardinality Categorical features

```
[65]: # --- Record Start Time ---
start_time = time.time()
print("\n--- Block 2: Pre-Processing the Features ---")

# If X_train is empty, skip preprocessing
if 'X_train' in globals() and not X_train.empty:

    # --- 1. Define Feature Lists (based on X_train columns) ---

    _features_for_training = [
        'AGE_AT_ADMISSION',      # Numerical continuous
        'ADMISSION_TYPE',        # Categorical low cardinality
        'INSURANCE',              # Categorical low cardinality
        'PRIMARY_ICD9_CODE',      # Categorical high cardinality
        'FIRST_CAREUNIT',         # Categorical low cardinality
        'ADMISSION_HOUR'         # Numerical cyclical
    ]

    # Filter features to only those present in the training data from Block 1
    features_for_training_final = [f for f in _features_for_training if f in
    ↪X_train.columns]
    print(f"Features selected for preprocessing: {features_for_training_final}")

    numerical_cont_features = [f for f in ['AGE_AT_ADMISSION'] if f in
    ↪features_for_training_final]
    numerical_cycl_features = [f for f in ['ADMISSION_HOUR'] if f in
    ↪features_for_training_final] # e.g., 0-23
    categorical_low_card_features = [f for f in ['ADMISSION_TYPE', 'INSURANCE',
    ↪'FIRST_CAREUNIT'] if f in features_for_training_final]
```

```

categorical_high_card_features = [f for f in ['PRIMARY_ICD9_CODE'] if f in_
↪features_for_training_final]

print(f"Continuous numerical features: {numerical_cont_features}")
print(f"Cyclical numerical features: {numerical_cycl_features}")
print(f"Low/Medium cardinality categorical features:
↪{categorical_low_card_features}")
print(f"High cardinality categorical features:
↪{categorical_high_card_features}")

# --- 2. Feature Engineering & Preprocessing Pipeline Components ---

# a) Continuous Numerical Features: Scale
numerical_cont_transformer = Pipeline(steps=[
    ('scaler', StandardScaler())
])

# b) Cyclical Numerical Features: Encode hour using sine and cosine
def sin_transformer(X_in):
    # Ensure input is 2D for FunctionTransformer
    X_proc = X_in.copy()
    if isinstance(X_proc, pd.Series): X_proc = X_proc.to_frame()
    elif X_proc.ndim == 1: X_proc = X_proc.reshape(-1, 1)
    return np.sin(2 * np.pi * X_proc / 24.0)

def cos_transformer(X_in):
    X_proc = X_in.copy()
    if isinstance(X_proc, pd.Series): X_proc = X_proc.to_frame()
    elif X_proc.ndim == 1: X_proc = X_proc.reshape(-1, 1)
    return np.cos(2 * np.pi * X_proc / 24.0)

# c) Low/Medium Cardinality Categorical Features: One-hot encode
categorical_low_card_transformer = Pipeline(steps=[
    ('onehot', OneHotEncoder(handle_unknown='ignore', sparse_output=False))
])

# d) High Cardinality Categorical Features: Feature Hashing
# Custom Transformer for FeatureHasher Input
class Dictifier(BaseEstimator, TransformerMixin):
    def __init__(self, col_name):
        self.col_name = col_name
    def fit(self, X, y=None):
        return self
    def transform(self, X_in):
        # X_in is expected to be a DataFrame/Series for the specific column
        if isinstance(X_in, pd.DataFrame): series = X_in.iloc[:, 0]
        elif isinstance(X_in, pd.Series): series = X_in

```

```

        else: series = pd.Series(X_in.flatten()) # Fallback for numpy array
        return [{self.col_name: str(val)} for val in series]

# --- Create the list of transformers for ColumnTransformer ---
transformers_list = []

if numerical_cont_features:
    transformers_list.append(('num_cont', numerical_cont_transformer,
↪numerical_cont_features))
    if numerical_cycl_features: # Assuming only one cyclical feature
↪'ADMISSION_HOUR' for this example
        transformers_list.append(('hour_sin',
↪FunctionTransformer(sin_transformer, validate=False),
↪numerical_cycl_features))
        transformers_list.append(('hour_cos',
↪FunctionTransformer(cos_transformer, validate=False),
↪numerical_cycl_features))
    if categorical_low_card_features:
        transformers_list.append(('cat_low', categorical_low_card_transformer,
↪categorical_low_card_features))

if categorical_high_card_features:
    n_hash_features = 50 # Number of features for the hasher, adjust as
↪needed
    for i, col_name in enumerate(categorical_high_card_features):
        print(f" - Adding hasher for high-cardinality feature: {col_name}")
        transformer_name = f'cat_high_{col_name.replace(" ", "_").lower()}'
↪# Unique name
        high_card_pipeline = Pipeline(steps=[
            ('dictifier', Dictifier(col_name=col_name)),
            ('hasher', FeatureHasher(n_features=n_hash_features,
↪input_type='dict'))
        ])
        transformers_list.append((transformer_name, high_card_pipeline,
↪[col_name])) # Pass as list
    else:
        print("No high cardinality features specified or found for hashing.")

# --- Define the main ColumnTransformer ---
if transformers_list:
    preprocessor = ColumnTransformer(
        transformers=transformers_list,
        remainder='drop', # Drop any columns not specified in
↪features_for_training_final
        n_jobs=-1
    )

```

```

        print("Preprocessor defined.")
    else:
        print("No features to preprocess. Preprocessor not created.")
        preprocessor = 'passthrough' # Or handle as an error/empty pipeline

else:
    print("X_train is empty or not defined. Skipping feature preprocessing.")
    features_for_training_final = []
    preprocessor = 'passthrough' # To avoid error in next block

# --- Record End Time ---
end_time = time.time()
execution_time = end_time - start_time
print(f"\nExecution time: {execution_time:.2f} seconds")

```

--- Block 2: Pre-Processing the Features ---

Features selected for preprocessing: ['AGE_AT_ADMISSION', 'ADMISSION_TYPE', 'INSURANCE', 'PRIMARY_ICD9_CODE', 'FIRST_CAREUNIT', 'ADMISSION_HOUR']
 Continuous numerical features: ['AGE_AT_ADMISSION']
 Cyclical numerical features: ['ADMISSION_HOUR']
 Low/Medium cardinality categorical features: ['ADMISSION_TYPE', 'INSURANCE', 'FIRST_CAREUNIT']
 High cardinality categorical features: ['PRIMARY_ICD9_CODE']
 - Adding hasher for high-cardinality feature: PRIMARY_ICD9_CODE
 Preprocessor defined.

Execution time: 0.00 seconds

6.3 Applying the XGBoost Classifier Model

Here we will apply the XGBoost model to the dataset. We will do this by:

1. Defining the model pipeline
2. Training the model
3. Making predictions using the trained model
4. Post-Process predictions

```

[67]: # --- Record Start Time ---
start_time_block3_weighted = time.time()
print("\n--- Block 3: Applying the XGBoost Model (Classification) - MANUAL_
↳ PREDICTION FALLBACK & SAMPLE WEIGHTS ---")

# Ensure necessary variables are available
if ('X_train' not in globals() or X_train.empty or
    'y_train_cat' not in globals() or y_train_cat.empty or
    'preprocessor' not in globals() or
    'features_for_training_final' not in globals() or

```

```

'num_classes' not in globals() or
'train_sample_weights' not in globals()):
    print("Necessary data (X_train, y_train_cat, preprocessor, features,
↳num_classes, train_sample_weights) not available or y_train_cat is empty.
↳Skipping model training.")
    y_pred_val_final = pd.Series(dtype=int)
    y_pred_test_final = pd.Series(dtype=int)
    xgb_model_pipeline = None
    fitted_pipeline_preprocessor = None
    fitted_pipeline_classifier = None
else:
    if y_train_cat.empty:
        print(" y_train_cat is EMPTY. Cannot proceed with model training.")
        xgb_model_pipeline = None
        fitted_pipeline_preprocessor = None
        fitted_pipeline_classifier = None
    elif train_sample_weights is None or train_sample_weights.size == 0:
        print(" train_sample_weights is None or empty. Cannot proceed with
↳weighted training.")
        xgb_model_pipeline = None
        fitted_pipeline_preprocessor = None
        fitted_pipeline_classifier = None
    else:
        print(f" y_train_cat.shape: {y_train_cat.shape}, dtype: {y_train_cat.
↳dtype}")
        if y_train_cat.isnull().any():
            print(f" WARNING: y_train_cat contains NaNs! Count: {y_train_cat.
↳isnull().sum()}")
            y_train_cat = y_train_cat.astype(int) # Ensure it's int
            print(f" Type of target (y_train_cat): {type_of_target(y_train_cat)}")
            print(f" y_train_cat.min(): {y_train_cat.min()}, y_train_cat.max():
↳{y_train_cat.max()}")

            current_X_train_cols = X_train.columns.tolist()
            valid_features_for_model = [f for f in features_for_training_final if f
↳in current_X_train_cols]
            if not valid_features_for_model:
                print("WARNING: No valid features for model found in X_train.
↳Preprocessing might be incorrect or use no features.")

            X_train_processed_standalone = None
            try:
                if preprocessor == 'passthrough':
                    X_train_processed_standalone =
↳X_train[valid_features_for_model].copy() if valid_features_for_model else
↳X_train.copy()

```

```

        elif hasattr(preprocessor, 'fit_transform'):
            temp_preprocessor_standalone = clone(preprocessor)
            print(" Fitting and transforming X_train with a cloned_
↳preprocessor for standalone test...")
            X_train_processed_standalone = temp_preprocessor_standalone.
↳fit_transform(X_train[valid_features_for_model])
            print(f" X_train_processed_standalone shape:_
↳{X_train_processed_standalone.shape}")
        else:
            print(" Preprocessor is not 'passthrough' and does not have_
↳'fit_transform'. Cannot preprocess for standalone test.")
            except Exception as e_preprocess_standalone:
                print(f" ERROR during X_train preprocessing for standalone test:_
↳{e_preprocess_standalone}")

        if X_train_processed_standalone is not None:
            standalone_xgb = XGBClassifier(objective='multi:softmax',_
↳num_class=num_classes, n_estimators=50, learning_rate=0.1, max_depth=3,_
↳random_state=42, n_jobs=-1, eval_metric='mlogloss')
            try:
                # Pass sample_weight to standalone fit
                standalone_xgb.fit(X_train_processed_standalone, y_train_cat,_
↳sample_weight=train_sample_weights)
                print(" Standalone XGBClassifier fitted successfully with_
↳sample_weight!")
            except Exception as e_standalone:
                print(f" ERROR fitting standalone XGBClassifier with_
↳sample_weight: {e_standalone}")
            else:
                print(" Skipping Standalone XGBClassifier test as X_train could_
↳not be processed for it.")

        # preprocessor is the definition from Block 2. Pipeline will fit it.
        xgb_model_pipeline = Pipeline(steps=[
            ('preprocessor', preprocessor),
            ('classifier', XGBClassifier(objective='multi:softmax',_
↳num_class=num_classes, n_estimators=100, learning_rate=0.1, max_depth=5,_
↳subsample=0.8, colsample_bytree=0.8, random_state=42, n_jobs=-1,_
↳eval_metric='mlogloss'))
        ])

        fitted_pipeline_preprocessor = None
        fitted_pipeline_classifier = None
        pipeline_fit_successful = False

    try:

```

```

        print(f" Attempting to fit the pipeline on_
↳X_train[valid_features_for_model] (shape: {X_train[valid_features_for_model].
↳shape if valid_features_for_model else X_train.shape}) and y_train_cat...")
        fit_params = {'classifier__sample_weight': train_sample_weights}
        xgb_model_pipeline.fit(X_train, y_train_cat, **fit_params)
        print(" Pipeline fitted successfully with sample_weight!")
        pipeline_fit_successful = True
        fitted_pipeline_preprocessor = xgb_model_pipeline.
↳named_steps['preprocessor']
        fitted_pipeline_classifier = xgb_model_pipeline.
↳named_steps['classifier']

    except Exception as e_pipeline_fit:
        print(f" ERROR during pipeline.fit() with sample_weight:_
↳{e_pipeline_fit}")
        xgb_model_pipeline = None # Ensure pipeline is None if fit fails

    # Initialize prediction variables
    y_pred_val_final = pd.Series(dtype=int)
    y_pred_test_final = pd.Series(dtype=int)

    if pipeline_fit_successful and fitted_pipeline_preprocessor and_
↳fitted_pipeline_classifier:
        # Validation set predictions
        if 'X_val' in globals() and not X_val.empty:
            if 'y_val_cat' in globals() and not y_val_cat.empty: # Check if_
↳there's a target to compare
                try:
                    # Use the same features for transform as were used for_
↳training the preprocessor
                    X_val_to_transform = X_val # Pass the X_val with all_
↳original features
                    X_val_processed = fitted_pipeline_preprocessor.
↳transform(X_val_to_transform)
                    print(f" X_val_processed shape: {X_val_processed.
↳shape}")
                    y_pred_val_cat_manual = fitted_pipeline_classifier.
↳predict(X_val_processed)
                    y_pred_val_final = pd.Series(y_pred_val_cat_manual,_
↳index=X_val.index)
                    print(f" Manual predictions on validation set_
↳successful. Shape: {y_pred_val_final.shape}")
                    except Exception as e_manual_pred_val:
                        print(f" ERROR during manual prediction on_
↳validation set: {e_manual_pred_val}")
                else:

```

```

        print(" Validation set target (y_val_cat) is empty.␣
↳Skipping manual validation predictions.")
    else:
        print(" Validation set (X_val) is empty or not available.␣
↳Skipping manual validation predictions.")

    # Test set predictions
    if 'X_test' in globals() and not X_test.empty:
        if 'y_test_cat' in globals() and not y_test_cat.empty: # Check␣
↳if there's a target to compare
            try:
                X_test_to_transform = X_test # Pass the X_test with all␣
↳original features
                X_test_processed = fitted_pipeline_preprocessor.
↳transform(X_test_to_transform)
                print(f"    X_test_processed shape: {X_test_processed.
↳shape}")
                y_pred_test_cat_manual = fitted_pipeline_classifier.
↳predict(X_test_processed)
                y_pred_test_final = pd.Series(y_pred_test_cat_manual,␣
↳index=X_test.index)
                print(f"    Manual predictions on test set successful.␣
↳Shape: {y_pred_test_final.shape}")
            except Exception as e_manual_pred_test:
                print(f"    ERROR during manual prediction on test set:␣
↳{e_manual_pred_test}")
            else:
                print(" Test set target (y_test_cat) is empty. Skipping␣
↳manual test predictions.")
            else:
                print(" Test set (X_test) is empty or not available. Skipping␣
↳manual test predictions.")

    elif xgb_model_pipeline is not None:
        try:
            if 'X_val' in globals() and not X_val.empty and 'y_val_cat' in␣
↳globals() and not y_val_cat.empty:
                y_pred_val_cat_pipeline = xgb_model_pipeline.predict(X_val)␣
↳# Use full X_val
                y_pred_val_final = pd.Series(y_pred_val_cat_pipeline,␣
↳index=X_val.index)
                print(f" Pipeline predictions on validation set successful.
↳ Shape: {y_pred_val_final.shape}")
                if 'X_test' in globals() and not X_test.empty and 'y_test_cat'␣
↳in globals() and not y_test_cat.empty:

```



```

        y_pred_test_cat_pipeline = xgb_model_pipeline.
↪predict(X_test) # Use full X_test
        y_pred_test_final = pd.Series(y_pred_test_cat_pipeline,
↪index=X_test.index)
        print(f" Pipeline predictions on test set successful.
↪Shape: {y_pred_test_final.shape}")
        except Exception as e_pipeline_predict:
            print(f" ERROR during pipeline.predict() fallback:
↪{e_pipeline_predict}")
        else:
            print("\nPipeline did not fit successfully. No predictions will be
↪made.")

# --- Record End Time ---
end_time = time.time()
execution_time = end_time - start_time
print(f"\nXGBoost training & prediction execution time (MANUAL PREDICTION
↪FALLBACK & SAMPLE WEIGHTS block): {execution_time:.2f} seconds")

```

--- Block 3: Applying the XGBoost Model (Classification) - MANUAL PREDICTION FALLBACK & SAMPLE WEIGHTS ---

```

y_train_cat.shape: (35551,), dtype: int64
Type of target (y_train_cat): multiclass
y_train_cat.min(): 0, y_train_cat.max(): 11
Fitting and transforming X_train with a cloned preprocessor for standalone
test...
X_train_processed_standalone shape: (35551, 68)
Standalone XGBClassifier fitted successfully with sample_weight!
Attempting to fit the pipeline on X_train[valid_features_for_model] (shape:
(35551, 6)) and y_train_cat...
Pipeline fitted successfully with sample_weight!
X_val_processed shape: (4444, 68)
Manual predictions on validation set successful. Shape: (4444,)
X_test_processed shape: (4444, 68)
Manual predictions on test set successful. Shape: (4444,)

```

XGBoost training & prediction execution time (MANUAL PREDICTION FALLBACK & SAMPLE WEIGHTS block): 283.36 seconds

6.4 Vizualizing the Results

To end we use statistical measures and adequate plots to visualize the prediction results.

```

[68]: # --- Record Start Time ---
start_time = time.time()
print("\n--- Block 4: Evaluating and Visualizing Results ---")

```

```

# Check if necessary data for evaluation is available
if ('y_val_cat' not in globals() or
    'y_test_cat' not in globals() or
    'y_pred_val_final' not in globals() or
    'y_pred_test_final' not in globals() or
    'num_classes' not in globals()):
    print("Necessary data for evaluation/visualization is missing. Skipping.")
else:
    # Define descriptive labels for your bins for plotting
    bin_labels_for_plots = []
    if 'WINDOW_DAYS' in globals() and 'BIN_WIDTH' in globals() and num_classes_
    ↪ > 0 :
        for i in range(num_classes):
            lower_bound = WINDOW_DAYS + i * BIN_WIDTH
            upper_bound = WINDOW_DAYS + (i + 1) * BIN_WIDTH
            bin_labels_for_plots.append(f"[{lower_bound:.1f}-{upper_bound:.
    ↪ 1f})")
    else: # Fallback labels
        bin_labels_for_plots = [f"Cat {i}" for i in range(num_classes)]

    print(f"Using bin labels for plots: {bin_labels_for_plots}")

    # --- 1. Evaluate the Model ---
    # Validation Set Metrics (if y_val_cat and y_pred_val_final are not empty)
    if not y_val_cat.empty and not y_pred_val_final.empty:
        # Ensure labels in classification report and confusion matrix cover all
    ↪ possible classes
        unique_labels_present_val = np.union1d(y_val_cat.unique(),
    ↪ y_pred_val_final.unique())
        # Ensure all labels from 0 to num_classes-1 are considered for
    ↪ consistency if some classes have 0 instances
        report_labels_val = list(range(num_classes))

        accuracy_val = accuracy_score(y_val_cat, y_pred_val_final)
        kappa_val = cohen_kappa_score(y_val_cat, y_pred_val_final,
    ↪ labels=report_labels_val if report_labels_val else None)
        report_val = classification_report(y_val_cat, y_pred_val_final,
    ↪ labels=report_labels_val, target_names=bin_labels_for_plots, zero_division=0)

        print("\n--- Validation Set Metrics (XGBoost - Categorical) ---")
        print(f"Accuracy: {accuracy_val:.3f}")
        print(f"Cohen's Kappa: {kappa_val:.3f}")
        print("Classification Report (Validation):\n", report_val)
    else:

```

```

        print("\nValidation data (actual or predicted) is empty. Skipping validation metrics.")

    # Test Set Metrics (if y_test_cat and y_pred_test_final are not empty)
    if not y_test_cat.empty and not y_pred_test_final.empty:
        unique_labels_present_test = np.union1d(y_test_cat.unique(),
        y_pred_test_final.unique())
        report_labels_test = list(range(num_classes))

        accuracy_test = accuracy_score(y_test_cat, y_pred_test_final)
        kappa_test = cohen_kappa_score(y_test_cat, y_pred_test_final,
        labels=report_labels_test if report_labels_test else None)
        report_test = classification_report(y_test_cat, y_pred_test_final,
        labels=report_labels_test, target_names=bin_labels_for_plots,
        zero_division=0)
        cm_test = confusion_matrix(y_test_cat, y_pred_test_final,
        labels=report_labels_test)

        print("\n--- Test Set Metrics (XGBoost - Categorical) ---")
        print(f"Accuracy: {accuracy_test:.3f}")
        print(f"Cohen's Kappa: {kappa_test:.3f}")
        print("Classification Report (Test):\n", report_test)
        # print("Confusion Matrix (Test):\n", cm_test)

    # --- 2. Visualize Results (Test Set) ---
    # a) Confusion Matrix Heatmap
    if num_classes > 0:
        plt.figure(figsize=(min(10, num_classes + 2), min(8, num_classes +
        1)))

        sns.heatmap(cm_test, annot=True, fmt='d', cmap='Blues',
                    xticklabels=bin_labels_for_plots,
                    yticklabels=bin_labels_for_plots)
        plt.xlabel("Predicted Category")
        plt.ylabel("Actual Category")
        plt.title("Confusion Matrix (Test Set - XGBoost)")
        plt.tight_layout()
        plt.show()

    # b) Actual vs. Predicted Class Distribution
    plt.figure(figsize=(max(8, num_classes * 0.8), 6))
    actual_counts = y_test_cat.value_counts().
    reindex(range(num_classes), fill_value=0).sort_index()
    predicted_counts = y_pred_test_final.value_counts().
    reindex(range(num_classes), fill_value=0).sort_index()

```

```

df_counts = pd.DataFrame({'Actual': actual_counts, 'Predicted':
↪predicted_counts})
    if not df_counts.empty:
        # Use the generated bin_labels_for_plots for the x-axis ticks
        df_counts.index = [bin_labels_for_plots[i] for i in df_counts.
↪index if i < len(bin_labels_for_plots)]

        df_counts.plot(kind='bar', width=0.8)
        plt.title("Actual vs. Predicted Class Distribution (Test Set)")
        plt.xlabel("ICU LOS Category")
        plt.ylabel("Number of Cases")
        plt.xticks(rotation=45, ha="right")
        plt.legend()
        plt.grid(axis='y', linestyle='--')
        plt.tight_layout()
        plt.show()
    else:
        print("Cannot plot class distribution: Counts data is empty.")
    else:
        print("Number of classes is 0, cannot generate plots.")
    else:
        print("\nTest data (actual or predicted) is empty. Skipping test_
↪metrics and visualization.")

# --- Record End Time ---
end_time = time.time()
execution_time = end_time - start_time
print(f"\nExecution time: {execution_time:.2f} seconds")

```

--- Block 4: Evaluating and Visualizing Results ---

Using bin labels for plots: ['[1.0-3.0)', '[3.0-5.0)', '[5.0-7.0)', '[7.0-9.0)', '[9.0-11.0)', '[11.0-13.0)', '[13.0-15.0)', '[15.0-17.0)', '[17.0-19.0)', '[19.0-21.0)', '[21.0-23.0)', '[23.0-25.0)']

--- Validation Set Metrics (XGBoost - Categorical) ---

Accuracy: 0.438

Cohen's Kappa: 0.083

Classification Report (Validation):

	precision	recall	f1-score	support
[1.0-3.0)	0.61	0.71	0.66	2478
[3.0-5.0)	0.25	0.16	0.20	815
[5.0-7.0)	0.12	0.05	0.07	378
[7.0-9.0)	0.09	0.03	0.05	238
[9.0-11.0)	0.05	0.03	0.03	156

[11.0-13.0)	0.03	0.02	0.02	95
[13.0-15.0)	0.04	0.04	0.04	96
[15.0-17.0)	0.05	0.09	0.07	64
[17.0-19.0)	0.01	0.03	0.01	38
[19.0-21.0)	0.05	0.17	0.08	35
[21.0-23.0)	0.03	0.13	0.05	23
[23.0-25.0)	0.04	0.14	0.06	28
accuracy			0.44	4444
macro avg	0.11	0.13	0.11	4444
weighted avg	0.41	0.44	0.42	4444


--- Test Set Metrics (XGBoost - Categorical) ---

Accuracy: 0.426


Cohen's Kappa: 0.060

Classification Report (Test):

	precision	recall	f1-score	support
[1.0-3.0)	0.60	0.71	0.65	2446
[3.0-5.0)	0.21	0.13	0.16	859
[5.0-7.0)	0.06	0.03	0.04	388
[7.0-9.0)	0.06	0.02	0.03	221
[9.0-11.0)	0.04	0.02	0.03	141
[11.0-13.0)	0.08	0.04	0.05	119
[13.0-15.0)	0.09	0.12	0.10	75
[15.0-17.0)	0.03	0.05	0.04	59
[17.0-19.0)	0.04	0.13	0.06	38
[19.0-21.0)	0.03	0.09	0.04	34
[21.0-23.0)	0.04	0.09	0.06	35
[23.0-25.0)	0.03	0.10	0.04	29
accuracy			0.43	4444
macro avg	0.11	0.13	0.11	4444
weighted avg	0.38	0.43	0.40	4444



<Figure size 960x600 with 0 Axes>



Execution time: 2.44 seconds

7 Final Conclusions

7.0.1 Features Distribution

We were able to get some interesting conclusions from the distribution analysis for each feature, namely: - Most of patients only stay in ICU for 5 days or less, with a decreasing trend of the number of patients as the number of days increases.

7.0.2 Features Correlation

We were also able to draw interesting conclusions from the correlation between different features, namely: - The correlation between Diagnosis and ICU Length of Stay is evident, with Twin mate, liveborn, delivered in hospital, with cesarean section having a much higher mean value compared to other types of Diagnosis - Regarding the connection between Age and Length of ICU Stay, we can conclude that the vast majority of ICU stays are relatively short across all adult age groups, although the spread of the values of Length of ICU Stay appears to widen slightly for older patients - All Insurance systems have a similar distribution and mean for ICU Length of Stay, with a clear exception of Self Pay, whose average value is significantly lower. - Patients who have EMERGENCY as their Type of Admission tend to EXPIRE(die) much more

frequently. - NICU (Neonatal Intensive Care Unit) has the lowest number of deaths and MICU (Medical Intensive Care Unit) has the highest number of deaths.

7.0.3 Predictor Model

- The **accuracy** value in the **Validation** and in **Testing** set are similar and have medium-low values, which mean the model wasn't very successful at predicting the patients **ICU Length of Stay** during the validation phase
- The **Cohen's Kappa** value for both **Validation** and **Testing** is very low, which indicates the model is not much better than a random chance model at predicting the correct classes
- The **heatmap** shows that the model is very good at predicting correctly the majority class [1-3] days, but mostly fails at predicting other classes correctly
- The **bar plot** shows that the model predictions distribution is pretty similar to the actual data distribution, which is a good sign for possible future improvements
- In summary, the predictions for the majority class are accurate but for minority classes are not, which is a very common problem in Machine Learning prediction models.
- In future works, we could try to improve these bad results by performing better **feature engineering**, performing **fine-tuning** in the parameters of the predictor model, choose more wisely the features to use and possibly choose a more adequate model for the dataset.

7.0.4 BigQuery and Execution Time

- The BigQuery proved to be a reliable and intuitive platform to perform **Machine Learning** projects
- The queries that were timed (not all were because all the results would be very similar) executed almost instantly (range from < 1 second to ~3/4 seconds)
- All the steps for executing the **XGBoost Classifier** model were also (and surprisingly!) very fast, with all executions running almost instantly

7.0.5 Dask and Execution Time

- Dask can support very large datasets while maintaining very quick and efficient processing speeds.
- However, because of the way we treated the data, we weren't able to see many improvements when compared to other libraries like Pandas, with most executions ranging from 2~40 seconds.

7.0.6 Project Developed by:

- Francisco Macieira → up202207166@edu.fc.up.pt
- Manuel Silva → up202108874@edu.fe.up.pt
- Nuno Gomes → up202206195@edu.fc.up.pt